

HEP C++ course

D.Chamont, B. Gruber, S. Hageboeck, S. Ponce
sebastien.ponce@cern.ch

CERN

October 15, 2025



Foreword

What this course is not

- It is not for absolute beginners
- It is not for experts
- It is not complete at all (would need 3 weeks...)
 - although it is already too long for the time we have
 - 564 slides, 707 pages, 32 exercises...

How I see it

Adaptative pick what you want

Interactive tell me what to skip/insist on

Practical let's spend time on real code

Where to find latest version ?

- full sources at github.com/hsf-training/cpluspluscourse
- latest pdf on [GitHub](#)

More courses

The HSF Software Training Center

A set of course modules on more software engineering aspects prepared from within the HEP community

- Unix shell
- Python
- Version control (git, gitlab, github)
- ...

<https://hepsoftwarefoundation.org/training/curriculum.html>



Outline

- 1 History and goals
- 2 Language basics
- 3 Object orientation (OO)
- 4 Core modern C++
- 5 Expert C++
- 6 Useful tools
- 7 Concurrency
- 8 C++ and python



Detailed outline

- 1 History and goals
 - History
 - Why we use it?

- 2 Language basics
 - Core syntax and types
 - Arrays and Pointers
 - Scopes / namespaces
 - Class and enum types
 - References
 - Functions
 - Operators
 - Control structures
 - Headers and interfaces
 - Auto keyword
 - Inline keyword
 - Assertions

- 3 Object orientation (OO)
 - Objects and Classes
 - Inheritance
 - Constructors/destructors
 - Static members
 - Allocating objects
 - Advanced OO

- 4 Core modern C++
 - Constness
 - Constant Expressions
 - Exceptions
 - Move semantics
 - Copy elision
 - Templates
 - Lambdas
 - The STL
 - Standard library
 - Ranges
 - RAII and smart pointers
 - Initialization

- 5 Expert C++
 - The `<=>` operator
 - Variadic templates
 - Perfect forwarding
 - SFINAE
 - Concepts
 - Modules

- 6 Coroutines
 - Useful tools
 - C++ editor
 - Version control
 - Code formatting
 - The Compiling Chain
 - Web tools
 - Debugging
 - Sanitizers
 - The Valgrind family
 - Static code analysis
 - Profiling
 - Doxygen

- 7 Concurrency
 - Threads and async
 - Mutexes
 - Atomic types
 - Thread-local storage
 - Condition Variables

- 8 C++ and python
 - Writing a module
 - Marrying C++ and C
 - The ctypes module
 - The cppy project

History and goals

1 History and goals

- History
- Why we use it?

2 Language basics

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

6 Useful tools

7 Concurrency

8 C++ and python


History

1 History and goals

- History
- Why we use it?



C/C++ origins



C inventor
Dennis M. Ritchie



C++ inventor
Bjarne Stroustrup

- Both C and C++ are born in Bell Labs
- C++ *almost* embeds C
- C and C++ are still under development
- We will discuss all C++ specs up to C++ 20 (only partially)
- Each slide will be marked with first spec introducing the feature

C++ 11, C++ 14, C++ 17, C++ 20, C++ 23, C++ 26...

Status

- A new C++ specification every 3 years
 - C++ 23 complete since 11th of Feb. 2023
 - work on C++ 26 is ongoing
- Bringing each time a lot of goodies



C++ 11, C++ 14, C++ 17, C++ 20, C++ 23, C++ 26...

Status

- A new C++ specification every 3 years
 - C++ 23 complete since 11th of Feb. 2023
 - work on C++ 26 is ongoing
- Bringing each time a lot of goodies

How to use C++ XX features

- Use a compatible compiler
- add `-std=c++xx` to compilation flags
- e.g. `-std=c++17`

C++	gcc	clang
11	≥ 4.8	≥ 3.3
14	≥ 4.9	≥ 3.4
17	≥ 7.3	≥ 5
20	$\geq 11, 14$	≥ 19
23	> 15	> 20

Table: Minimum versions of gcc and clang for a given C++ version



Why we use it?

1 History and goals

- History
- Why we use it?



Why is C++ our language of choice?

Adapted to large projects

- statically and strongly typed
- object oriented
- widely used (and taught)
- many available libraries



Why is C++ our language of choice?

Adapted to large projects

- statically and strongly typed
- object oriented
- widely used (and taught)
- many available libraries

Fast

- compiled to native machine code
 - unlike Java, C#, Python, ...
- allows to go close to hardware when needed



Why is C++ our language of choice?

Adapted to large projects

- statically and strongly typed
- object oriented
- widely used (and taught)
- many available libraries

Fast

- compiled to native machine code
 - unlike Java, C#, Python, ...
- allows to go close to hardware when needed

What we get

- the most powerful language
- the most complicated one
- the most error prone?



Language basics

- **Inline keyword**
- **Assertions**

1 History and goals

2 Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

6 Useful tools

7 Concurrency

8 C++ and python



Core syntax and types

2 Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Hello World

C++ 98

```
1 #include <iostream>
2
3 // This is a function
4 void print(int i) {
5     std::cout << "Hello, world " << i << std::endl;
6 }
7
8 int main(int argc, char** argv) {
9     int n = 3;
10    for (int i = 0; i < n; i++) {
11        print(i);
12    }
13    return 0;
14 }
```



Comments

C++ 98

```
1 // simple comment until end of line
2 int i;
3
4 /* multiline comment
5  * in case we need to say more
6  */
7 double /* or something in between */ d;
8
9 /**
10  * Best choice : doxygen compatible comments
11  * \brief checks whether i is odd
12  * \param i input
13  * \return true if i is odd, otherwise false
14  * \see https://www.doxygen.nl/manual/docblocks.html
15 */
16 bool isOdd(int i);
```

Basic types(1)

C++ 98

```
1  bool b = true;           // boolean, true or false
2
3  char c = 'a';            // min 8 bit integer
4
5
6  signed char c = 4;       // min 8 bit signed integer
7  unsigned char c = 4;      // min 8 bit unsigned integer
8
9  char* s = "a C string";  // array of chars ended by \0
10 string t = "a C++ string"; // class provided by the STL
11
12 short int s = -444;       // min 16 bit signed integer
13 unsigned short s = 444;    // min 16 bit unsigned integer
14 short s = -444;          // int is optional
```



Basic types(2)

C++ 98

```
1 int i = -123456;           // min 16, usually 32 bit
2 unsigned int i = 1234567; // min 16, usually 32 bit
3
4 long l = 0L               // min 32 bit
5 unsigned long l = 0UL;    // min 32 bit
6
7 long long ll = 0LL;       // min 64 bit
8 unsigned long long l = 0ULL; // min 64 bit
9
10 float f = 1.23f;         // 32 (1+8+23) bit float
11 double d = 1.23E34;      // 64 (1+11+52) bit float
12 long double ld = 1.23E34L // min 64 bit float
```



Portable numeric types

C++ 98

```
1 #include <cstdint> // defines the following:  
2  
3 std::int8_t c = -3;           // 8 bit signed integer  
4 std::uint8_t c = 4;          // 8 bit unsigned integer  
5  
6 std::int16_t s = -444;       // 16 bit signed integer  
7 std::uint16_t s = 444;        // 16 bit unsigned integer  
8  
9 std::int32_t s = -674;       // 32 bit signed integer  
10 std::uint32_t s = 674;        // 32 bit unsigned integer  
11  
12 std::int64_t s = -1635;      // 64 bit signed integer  
13 std::uint64_t s = 1635;       // 64 bit unsigned int
```



Integer literals

C++ 98

```
1 int i = 1234;           // decimal      (base 10)
2 int i = 02322;          // octal        (base 8)
3 int i = 0x4d2;          // hexadecimal (base 16)
4 int i = 0X4D2;          // hexadecimal (base 16)
5 int i = 0b10011010010;  // binary       (base 2) C++14
6
7 int i = 123'456'789;    // digit separators, C++14
8 int i = 0b100'1101'0010; // digit separators, C++14
9
10 42                  // int
11 42u,   42U           // unsigned int
12 42l,   42L           // long
13 42ul,  42UL          // unsigned long
14 42ll,  42LL          // long long
15 42ull, 42ULL         // unsigned long long
```

Floating-point literals

C++ 98

```
1 double d = 12.34;
2 double d = 12. ;
3 double d = .34;
4 double d = 12e34;           // 12 * 10^34
5 double d = 12E34;          // 12 * 10^34
6 double d = 12e-34;         // 12 * 10^-34
7 double d = 12.34e34;       // 12.34 * 10^34
8
9 double d = 123'456.789'101; // digit separators, C++14
10
11 double d = 0x4d2.4p3;     // hexfloat, 0x4d2.4 * 2^3
12                           // = 1234.25 * 2^3 = 9874
13
14 3.14f, 3.14F,    // float
15 3.14,   3.14,    // double
16 3.14l, 3.14L,    // long double
```

Useful aliases

C++ 11 Adv

```
1 #include <cstddef> // (and others) defines:  
2  
3 // unsigned integer, can hold any variable's size  
4 std::size_t s = sizeof(int);  
5  
6 #include <cstdint> // defines:  
7  
8 // signed integer, can hold any diff between two pointers  
9 std::ptrdiff_t c = &s - &s;  
10  
11 // signed/unsigned integer, can hold any pointer value  
12 std::intptr_t i = reinterpret_cast<intptr_t>(&s);  
13 std::uintptr_t i = reinterpret_cast<uintptr_t>(&s);
```



Extended floating-point types

C++ 23 Adv

Extended floating-point types

- *Optional* types, which may be provided
- Custom conversion and promotion rules

```
1 #include <stdfloat> // may define these:  
2  
3 std::float16_t = 3.14f16; // 16 (1+5+10) bit float  
4 std::float32_t = 3.14f32; // like float (1+8+23)  
5 // but different type  
6 std::float64_t = 3.14f64; // like double (1+11+52)  
7 // but different type  
8 std::float128_t = 3.14f128; // 128 (1+15+112) bit float  
9 std::bfloat16_t = 3.14bf16; // 16 (1+8+7) bit float  
10  
11 // also F16, F32, F64, F128 or BF16 suffix possible
```



Arrays and Pointers

2 Language basics

- Core syntax and types
- **Arrays and Pointers**
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Static arrays

C++ 98

```
1 int ai[4] = {1,2,3,4};  
2 int ai[] = {1,2,3,4}; // identical  
3  
4 char ac[3] = {'a','b','c'}; // char array  
5 char ac[4] = "abc"; // valid C string  
6 char ac[4] = {'a','b','c',0}; // same valid string  
7  
8 int i = ai[2]; // i = 3  
9 char c = ac[8]; // at best garbage, may segfault  
10 int i = ai[4]; // also garbage !
```



Pointers

C++ 98

```
1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;
```

Pointers

C++ 98

```
1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = pai + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;
```

Memory layout

	0x3028
	0x3024
	0x3020
	0x301C
	0x3018
	0x3014
	0x3010
	0x300C
	0x3008
	0x3004
i = 4	0x3000

Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = pai + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
	0x3020
	0x301C
	0x3018
	0x3014
	0x3010
	0x300C
	0x3008
pi = 0x3000	0x3004
i = 4	0x3000



Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = pai + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
	0x3020
	0x301C
	0x3018
	0x3014
	0x3010
	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000




Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
	0x3020
	0x301C
	0x3018
ai[2] = 3	0x3014
ai[1] = 2	0x3010
ai[0] = 1	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000



Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
	0x3020
	0x301C
pai = 0x300C	0x3018
ai[2] = 3	0x3014
ai[1] = 2	0x3010
ai[0] = 1	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000

Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
	0x3020
paj = 0x3010	0x301C
pai = 0x300C	0x3018
ai[2] = 3	0x3014
ai[1] = 2	0x3010
ai[0] = 1	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000

Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

	0x3028
	0x3024
k = 3	0x3020
paj = 0x3010	0x301C
pai = 0x300C	0x3018
ai[2] = 3	0x3014
ai[1] = 2	0x3010
ai[0] = 1	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000

Pointers

C++ 98

```

1 int i = 4;
2 int *pi = &i;
3 int j = *pi + 1;
4
5 int ai[] = {1,2,3};
6 int *pai = ai; // decay to ptr
7 int *paj = paj + 1;
8 int k = *paj + 1;
9
10 // compile error
11 int *pak = k;
12
13 // seg fault !
14 int *pak = (int*)k;
15 int l = *pak;

```

Memory layout

?? ←	0x3028
pak = 3	0x3024
k = 3	0x3020
paj = 0x3010	0x301C
pai = 0x300C	0x3018
ai[2] = 3	0x3014
ai[1] = 2	0x3010
ai[0] = 1	0x300C
j = 5	0x3008
pi = 0x3000	0x3004
i = 4	0x3000

A pointer to nothing

- if a pointer doesn't point to anything, set it to `nullptr`
 - useful to e.g. mark the end of a linked data structure
 - or absence of an optional function argument (pointer)
- same as setting it to 0 or `NULL` (before C++ 11)
- triggers compilation error when assigned to integer



A pointer to nothing

- if a pointer doesn't point to anything, set it to `nullptr`
 - useful to e.g. mark the end of a linked data structure
 - or absence of an optional function argument (pointer)
- same as setting it to 0 or `NULL` (before C++ 11)
- triggers compilation error when assigned to integer

Example code

```
1 int* ip = nullptr;
2 int i = NULL;      // compiles, bug?
3 int i = nullptr;   // ERROR
```



Dynamic arrays using C

C++ 98

```
1 #include <cstdlib>
2 #include <cstring>
3
4 int *bad;           // pointer to random address
5 int *ai = nullptr; // better, deterministic, testable
6
7 // allocate array of 10 ints (uninitialized)
8 ai = (int*) malloc(10*sizeof(int));
9 memset(ai, 0, 10*sizeof(int)); // and set them to 0
10
11 ai = (int*) calloc(10, sizeof(int)); // both in one go
12
13 free(ai); // release memory
```

Good practice: Don't use C's memory management

Use `std::vector` and friends or smart pointers

Manual dynamic arrays using C++

C++ 98

```
1 #include <cstdlib>
2 #include <cstring>
3
4 // allocate array of 10 ints
5 int* ai = new int[10];    // uninitialized
6 int* ai = new int[10]{};   // zero-initialized
7
8 delete[] ai; // release array memory
9
10 // allocate a single int
11 int* pi = new int;
12 int* pi = new int{};
13 delete pi; // release scalar memory
```

Good practice: Don't use manual memory management

Use `std::vector` and friends or smart pointers

Scopes / namespaces

2 Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Definition

Portion of the source code where a given name is valid

Typically :

- simple block of code, within {}
- function, class, namespace
- the global scope, i.e. translation unit (.cpp file + all includes)

Example

```
1 { int a;  
2   { int b;  
3     } // end of b scope  
4 } // end of a scope
```



Scope and lifetime of variables

C++ 98

Variable life time


- Variables are (statically) allocated when defined
- Variables are freed at the end of a scope

Good practice: Initialisation

- Initialise variables when allocating them!
- This prevents bugs reading uninitialized memory

```
1 int a = 1;  
2 {  
3     int b[4];  
4     b[0] = a;  
5 }  
6 // Doesn't compile here:  
7 // b[1] = a + 1;
```

Memory layout



0x3010
0x300C
0x3008
0x3004
0x3000
a = 1



Scope and lifetime of variables

C++ 98

Variable life time

- Variables are (statically) allocated when defined
- Variables are freed at the end of a scope

Good practice: Initialisation

- Initialise variables when allocating them!
- This prevents bugs reading uninitialised memory

```
1 int a = 1;  
2 {  
3     int b[4];  
4     b[0] = a;  
5 }  
6 // Doesn't compile here:  
7 // b[1] = a + 1;
```

Memory layout

b[3] = ?
b[2] = ?
b[1] = ?
b[0] = ?
a = 1

0x3010
0x300C
0x3008
0x3004
0x3000



Scope and lifetime of variables

C++ 98

Variable life time

- Variables are (statically) allocated when defined
- Variables are freed at the end of a scope

Good practice: Initialisation

- Initialise variables when allocating them!
- This prevents bugs reading uninitialised memory

```
1 int a = 1;  
2 {  
3     int b[4];  
4     b[0] = a;  
5 }  
6 // Doesn't compile here:  
7 // b[1] = a + 1;
```

Memory layout

b[3] = ?
b[2] = ?
b[1] = ?
b[0] = 1
a = 1

0x3010
0x300C
0x3008
0x3004
0x3000



Scope and lifetime of variables

C++ 98

Variable life time

- Variables are (statically) allocated when defined
- Variables are freed at the end of a scope

Good practice: Initialisation

- Initialise variables when allocating them!
- This prevents bugs reading uninitialised memory

```
1 int a = 1;  
2 {  
3     int b[4];  
4     b[0] = a;  
5 }  
6 // Doesn't compile here:  
7 // b[1] = a + 1;
```

Memory layout

?	0x3010
?	0x300C
?	0x3008
1	0x3004
a = 1	0x3000



Namespaces

C++ 98

- Namespaces allow to segment your code to avoid name clashes
- They can be embedded to create hierarchies (separator is '::')

```

1  int a;                                14  namespace p { // reopen p
2  namespace n {                           15    void f() {
3    int a; // no clash                  16      p::a = 6;
4  }                                         17      a = 6; //same as above
5  namespace p {                           18      ::a = 1;
6    int a; // no clash                  19      p::inner::a = 8;
7    namespace inner {                   20      inner::a = 8;
8      int a; // no clash                21      n::a = 3;
9    }                                     22    }
10 }                                       23  }
11 void f() {                            24  using namespace p::inner;
12   n::a = 3;                           25  void g() {
13 }                                     26    a = -1; // err: ambiguous
14                                         27  }

```



Nested namespaces

C++ 17

C++ 98: Old way to declare nested namespaces

```
1 namespace A {  
2     namespace B {  
3         namespace C {  
4             //...  
5         }  
6     }  
7 }
```

C++ 17: Nested declaration

```
1 namespace A::B::C {  
2     //...  
3 }
```

C++ 17: Namespace alias

```
1 namespace ABC = A::B::C;
```



Unnamed / anonymous namespaces

C++ 98 Adv

A namespace without a name!

```
1 namespace {  
2     int localVar;  
3 }
```

Purpose

- groups a number of declarations
- visible only in the current translation unit
- but not reusable outside
- allows much better compiler optimizations and checking
 - e.g. unused function warning
 - context dependent optimizations

Supersedes static

```
4 static int localVar; // equivalent C code
```



Using namespace directives

C++ 98 Adv

Avoid “using namespace” directives

- Make all members of a namespace visible in current scope
- Risk of name clashes or ambiguities

```
1 using namespace std;  
2 cout << "We can print now\n"; // uses std::cout
```

Never use in headers at global scope!

```
1 #include "PoorlyWritten.h" // using namespace std;  
2 struct array { ... };  
3 array a; // Error: name clash with std::array
```

What to do instead

- Qualify names: std::vector, std::cout, ...
- Put things that belong together in the same namespace
- Use *using declarations* in local scopes: `using std::cout;`



Class and enum types

2 Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- **Class and enum types**
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



struct

C++ 98

“members” grouped together under one name

```
1 struct Individual {           14 Individual *ptr = &student;
2     unsigned char age;        15 ptr->age = 25;
3     float weight;            16 // same as: (*ptr).age = 25;
4 };
5
6 Individual student;
7 student.age = 25;
8 student.weight = 78.5f;
9
10 Individual teacher = {
11     45, 67.0f
12 };
```

struct

C++ 98

“members” grouped together under one name

```
1 struct Individual {           14 Individual *ptr = &student;
2     unsigned char age;        15 ptr->age = 25;
3     float weight;           16 // same as: (*ptr).age = 25;
4 };
5
6 Individual student;
7 student.age = 25;
8 student.weight = 78.5f;
9
10 Individual teacher = {
11     45, 67.0f
12 };
```

Memory layout

	0x3010
	0x300C
	0x3008
	0x3004
	0x3000



struct

C++ 98

“members” grouped together under one name


```

1 struct Individual {
2     unsigned char age;
3     float weight;
4 };
5
6 Individual student;
7 student.age = 25;
8 student.weight = 78.5f;
9
10 Individual teacher = {
11     45, 67.0f
12 };

```

14 Individual *ptr = &student;
 15 ptr->age = 25;
 16 // same as: (*ptr).age = 25;

Memory layout



struct

C++ 98

“members” grouped together under one name

```

1 struct Individual {
2     unsigned char age;
3     float weight;
4 };
5
6 Individual student;
7 student.age = 25;
8 student.weight = 78.5f;
9
10 Individual teacher = {
11     45, 67.0f
12 };
13
14 Individual *ptr = &student;
15 ptr->age = 25;
16 // same as: (*ptr).age = 25;

```

Memory layout

student teacher					0x3010
	←	67.0	→		0x300C
	45	?	?	?	0x3008
	←	78.5	→		0x3004
	25	?	?	?	0x3000



“members” grouped together under one name

```

1 struct Individual {
2     unsigned char age;
3     float weight;
4 };
5
6 Individual student;
7 student.age = 25;
8 student.weight = 78.5f;
9
10 Individual teacher = {
11     45, 67.0f
12 };
13
14 Individual *ptr = &student;
15 ptr->age = 25;
16 // same as: (*ptr).age = 25;

```

Memory layout

student teacher	←	0x3000	→	0x3010
	←	67.0	→	0x300C
	45	?	?	?
	←	78.5	→	0x3008
	25	?	?	?
	←		→	0x3004
	←		→	0x3000



union

C++ 98

“members” packed together at same memory location

```
1 union Duration {  
2     int seconds;  
3     short hours;  
4     char days;  
5 };  
6 Duration d1, d2, d3;  
7 d1.seconds = 259200;  
8 d2.hours = 72;  
9 d3.days = 3;  
10 d1.days = 3; // d1.seconds overwritten  
11 int a = d1.seconds; // d1.seconds is garbage
```



union

C++ 98

“members” packed together at same memory location

```
1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage
```

Memory layout

	0x300C
	0x3008
	0x3004
	0x3000




union

C++ 98

“members” packed together at same memory location

```
1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage
```

Memory layout



union

C++ 98

“members” packed together at same memory location

```
1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage
```

Memory layout

				0x300C
				0x3008
←	d2	72	→	?
←	—	d1	259200	—→

union

C++ 98

“members” packed together at same memory location

```
1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage
```

Memory layout

					0x300C	
d3	3	?	?	?	0x3008	
←	d2	72	→	?	?	0x3004
←	—	d1	259200	—	→	0x3000



union

C++ 98

“members” packed together at same memory location

```
1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage
```

Memory layout

					0x300C	
d3	3	?	?	?	0x3008	
←	d2	72	→	?	?	0x3004
d1	3	?	?	?	0x3000	



“members” packed together at same memory location

```

1 union Duration {
2     int seconds;
3     short hours;
4     char days;
5 };
6 Duration d1, d2, d3;
7 d1.seconds = 259200;
8 d2.hours = 72;
9 d3.days = 3;
10 d1.days = 3; // d1.seconds overwritten
11 int a = d1.seconds; // d1.seconds is garbage

```

Memory layout

					0x300C
d3	3	?	?	?	0x3008
←	d2	72	→	?	?
d1	3	?	?	?	0x3004
					0x3000

Good practice: Avoid unions

- Starting with C++ 17: prefer std::variant



Enums

C++ 98

- use to declare a list of related constants (enumerators)
- has an underlying integral type
- enumerator names leak into enclosing scope

```
1 enum VehicleType {  
2     BIKE,    // 0  
3     CAR,     // 1  
4     BUS,     // 2  
5 };  
6 VehicleType t = CAR;
```

```
8 enum VehicleType  
9     : int { // C++11  
10    BIKE = 3,  
11    CAR = 5,  
12    BUS = 7,  
13 };  
14 VehicleType t2 = BUS;
```

Scoped enumeration, aka enum class

C++ 11

Same syntax as enum, with scope

```
1 enum class VehicleType { Bus, Car };  
2 VehicleType t = VehicleType::Car;
```



Scoped enumeration, aka enum class

C++ 11

Same syntax as enum, with scope

```
1 enum class VehicleType { Bus, Car };  
2 VehicleType t = VehicleType::Car;
```

Only advantages

- scopes enumerator names, avoids name clashes
- strong typing, no automatic conversion to int

```
3 enum VType { Bus, Car }; enum Color { Red, Blue };  
4 VType t = Bus;  
5 if (t == Red) { /* We do enter */ }  
6 int a = 5 * Car; // Ok, a = 5  
7  
8 enum class VT { Bus, Car }; enum class Col { Red, Blue };  
9 VT t = VT::Bus;  
10 if (t == Col::Red) { /* Compiler error */ }  
11 int a = t * 5; // Compiler error
```

More sensible example

C++ 11

```
1 enum class ShapeType {
2     Circle,
3     Rectangle
4 };
5
6 struct Rectangle {
7     float width;
8     float height;
9 };
```



More sensible example

C++ 11

```
1 enum class ShapeType {           10 struct Shape {
2     Circle,                      11     ShapeType type;
3     Rectangle                     12     union {
4 };                                13         float radius;
5                               14         Rectangle rect;
6 struct Rectangle {              15     };
7     float width;                16 };
8     float height;               9 };
```



More sensible example

C++ 11

```
1 enum class ShapeType {           10 struct Shape {
2     Circle,                      11     ShapeType type;
3     Rectangle                     12     union {
4 };                                13         float radius;
5                               14         Rectangle rect;
6 struct Rectangle {              15     };
7     float width;                 16 };
8     float height;                17
9 };                                18
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
```



typedef and using

C++ 98 / C++ 11

Used to create type aliases

C++ 98

```
1 typedef std::uint64_t myint;  
2 myint count = 17;  
3 typedef float position[3];
```

C++ 11

```
4 using myint = std::uint64_t;  
5 myint count = 17;  
6 using position = float[3];  
7  
8 template <typename T> using myvec = std::vector<T>;  
9 myvec<int> myintvec;
```



References

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



References

C++ 98

References

- References allow for direct access to another object
- They can be used as shortcuts / better readability
- They can be declared **const** to allow only read access

Example:

```
1 int i = 2;
2 int &iref = i; // access to i
3 iref = 3;      // i is now 3
4
5 // const reference to a member:
6 struct A { int x; int y; } a;
7 const int &x = a.x; // direct read access to A's x
8 x = 4;            // doesn't compile
9 a.x = 4;          // fine
```



Pointers vs References

C++ 98

Specificities of reference

- Natural syntax
- Cannot be `nullptr`
- Must be assigned when defined, cannot be reassigned
- References to temporary objects must be `const`

Advantages of pointers

- Can be `nullptr`
- Can be initialized after declaration, can be reassigned



Pointers vs References

C++ 98

Specificities of reference

- Natural syntax
- Cannot be `nullptr`
- Must be assigned when defined, cannot be reassigned
- References to temporary objects must be `const`

Advantages of pointers

- Can be `nullptr`
- Can be initialized after declaration, can be reassigned

Good practice: References

- Prefer using references instead of pointers
- Mark references `const` to prevent modification

Functions

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- **Functions**
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Functions

C++ 98

```
1 // with return type           11 // no return
2 int square(int a) {          12 void log(char* msg) {
3     return a * a;            13     std::cout << msg;
4 }                           14 }
5                               15
6 // multiple parameters       16 // no parameter
7 int mult(int a,             17 void hello() {
8     int b) {                18     std::cout << "Hello World";
9     return a * b;            19 }
```



Functions

C++ 98

```
1 // with return type           11 // no return
2 int square(int a) {          12 void log(char* msg) {
3     return a * a;             13     std::cout << msg;
4 }                           14 }
5                               15
6 // multiple parameters       16 // no parameter
7 int mult(int a,              17 void hello() {
8     int b) {                  18     std::cout << "Hello World";
9     return a * b;             19 }
```

Functions and references to returned values

```
1 int result = square(2);
2 int & temp = square(2);        // Not allowed
3 int const & temp2 = square(2); // OK
```



Function default arguments

C++ 98

```
1 // must be the trailing    11 // multiple default
2 // argument                12 // arguments are possible
3 int add(int a,           13 int add(int a = 2,
4         int b = 2) {        14             int b = 2) {
5     return a + b;          15         return a + b;
6 }                           16     }
7 // add(1) == 3            17 // add() == 4
8 // add(3,4) == 7          18 // add(3) == 5
9
```



Functions: parameters are passed by value

C++ 98

```
1 struct BigStruct {....};  
2 BigStruct s;  
3  
4 // parameter by value  
5 void printVal(BigStruct p) {  
6     ...  
7 }  
8 printVal(s); // copy  
9  
10 // parameter by reference  
11 void printRef(BigStruct &q) {  
12     ...  
13 }  
14 printRef(s); // no copy
```

Memory layout

0x31E0
0x3190
0x3140
0x30F0
0x30A0
0x3050
0x3000




Functions: parameters are passed by value

C++ 98

```
1 struct BigStruct {....};  
2 BigStruct s;  
3  
4 // parameter by value  
5 void printVal(BigStruct p) {  
6     ...  
7 }  
8 printVal(s); // copy  
9  
10 // parameter by reference  
11 void printRef(BigStruct &q) {  
12     ...  
13 }  
14 printRef(s); // no copy
```

Memory layout



Functions: parameters are passed by value

C++ 98

```
1 struct BigStruct {....};  
2 BigStruct s;  
3  
4 // parameter by value  
5 void printVal(BigStruct p) {  
6     ...  
7 }  
8 printVal(s); // copy  
9  
10 // parameter by reference  
11 void printRef(BigStruct &q) {  
12     ...  
13 }  
14 printRef(s); // no copy
```

Memory layout

p	{	0x31E0
	pn = sn	0x3190
	...	0x3140
p	{	0x30F0
	p1 = s1	0x30A0
s	{	0x3050
	sn	0x3000
	...	
	s1	



Functions: parameters are passed by value

C++ 98

```
1 struct BigStruct {....};  
2 BigStruct s;  
3  
4 // parameter by value  
5 void printVal(BigStruct p) {  
6     ...  
7 }  
8 printVal(s); // copy  
9  
10 // parameter by reference  
11 void printRef(BigStruct &q) {  
12     ...  
13 }  
14 printRef(s); // no copy
```

Memory layout

	0x31E0
	0x3190
	0x3140
q = 0x3000	0x30F0
sn	0x30A0
...	0x3050
s1	0x3000



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
	0x3004
	0x3000



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
	0x3004
s.a = 1	0x3000



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
p.a = 1	0x3004
s.a = 1	0x3000



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
p.a = 2	0x3004
s.a = 1	0x3000



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
	0x3004
s.a = 1	0x3000




Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout




Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout



Functions: pass by value or reference?

C++ 98

```
1 struct SmallStruct {int a;};
2 SmallStruct s = {1};
3
4 void changeVal(SmallStruct p) {
5     p.a = 2;
6 }
7 changeVal(s);
8 // s.a == 1
9
10 void changeRef(SmallStruct &q) {
11     q.a = 2;
12 }
13 changeRef(s);
14 // s.a == 2
```

Memory layout

	0x3008
	0x3004
s.a = 2	0x3000



Pass by value, reference or pointer

C++ 98

Different ways to pass arguments to a function

- By default, arguments are passed by value (= copy)
good for small types, e.g. numbers
- Use references for parameters to avoid copies
good for large types, e.g. objects
- Use **const** for safety and readability whenever possible



Pass by value, reference or pointer

C++ 98

Different ways to pass arguments to a function

- By default, arguments are passed by value (= copy)
good for small types, e.g. numbers
- Use references for parameters to avoid copies
good for large types, e.g. objects
- Use **const** for safety and readability whenever possible

Syntax

```
1 struct T {...}; T a;
2 void fVal(T value);           fVal(a);    // by value
3 void fRef(const T &value);   fRef(a);    // by reference
4 void fPtr(const T *value);   fPtr(&a);   // by pointer
5 void fWrite(T &value);      fWrite(a);  // non-const ref
```



Overloading

C++ 98

Overloading

- We can have multiple functions with the same name
 - Must have different parameter lists
 - A different return type alone is not allowed
 - Form a so-called “overload set”
- Default arguments can cause ambiguities

```
1 int sum(int b);           // 1
2 int sum(int b, int c);    // 2, ok, overload
3 // float sum(int b, int c); // disallowed
4 sum(42); // calls 1
5 sum(42, 43); // calls 2
6 int sum(int b, int c, int d = 4); // 3, overload
7 sum(42, 43, 44); // calls 3
8 sum(42, 43);      // error: ambiguous, 2 or 3
```



Functions

C++ 98

Exercise: Functions

Familiarise yourself with pass by value / pass by reference.

- Go to exercises/functions
- Look at functions.cpp
- Compile it (make) and run the program (./functions)
- Work on the tasks that you find in functions.cpp



Functions: good practices

C++ 98

Good practice: Write readable functions

- Keep functions short
- Do one logical thing (single-responsibility principle)
- Use expressive names
- Document non-trivial functions

Example: Good

```
1 // Count number of dilepton events in data.  
2 // \param d Dataset to search.  
3 unsigned int countDileptons(Data &d) {  
4     selectEventsWithMuons(d);  
5     selectEventsWithElectrons(d);  
6     return d.size();  
7 }
```



Functions: good practices

C++ 98

Example: don't! Everything in one long function

```
1  unsigned int runJob() { 15      if (...) {
2    // Step 1: data           16          data.erase(...);
3    Data data;              17      }
4    data.resize(123456);    18      }
5    data.fill(...);        19
6                                20      // Step 4: dileptons
7    // Step 2: muons         21      int counter = 0;
8    for (...) {            22      for (...) {
9      if (...) {           23          if (...) {
10        data.erase(...);   24          counter++;
11      }                   25      }
12    }                     26      }
13    // Step 3: electrons    27
14    for (...) {           28      return counter;
15      }                   29    }
```



Operators

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Operators(1)

C++ 98

Binary and Assignment Operators

```
1 int i = 1 + 4 - 2;    // 3
2 i *= 3;                // 9, short for: i = i * 3;
3 i /= 2;                // 4
4 i = 23 % i;            // modulo => 3
```



Operators(1)

C++ 98

Binary and Assignment Operators

```
1 int i = 1 + 4 - 2;    // 3
2 i *= 3;                // 9, short for: i = i * 3;
3 i /= 2;                // 4
4 i = 23 % i;            // modulo => 3
```

Increment / Decrement Operators

```
1 int i = 0; i++; // i = 1
2 int j = ++i; // i = 2, j = 2
3 int k = i++; // i = 3, k = 2
4 int l = --i; // i = 2, l = 2
5 int m = i--; // i = 1, m = 2
```



Operators(1)

C++ 98

Binary and Assignment Operators

```
1 int i = 1 + 4 - 2;    // 3
2 i *= 3;                // 9, short for: i = i * 3;
3 i /= 2;                // 4
4 i = 23 % i;            // modulo => 3
```

Increment / Decrement Operators

Use wisely

```
1 int i = 0; i++; // i = 1
2 int j = ++i; // i = 2, j = 2
3 int k = i++; // i = 3, k = 2
4 int l = --i; // i = 2, l = 2
5 int m = i--; // i = 1, m = 2
```



Operators(2)

C++ 98

Bitwise and Assignment Operators

```
1 unsigned i = 0xee & 0x55;      // 0x44
2 i |= 0xee;                      // 0xee
3 i ^= 0x55;                      // 0xbb
4 unsigned j = ~0xee;            // 0xffffffff11
5 unsigned k = 0x1f << 3;        // 0xf8
6 unsigned l = 0x1f >> 2;        // 0x7
```



Operators(2)

C++ 98

Bitwise and Assignment Operators

```
1 unsigned i = 0xee & 0x55;    // 0x44
2 i |= 0xee;                  // 0xee
3 i ^= 0x55;                  // 0xbb
4 unsigned j = ~0xee;         // 0xffffffff11
5 unsigned k = 0x1f << 3;     // 0xf8
6 unsigned l = 0x1f >> 2;     // 0x7
```

Logical Operators

```
1 bool a = true;
2 bool b = false;
3 bool c = a && b;      // false
4 bool d = a || b;       // true
5 bool e = !d;          // false
```

Operators(3)

C++ 98

Comparison Operators

```
1 bool a = (3 == 3); // true
2 bool b = (3 != 3); // false
3 bool c = (4 < 4); // false
4 bool d = (4 <= 4); // true
5 bool e = (4 > 4); // false
6 bool f = (4 >= 4); // true
7 auto g = (5 <=> 5); // C++20 (later)
```



Operators(3)

C++ 98

Comparison Operators

```
1 bool a = (3 == 3); // true
2 bool b = (3 != 3); // false
3 bool c = (4 < 4); // false
4 bool d = (4 <= 4); // true
5 bool e = (4 > 4); // false
6 bool f = (4 >= 4); // true
7 auto g = (5 <=> 5); // C++20 (later)
```

Precedences

```
c &= 1+(++b) | (a--) * 4%5^7; // ???
```

Details can be found on [cppreference](#)



Operators(3)

C++ 98

Comparison Operators

```
1 bool a = (3 == 3); // true
2 bool b = (3 != 3); // false
3 bool c = (4 < 4); // false
4 bool d = (4 <= 4); // true
5 bool e = (4 > 4); // false
6 bool f = (4 >= 4); // true
7 auto g = (5 <=> 5); // C++20 (later)
```

Precedences

Avoid

c &= 1+(++b) | (a--) * 4%5^7; // ???

Details can be found on [cppreference](#)



Operators(3)

C++ 98

Comparison Operators

```
1 bool a = (3 == 3); // true
2 bool b = (3 != 3); // false
3 bool c = (4 < 4); // false
4 bool d = (4 <= 4); // true
5 bool e = (4 > 4); // false
6 bool f = (4 >= 4); // true
7 auto g = (5 <=> 5); // C++20 (later)
```

Precedences

Avoid - use parentheses

```
c &= 1+(++b) | (a--) * 4%5^7; // ???
```

Details can be found on [cppreference](#)



Control structures

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- **Control structures**
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Control structures: if

C++ 98

if syntax

```
1 if (condition1) {  
2     Statement1; Statement2;  
3 } else if (condition2)  
4     OnlyOneStatement;  
5 else {  
6     Statement3;  
7     Statement4;  
8 }
```

- The **else** and **else if** clauses are optional
- The **else if** clause can be repeated
- Braces are optional if there is a single statement



Control structures: if

C++ 98

Practical example

```
1 int collatz(int a) {
2     if (a <= 0) {
3         std::cout << "not supported\n";
4         return 0;
5     } else if (a == 1) {
6         return 1;
7     } else if (a%2 == 0) {
8         return collatz(a/2);
9     } else {
10        return collatz(3*a+1);
11    }
12 }
```



Control structures: conditional operator

C++ 98

Syntax

```
test ? expression1 : expression2;
```

- If test is **true** expression1 is returned
- Else, expression2 is returned



Control structures: conditional operator

C++ 98

Syntax

```
test ? expression1 : expression2;
```

- If test is **true** expression1 is returned
- Else, expression2 is returned

Practical example

```
1 const int charge = isLepton ? -1 : 0;
```



Control structures: conditional operator

C++ 98

Syntax

```
test ? expression1 : expression2;
```

- If test is **true** expression1 is returned
- Else, expression2 is returned

Practical example

```
1 const int charge = isLepton ? -1 : 0;
```

Do not abuse it

```
1 int collatz(int a) {  
2     return a==1 ? 1 : collatz(a%2==0 ? a/2 : 3*a+1);  
3 }
```

- Explicit **ifs** are generally easier to read
- Use the ternary operator with short conditions and expressions
- Avoid nesting



Control structures: switch

C++ 98

Syntax

```
1  switch(identifier) {  
2      case c1 : statements1; break;  
3      case c2 : statements2; break;  
4      case c3 : statements3; break;  
5      ...  
6      default : statementsn; break;  
7  }
```

- The **break** statement is not mandatory but...
- Cases are entry points, not independent pieces
- Execution “falls through” to the next case without a **break!**
- The **default** case may be omitted



Control structures: switch

C++ 98

Syntax

```
1  switch(identifier) {  
2      case c1 : statements1; break;  
3      case c2 : statements2; break;  
4      case c3 : statements3; break;  
5      ...  
6      default : statementsn; break;  
7  }
```

- The **break** statement is not mandatory but...
- Cases are entry points, not independent pieces
- Execution “falls through” to the next case without a **break!**
- The **default** case may be omitted

Use break

Avoid **switch** statements with fall-through cases



Control structures: switch

C++ 98

Practical example

```
1 enum class Lang { French, German, English, Other };
2 Lang language = ...;
3 switch (language) {
4     case Lang::French:
5         std::cout << "Bonjour";
6         break;
7     case Lang::German:
8         std::cout << "Guten Tag";
9         break;
10    case Lang::English:
11        std::cout << "Good morning";
12        break;
13    default:
14        std::cout << "I do not speak your language";
15 }
```

[[fallthrough]] attribute

C++ 17

New compiler warning

Since C++ 17, compilers are encouraged to warn on fall-through

C++ 17

```
1 switch (c) {
2     case 'a':
3         f();      // Warning emitted
4     case 'b': // Warning probably suppressed
5     case 'c':
6         g();
7         [[fallthrough]]; // Warning suppressed
8     case 'd':
9         h();
10 }
```



Init-statements for if and switch

C++ 17

Purpose

Allows to limit variable scope in **if** and **switch** statements

C++ 17

```
1 if (Value val = GetValue(); condition(val)) {  
2     f(val); // ok  
3 } else  
4     g(val); // ok  
5 h(val); // error, no `val` in scope here
```



Init-statements for if and switch

C++ 17

Purpose

Allows to limit variable scope in **if** and **switch** statements

C++ 17

```
1 if (Value val = GetValue(); condition(val)) {  
2     f(val); // ok  
3 } else  
4     g(val); // ok  
5 h(val); // error, no `val` in scope here
```

C++ 98

Don't confuse with a variable declaration as condition:

```
7 if (Value* val = GetValuePtr())  
8     f(*val);
```



Control structures: for loop

C++ 98

for loop syntax

```
1 for(initializations; condition; increments) {  
2     statements;  
3 }
```

- Multiple initializations / increments are comma separated
- Initializations can contain declarations
- Braces are optional if loop body is a single statement



Control structures: for loop

C++ 98

for loop syntax

```
1 for(initializations; condition; increments) {  
2     statements;  
3 }
```

- Multiple initializations / increments are comma separated
- Initializations can contain declarations
- Braces are optional if loop body is a single statement

Practical example

```
4 for(int i = 0, j = 0 ; i < 10 ; i++, j = i*i) {  
5     std::cout << i << "^2 is " << j << '\n';  
6 }
```



Control structures: for loop

C++ 98

for loop syntax

```
1 for(initializations; condition; increments) {  
2     statements;  
3 }
```

- Multiple initializations / increments are comma separated
- Initializations can contain declarations
- Braces are optional if loop body is a single statement

Practical example

```
4 for(int i = 0, j = 0 ; i < 10 ; i++, j = i*i) {  
5     std::cout << i << "^2 is " << j << '\n';  
6 }
```

Good practice: Don't abuse the for syntax

- The **for** loop head should fit in 1-3 lines

Range-based loops

C++ 11

Reason of being

- Simplifies loops over “ranges” tremendously
- Especially with STL containers and ranges

Syntax

```
1 for ( type iteration_variable : range ) {  
2     // body using iteration_variable  
3 }
```

Example code

```
4 int v[4] = {1,2,3,4};  
5 int sum = 0;  
6 for (int a : v) { sum += a; }
```



Init-statements for range-based loops

C++ 20

Purpose

Allows to limit variable scope in range-based loops

C++ 17

```
1 std::array data = {"hello", ", ", "world"};
2 std::size_t i = 0;
3 for (auto& d : data) {
4     std::cout << i++ << ' ' << d << '\n';
5 }
```

C++ 20

```
6 for (std::size_t i = 0; auto const & d : data) {
7     std::cout << i++ << ' ' << d << '\n';
8 }
```



Control structures: while loop

C++ 98

while loop syntax

```
1  while(condition) {  
2      statements;  
3  }  
  
4  
  
5  do {  
6      statements;  
7  } while(condition);
```

- Braces are optional if the body is a single statement



Control structures: while loop

C++ 98

while loop syntax

```
1  while(condition) {  
2      statements;  
3  }  
  
4  
  
5  do {  
6      statements;  
7  } while(condition);
```

- Braces are optional if the body is a single statement

Bad example

```
1  while (n != 1)  
2      if (0 == n%2) n /= 2;  
3      else n = 3 * n + 1;
```



Control structures: jump statements

C++ 98

`break` Exits the loop and continues after it

`continue` Goes immediately to next loop iteration

`return` Exits the current function

`goto` Can jump anywhere inside a function, avoid!



Control structures: jump statements

C++ 98

break Exits the loop and continues after it

continue Goes immediately to next loop iteration

return Exits the current function

goto Can jump anywhere inside a function, avoid!

Bad example

```
1 while (1) {  
2     if (n == 1) break;  
3     if (0 == n%2) {  
4         std::cout << n << '\n';  
5         n /= 2;  
6         continue;  
7     }  
8     n = 3 * n + 1;  
9 }
```



Control structures

C++ 11

Exercise: Control structures

Familiarise yourself with different kinds of control structures.
Re-implement them in different ways.

- Go to exercises/control
- Look at control.cpp
- Compile it (make) and run the program (./control)
- Work on the tasks that you find in README.md



Headers and interfaces

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- **Headers and interfaces**
- Auto keyword
- Inline keyword
- Assertions



Headers and interfaces

C++ 98

Interface

Set of declarations defining some functionality

- Put in a so-called “header file”
- The implementation exists somewhere else

Header: hello.hpp

```
void printHello();
```

Usage: myfile.cpp

```
1 #include "hello.hpp"
2 int main() {
3     printHello();
4 }
```

Preprocessor

C++ 98

```
1 // file inclusion
2 #include "hello.hpp"
3 // macro constants and function-style macros
4 #define MY_GOLDEN_NUMBER 1746
5 #define CHECK_GOLDEN(x) if ((x) != MY_GOLDEN_NUMBER) \
6     std::cerr << #x " was not the golden number\n";
7 // compile time or platform specific configuration
8 #if defined(USE64BITS) || defined(__GNUG__)
9     using myint = std::uint64_t;
10 #elif
11     using myint = std::uint32_t;
12 #endif
```



Preprocessor

C++ 98

```
1 // file inclusion
2 #include "hello.hpp"
3 // macro constants and function-style macros
4 #define MY_GOLDEN_NUMBER 1746
5 #define CHECK_GOLDEN(x) if ((x) != MY_GOLDEN_NUMBER) \
6     std::cerr << #x " was not the golden number\n";
7 // compile time or platform specific configuration
8 #if defined(USE64BITS) || defined(__GNUG__)
9     using myint = std::uint64_t;
10 #elif
11     using myint = std::uint32_t;
12 #endif
```

Good practice: Use preprocessor only in very restricted cases

- Conditional inclusion of headers
- Customization for specific compilers/platforms



Header include guards

C++ 98

Problem: redefinition by accident

- Headers may define new names (e.g. types)
- Multiple (transitive) inclusions of a header would define those names multiple times, which is a compile error
- Solution: guard the content of your headers!

Include guards

```
1 #ifndef MY_HEADER_INCLUDED
2 #define MY_HEADER_INCLUDED
3 ... // header file content
4 #endif
```

Pragma once (non-standard)

```
1 #pragma once
2 ... // header file content
```

Header / source separation

C++ 98

Good practice: Headers and source files

- Headers should contain declarations of functions / classes
 - Only create them if interface is used somewhere else
- Might be included/compiled many times
- Good to keep them short
- Minimise `#include` statements
- Put long code in implementation files. Exceptions:
 - Short functions
 - Templates and constexpr functions



Auto keyword

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Auto keyword

C++ 11

Reason of being

- Many type declarations are redundant
- They are often a source for compiler warnings and errors
- Using auto prevents unwanted/unnecessary type conversions

```
1 std::vector<int> v;
2 float a = v[3];      // conversion intended?
3 int b = v.size();    // bug? unsigned to signed
```



Auto keyword

C++ 11

Reason of being

- Many type declarations are redundant
- They are often a source for compiler warnings and errors
- Using auto prevents unwanted/unnecessary type conversions

```
1 std::vector<int> v;
2 float a = v[3];      // conversion intended?
3 int b = v.size();    // bug? unsigned to signed
```

Practical usage

```
1 std::vector<int> v;
2 auto a = v[3];
3 const auto b = v.size(); // std::size_t
4 int sum{0};
5 for (auto n : v) { sum += n; }
```



Loops, references, auto

C++ 98

Exercise: Loops, references, auto

Familiarise yourself with range-based for loops and references

- Go to exercises/loopsRefsAuto
- Look at loopsRefsAuto.cpp
- Compile it (make) and run the program (./loopsRefsAuto)
- Work on the tasks that you find in loopsRefsAuto.cpp



Inline keyword

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- **Inline keyword**
- Assertions



Inline keyword

C++ 98 Adv

Inline functions originally

- Applies to a function to tell the compiler to inline it
 - That is, replace function calls by the function's content (similar to how a macro works)
- Only a hint, compiler can still choose to not inline
- Avoids call overhead at the cost of increasing binary size

Major side effect

- The linker reduces the duplicated functions into one
- An inline function definition can thus live in header files

```
1 inline int mult(int a, int b) {  
2     return a * b;  
3 }
```



Inline keyword

C++ 98 Adv

Inline functions nowadays

- Compilers can judge far better when to inline or not
 - thus primary purpose is gone
- Putting functions into headers became main purpose
- Many types of functions are marked **inline** by default:
 - function templates
 - **constexpr** functions
 - class member functions



Inline keyword

C++ 17 Adv

Inline variables

- Global (or **static** member) variable specified as **inline**
- Same side effect, linker merges all occurrences into one
- Allows to define global variables/constants in headers

```
1 // global.h
2 inline int count = 0;
3 inline const std::string filename = "output.txt";
4 // a.cpp
5 #include "global.h"
6 int f() { return count; }
7 // b.cpp
8 #include "global.h"
9 void g(int i) { count += i; }
```

- Avoid global variables! Global constants are fine.



Assertions

2

Language basics

- Core syntax and types
- Arrays and Pointers
- Scopes / namespaces
- Class and enum types
- References
- Functions
- Operators
- Control structures
- Headers and interfaces
- Auto keyword
- Inline keyword
- Assertions



Assertions

C++ 98 Adv

Checking invariants in a program

- An invariant is a property that is guaranteed to be true during certain phases of a program, and the program might crash or yield wrong results if it is violated
 - “Here, ‘a’ should always be positive”
- This can be checked using `assert`
- The program will be aborted if the assertion fails

assert in practice - godbolt

```
1 #include <cassert>
2 double f(double a) {
3     // [...] do stuff with a
4     // [...] that should leave it positive
5     assert(a > 0.);
6     return std::sqrt(a);
7 }
```



Assertions

C++ 98 Adv

Checking invariants in a program

- An invariant is a property that is guaranteed to be true during certain phases of a program, and the program might crash or yield wrong results if it is violated
 - “Here, ‘a’ should always be positive”
- This can be checked using `assert`
- The program will be aborted if the assertion fails

assert in practice - godbolt

```
% ./testAssert
Assertion failed: (a > 0.), function f,
file testAssert.cpp, line 5.
Abort trap: 6
```



Assertions

C++ 98 Adv

Good practice: Assert

- Assertions are mostly for developers and debugging
- Use them to check important invariants of your program
- Prefer handling user-facing errors with helpful error messages/exceptions
- Assertions can impact the speed of a program
 - Assertions are disabled when the macro `NDEBUG` is defined
 - Decide if you want to disable them when you release code

Disabling assertions

Compile a program with `NDEBUG` defined:

```
g++ -DNDEBUG -O2 -W [...] test.cpp -o test.exe
```

```
1 double f(double a) {  
2     assert(a > 0.); // no effect  
3     return std::sqrt(a);  
4 }
```



Static Assert

C++ 11 Adv

Checking invariants at compile time

- To check invariants at compile time, use `static_assert`
- The assertion can be any constant expression (see later)
- The message argument is optional in C++ 17 and later

static_assert

```
1 double f(UserType a) {
2     static_assert(
3         std::is_floating_point_v<UserType>, // C++17
4         "This function expects floating-point types.");
5     return std::sqrt(a);
6 }
```

```
a.cpp: In function 'double f(UserType)':  
a.cpp:3:9: error: static assertion failed: This function  
expecting floating-point types.
```

```
2 | static_assert(
|     std::is_floating_point_v<UserType>,
|     ~~~~~^~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~
```



Object orientation (OO)

1 History and goals

- Operator overloading
- Function objects
- Name Lookups

2 Language basics

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting

4 Core modern C++

5 Expert C++

6 Useful tools

7 Concurrency

8 C++ and python

Objects and Classes

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



What are classes and objects

C++ 98

Classes (or “user-defined types”)

C structs on steroids

- with inheritance
- with access control
- including methods (aka. member functions)

Objects

- instances of classes

A class encapsulates state and behavior of “something”

- shows an interface
- provides its implementation
 - status, properties
 - possible interactions
 - construction and destruction

My first class

C++ 98

```
1 struct MyFirstClass {
2     int a;
3     void squareA() {
4         a *= a;
5     }
6     int sum(int b) {
7         return a + b;
8     }
9 };
10
11 MyFirstClass myObj;
12 myObj.a = 2;
13
14 // let's square a
15 myObj.squareA();
```

MyFirstClass

```
int a;
void squareA();
int sum(int b);
```

Separating the interface

C++ 98

Header: MyClass.hpp

```

1 #pragma once
2 struct MyClass {
3     int a;
4     void squareA();
5 };

```

User 1: main.cpp

```

1 #include "MyClass.hpp"
2 int main() {
3     MyClass mc;
4     ...
5 }

```

Implementation: MyClass.cpp

```

1 #include "MyClass.hpp"
2 void MyClass::squareA() {
3     a *= a;
4 }

```

User 2: fun.cpp

```

1 #include "MyClass.hpp"
2 void f(MyClass& mc) {
3     mc.squareA();
4 }

```



Implementing methods

C++ 98

Good practice: Implementing methods

- usually in .cpp, outside of class declaration
- using the class name as “namespace”
- short member functions can be in the header
- some functions (templates, **constexpr**) must be in the header

```
1 #include "MyFirstClass.hpp"  
2  
3 void MyFirstClass::squareA() {  
4     a *= a;  
5 }  
6 int MyFirstClass::sum(int b) {  
7     return a + b;  
8 }
```

Method overloading

C++ 98

The rules in C++

- overloading is authorized and welcome
- signature is part of the method identity
- but not the return type

```
1 struct MyFirstClass {  
2     int a;  
3     int sum(int b);  
4     int sum(int b, int c);  
5 };  
6  
7 int MyFirstClass::sum(int b) { return a + b; }  
8  
9 int MyFirstClass::sum(int b, int c) {  
10    return a + b + c;  
11 }
```



Inheritance


3 Object orientation (OO)

- Objects and Classes
- **Inheritance**
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups

First inheritance

C++ 98

```
1 struct MyFirstClass {  
2     int a;  
3     void squareA() { a *= a; }  
4 };  
5 struct MySecondClass :  
6     MyFirstClass {  
7     int b;  
8     int sum() { return a + b; }  
9 };  
10  
11 MySecondClass myObj2;  
12 myObj2.a = 2;  
13 myObj2.b = 5;  
14 myObj2.squareA();  
15 int i = myObj2.sum(); // i = 9
```



First inheritance


C++ 98

```

1  struct MyFirstClass {
2      int a;
3      void squareA() { a *= a; }
4  };
5  struct MySecondClass :
6      MyFirstClass {
7      int b;
8      int sum() { return a + b; }
9  };
10
11 MySecondClass myObj2;
12 myObj2.a = 2;
13 myObj2.b = 5;
14 myObj2.squareA();
15 int i = myObj2.sum(); // i = 9

```

Memory layout



Managing access to class members

C++ 98

public / private keywords

`private` allows access only within the class

`public` allows access from anywhere

- The default for class is `private`
- The default for struct is `public`

```
1  class MyFirstClass {          9  MyFirstClass obj;
2  public:                      10  obj.a = 5;    // error !
3  void setA(int x);           11  obj.setA(5); // ok
4  int getA();                  12  obj.squareA();
5  void squareA();             13  int b = obj.getA();
6  private:                     14
7  int a;                       15
8  };
```

Managing access to class members

C++ 98

public / private keywords

private allows access only within the class

public allows access from anywhere

- The default for class is **private**
- The default for struct is **public**

```
1  class MyFirstClass {           9  MyFirstClass obj;
2  public:                      10  obj.a = 5;    // error !
3  void setA(int x);          11  obj.setA(5); // ok
4  int getA();                  12  obj.squareA();
5  void squareA();            13  int b = obj.getA();
6  private:                    14
7  int a;                      15
8  };
```

This breaks MySecondClass !



Managing access to class members(2)

C++ 98

Solution is protected keyword

Gives access to classes inheriting from base class

```
1  class MyFirstClass {          13  class MySecondClass :  
2  public:  
3      void setA(int a);        14      public MyFirstClass {  
4      int getA();             15      public:  
5      void squareA();         16          int sum() {  
6  protected:  
7      int a;                  17              return a + b;  
8  };                         18      }  
                                19  private:  
                                20      int b;  
                                21  };
```



Managing inheritance privacy

C++ 98

Inheritance can be public, protected or private

It influences the privacy of inherited members for external code.

The code of the class itself is not affected

`public` privacy of inherited members remains unchanged

`protected` inherited public members are seen as protected

`private` all inherited members are seen as private

this is the default for classes if nothing is specified



Managing inheritance privacy

C++ 98

Inheritance can be public, protected or private

It influences the privacy of inherited members for external code.

The code of the class itself is not affected

`public` privacy of inherited members remains unchanged

`protected` inherited public members are seen as protected

`private` all inherited members are seen as private

this is the default for classes if nothing is specified

Net result for external code


- only public members of public inheritance are accessible

Net result for code in derived classes

- only public and protected members of public and protected parents are accessible

Managing inheritance privacy - public

C++ 98




```

1 void funcSecond() {
2     int a = priv;      // Error
3     int b = prot;     // OK
4     int c = pub;      // OK
5 }
6 void funcThird() {
7     int a = priv;      // Error
8     int b = prot;     // OK
9     int c = pub;      // OK
10 }
11 void extFunc(MyThirdClass t) {
12     int a = t.priv;   // Error
13     int b = t.prot;   // Error
14     int c = t.pub;    // OK
15 }
  
```



Managing inheritance privacy - protected

C++ 98




```

1 void funcSecond() {
2     int a = priv;      // Error
3     int b = prot;     // OK
4     int c = pub;      // OK
5 }
6 void funcThird() {
7     int a = priv;      // Error
8     int b = prot;     // OK
9     int c = pub;      // OK
10 }
11 void extFunc(MyThirdClass t) {
12     int a = t.priv;   // Error
13     int b = t.prot;   // Error
14     int c = t.pub;    // Error
15 }
  
```

Managing inheritance privacy - private

C++ 98



```

1 void funcSecond() {
2     int a = priv;      // Error
3     int b = prot;      // OK
4     int c = pub;       // OK
5 }
6 void funcThird() {
7     int a = priv;      // Error
8     int b = prot;      // Error
9     int c = pub;       // Error
10 }
11 void extFunc(MyThirdClass t) {
12     int a = t.priv;   // Error
13     int b = t.prot;   // Error
14     int c = t.pub;    // Error
15 }
  
```

Final class

C++ 11 Adv

Idea

- make sure you cannot inherit from a given class
- by declaring it final

Practically

```
1 struct Base final {
2     ...
3 };
4 struct Derived : Base { // compiler error
5     ...
6 };
```



Constructors/destructors

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



Class constructors and destructors

C++ 98

Concept

- special functions called when building/destroying an object
- a class can have several constructors, but only one destructor
- the constructors have the same name as the class
- same for the destructor with a leading ~

```
1  class C {                      10 // note: special notation for
2  public:                         11 // initialization of members
3    C();                           12 C::C() : a(0) {}
4    C(int a);                     13 C::C(int a) : a(a) {}
5    ~C();                          14 C::~C() {}
6    ...                            15
7  protected:                      16
8    int a;                         17
9  };
```

Class constructors and destructors

C++ 98

```
1 class Vector {
2 public:
3     Vector(int n);
4     ~Vector();
5     void setN(int n, int value);
6     int getN(int n);
7 private:
8     int len;
9     int* data;
10 };
11 Vector::Vector(int n) : len(n) {
12     data = new int[n];
13 }
14 Vector::~Vector() {
15     delete[] data;
16 }
```

Constructors and inheritance

C++ 98

```
1 struct First {
2     int a;
3     First() {} // leaves a uninitialized
4     First(int a) : a(a) {}
5 };
6 struct Second : First {
7     int b;
8     Second();
9     Second(int b);
10    Second(int a, int b);
11 };
12 Second::Second() : First(), b(0) {}
13 Second::Second(int b) : b(b) {} // First() implicitly
14 Second::Second(int a, int b) : First(a), b(b) {}
```



Copy constructor

C++ 11

Concept

- special constructor called for replicating an object
- takes a single parameter of type **const &** to class
- provided by the compiler if not declared by the user
- in order to forbid copy, use = **delete** (see next slides)
 - or private copy constructor with no implementation in C++ 98



Copy constructor

C++ 11

Concept

- special constructor called for replicating an object
- takes a single parameter of type **const &** to class
- provided by the compiler if not declared by the user
- in order to forbid copy, use = **delete** (see next slides)
 - or private copy constructor with no implementation in C++ 98

```
1 struct C {  
2     C();  
3     C(const C &other);  
4 };
```



Copy constructor

C++ 11

Concept

- special constructor called for replicating an object
- takes a single parameter of type **const &** to class
- provided by the compiler if not declared by the user
- in order to forbid copy, use = **delete** (see next slides)
 - or private copy constructor with no implementation in C++ 98

```
1 struct C {  
2     C();  
3     C(const C &other);  
4 };
```

Good practice: The rule of 3/5 (C++ 98/11) - [cppreference](#)

if a class needs a custom destructor, a copy/move constructor or a copy/move assignment operator, it should have all three/five.



Class Constructors and Destructors

C++ 98

```
1 class Vector {
2 public:
3     Vector(int n);
4     Vector(const Vector &other);
5     ~Vector();
6 private:
7     int len; int* data;
8 };
9 Vector::Vector(int n) : len(n) {
10    data = new int[n];
11 }
12
13
14
15
16 Vector::~Vector() { delete[] data; }
```



Class Constructors and Destructors

C++ 98

```
1 class Vector {
2 public:
3     Vector(int n);
4     Vector(const Vector &other);
5     ~Vector();
6 private:
7     int len; int* data;
8 };
9 Vector::Vector(int n) : len(n) {
10    data = new int[n];
11 }
12 Vector::Vector(const Vector &other) : len(other.len) {
13    data = new int[len];
14    std::copy(other.data, other.data + len, data);
15 }
16 Vector::~Vector() { delete[] data; }
```



Explicit unary constructor

C++ 98

Concept

- A constructor with a single non-default parameter can be used by the compiler for an implicit conversion.

Example - godbolt

```
1 void print(const Vector & v) {  
2     std::cout << "printing v elements...\n";  
3 }  
4  
5 int main {  
6     // calls Vector::Vector(int n) to construct a Vector  
7     // then calls print with that Vector  
8     print(3);  
9 }
```



Explicit unary constructor

C++ 98

Concept

- The keyword **explicit** forbids such implicit conversions.
- It is recommended to use it systematically, except in special cases.

```
1 class Vector {  
2 public:  
3     explicit Vector(int n);  
4     Vector(const Vector &other);  
5     ~Vector();  
6     ...  
7 };
```



Defaulted Constructor

C++ 11

Idea

- avoid empty default constructors like `ClassName() {}`
- declare them as = **default**

Details

- without a user-defined constructor, a default one is provided
- any user-defined constructor disables the default one
- but the default one can be requested explicitly
- rule can be more subtle depending on data members

Practically

```
1 Class() = default; // provide default if possible
2 Class() = delete; // disable default constructor
```

Delegating constructor

C++ 11

Idea

- avoid replication of code in several constructors
- by delegating to another constructor, in the initialization list

Practically

```
1 struct Delegate {  
2     int m_i;  
3     explicit Delegate(int i) : m_i(i) {  
4         ... complex initialization ...  
5     }  
6     Delegate() : Delegate(42) {}  
7 };
```



Constructor inheritance

C++ 11

Idea

- avoid having to redeclare parent's constructors
- by stating that we inherit all parent constructors
- derived class can add more constructors

Practically

```
1 struct Base {  
2     explicit Base(int a); // ctor 1  
3 };  
4 struct Derived : Base {  
5     using Base::Base;  
6     Derived(int a, int b); // ctor 2  
7 };  
8 Derived d{5}; // calls ctor 1  
9 Derived d{5, 6}; // calls ctor 2
```

Member initialization

C++ 11

Idea

- avoid redefining same default value for members n times
- by defining it once at member declaration time

Practically

```
1 struct Base {  
2     int a{5}; // or: int a = 5;  
3         Base() = default;  
4     explicit Base(int _a) : a(_a) {}  
5 };  
6 struct Derived : Base {  
7     int b{6};  
8     using Base::Base;  
9 };  
10 Derived d1; // a = 5, b = 6  
11 Derived d2{7}; // a = 7, b = 6
```



Calling constructors

C++ 11

After object declaration, arguments within {}

```
1 struct A {  
2     int i;  
3     float f;  
4     A();  
5     A(int);  
6     A(int, int);  
7 };  
8  
9 A a{1,2};      // A::A(int, int)  
10 A a{1};        // A::A(int)  
11 A a{};         // A::A()  
12 A a;           // A::A()  
13 A a = {1,2};  // A::A(int, int)
```

Calling constructors the old way

C++ 98

Arguments are given within (), aka C++ 98 nightmare

```
1 struct A {  
2     int i;  
3     float f;  
4     A();  
5     A(int);  
6     A(int, int);  
7 };  
8  
9 A a(1,2);      // A::A(int, int)  
10 A a(1);        // A::A(int)  
11 A a();         // declaration of a function!  
12 A a;           // A::A()  
13 A a = (1,2);  // A::A(int), comma operator!  
14 A a = {1,2};  // not allowed
```

Constructing arrays and vectors

C++ 11

List of items given within {}

```
10 int ip[3]{1,2,3};  
11 int* ip = new int[3]{1,2,3}; // not allowed in C++98  
12 std::vector<int> v{1,2,3}; // same
```



Static members

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- **Static members**
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



Static members

C++ 98

Concept

- members attached to a class rather than to an object
- usable with or without an instance of the class
- identified by the **static** keyword

Static.hpp

```
1 class Text {
2 public:
3     static std::string upper(std::string);
4 private:
5     static int callsToUpper; // add `inline` in C++17
6 };
```



Static members

C++ 98

Concept

- members attached to a class rather than to an object
- usable with or without an instance of the class
- identified by the **static** keyword

Static.cpp

```
1 #include "Static.hpp"
2 int Text::callsToUpper = 0; // required before C++17
3
4 std::string Text::upper(std::string lower) {
5     callsToUpper++;
6     // convert lower to upper case
7     // return ...;
8 }
9 std::string uppers = Text::upper("my text");
10 // now Text::callsToUpper is 1
```



Allocating objects

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



Process memory organization

C++ 98

4 main areas


the code segment for the machine code of the executable

the data segment for global variables

the heap for dynamically allocated variables

the stack for parameters of functions and local variables

Memory layout



The Stack

C++ 98

Main characteristics

- allocation on the stack stays valid for the duration of the current scope. It is destroyed when it is popped off the stack.
- memory allocated on the stack is known at compile time and can thus be accessed through a variable.
- the stack is relatively small, it is not a good idea to allocate large arrays, structures or classes
- each thread in a process has its own stack
 - allocations on the stack are thus “thread private”
 - and do not introduce any thread-safety issues



Object allocation on the stack

C++ 98

On the stack

- objects are created on variable definition (constructor called)
- objects are destructed when out of scope (destructor is called)

```
1 int f() {  
2     MyFirstClass a{3}; // constructor called  
3     ...  
4 } // destructor called  
5  
6 int g() {  
7     MyFirstClass a; // default constructor called  
8     ...  
9 } // destructor called
```



The Heap

C++ 98

Main characteristics

- Allocated memory stays allocated until it is specifically deallocated
 - beware memory leaks
- Dynamically allocated memory must be accessed through pointers
- large arrays, structures, or classes should be allocated here
- there is a single, shared heap per process
 - allows to share data between threads
 - introduces race conditions and thread-safety issues!



Object allocation on the heap

C++ 98

On the heap

- objects are created by calling `new` (constructor is called)
- objects are destructed by calling `delete` (destructor is called)

```
1 int f() {
2     // default constructor called
3     MyFirstClass *a = new MyFirstClass;
4     delete a; // destructor is called
5 }
6 int g() {
7     // constructor called
8     MyFirstClass *a = new MyFirstClass{3};
9 } // memory leak !!!
```

Good practice: Prefer smart pointers over new/delete

Prefer smart pointers to manage objects (discussed later)



Array allocation on the heap

C++ 98

Arrays on the heap

- arrays of objects are created by calling `new[]`
default constructor is called for each object of the array
- arrays of objects are destructed by calling `delete[]`
destructor is called for each object of the array

```
1 int f() {  
2     // default constructor called 10 times  
3     MyFirstClass *a = new MyFirstClass[10];  
4     ...  
5     delete[] a; // destructor called 10 times  
6 }
```

Good practice: Prefer containers over new-ed arrays

Prefer containers to manage collections of objects (discussed later)

Advanced OO

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



this keyword

C++ 98

How to know an object's address?


- Sometimes we need the address of the current object
- Or we need to pass our address / a reference to a different entity (for example to implement operators, see later)
- All class methods can use the keyword **this**
 - It returns the address of the current object
 - Its type is T^* in the methods of a class T

```

1 struct S {
2     int a,b;
3     // these two are the same:
4     int getB() { return b; }           // 5
5     int getB() { return this->b; }    // 5
6     void testAddress() {
7         S* addr = this; // 0x3000
8     }
9 } s{2,5};

```

Memory layout




Polymorphism

C++ 98

the concept

- objects actually have multiple types simultaneously
- and can be used as any of them

```
1 Polygon p;  
2  
3 int f(Drawable & d) {...}  
4 f(p); //ok  
5  
6 try {  
7     throw p;  
8 } catch (Shape & e) {  
    // will be caught  
}
```



Polymorphism

C++ 98

the concept

- objects actually have multiple types simultaneously
- and can be used as any of them

```

1  Polygon p;
2
3  int f_Drawable( Drawable & d) { ... }
4  f(p); //ok
5
6  try {
7      throw p;
8  } catch (Shape & e) {
9      // will be caught
10 }
```

Memory layout

	0x3020
	0x301C
Polygon	Polygon.nLines
	...
	Shape.b
	Shape.a
	...
	Drawable.b
	Drawable.a



Polymorphism

C++ 98

the concept

- objects actually have multiple types simultaneously
- and can be used as any of them

```

1  Polygon p;
2
3  int f_Drawable( Drawable & d) { ... }
4  f(p); //ok
5
6  try {
7      throw p;
8  } catch (Shape & e) {
9      // will be caught
10 }
```

Memory layout

	0x3020
	0x301C
Polygon.nLines	0x3018
...	0x3014
Shape.b	0x3010
Shape.a	0x300C
...	0x3008
Drawable.b	0x3004
Drawable.a	0x3000

Drawable



Polymorphism

C++ 98

the concept

- objects actually have multiple types simultaneously
- and can be used as any of them

```

1  Polygon p;
2
3  int f_Drawable( Drawable & d) { ... }
4  f(p); //ok
5
6  try {
7      throw p;
8  } catch (Shape & e) {
9      // will be caught
10 }
```

Memory layout

	0x3020
	0x301C
Polygon.nLines	0x3018
...	0x3014
Shape.b	0x3010
Shape.a	0x300C
...	0x3008
Drawable.b	0x3004
Drawable.a	0x3000

Shape




Inheritance privacy and polymorphism

C++ 98

Only public base classes are visible to outside code

- private and protected bases are not
- this may restrict usage of polymorphism

```
1 Polygon p;  
2  
3 int f(Drawable & d) {...}  
4 f(p); // Not ok anymore  
5  
6 try {  
7     throw p;  
8 } catch (Shape & e) {  
9     // ok, will be caught  
10 }
```




Method overriding

C++ 98

the idea

- a method of the parent class can be replaced in a derived class
- but which one is called?

```
1 Polygon p;  
2 p.draw(); // ?  
3  
4 Shape & s = p;  
5 s.draw(); // ?
```



Virtual methods

C++ 98

the concept

- methods can be declared **virtual**
- for these, the most derived object's implementation is used
(i.e. the dynamic type behind a pointer/reference)
- for non-virtual methods, the static type of the variable decides




Virtual methods

C++ 98

the concept

- methods can be declared **virtual**
- for these, the most derived object's implementation is used
(i.e. the dynamic type behind a pointer/reference)
- for non-virtual methods, the static type of the variable decides

```
1 Polygon p;  
2 p.draw(); // Polygon.draw  
3  
4 Shape & s = p;  
5 s.draw(); // Drawable.draw
```




Virtual methods

C++ 98

the concept

- methods can be declared **virtual**
- for these, the most derived object's implementation is used
(i.e. the dynamic type behind a pointer/reference)
- for non-virtual methods, the static type of the variable decides

```
1 Polygon p;  
2 p.draw(); // Polygon.draw  
3  
4 Shape & s = p;  
5 s.draw(); // Polygon.draw
```



Virtual methods - implications

C++ 11

Mechanics

- virtual methods are dispatched at run time
 - while non-virtual methods are bound at compile time
- they also imply extra storage and an extra indirection
 - practically, the object stores a pointer to the correct method
 - in a so-called “virtual table” (“vtable”)

Consequences

- virtual methods are “slower” than standard ones
- and they can rarely be inlined
- templates are an alternative for performance-critical cases



override keyword

C++ 11

Principle

- when overriding a virtual method, the **override** keyword should be used
- the **virtual** keyword is then optional in derived classes

Practically

```
1 struct Base {  
2     virtual void some_func(float);  
3 };  
4 struct Derived : Base {  
5     void some_func(float) override;  
6 };
```



Why was override keyword introduced?

C++ 11

To detect the mistake in the following code :

Without override (C++ 98)

```
1 struct Base {  
2     virtual void some_func(float);  
3 };  
4 struct Derived : Base {  
5     void some_func(double); // oops !  
6 };
```

- with **override**, you would get a compiler error
- if you forget **override** when you should have it, you may get a compiler warning



final keyword

C++ 11 Adv

Idea

- make sure you cannot further override a given virtual method
- by declaring it final

Practically

```
1 struct Base {  
2     virtual void some_func(float);  
3 };  
4 struct Intermediate : Base {  
5     void some_func(float) final;  
6 };  
7 struct Derived : Intermediate {  
8     void some_func(float) override; // error  
9 };
```

Pure Virtual methods

C++ 11

Concept

- unimplemented methods that must be overridden
- marked by = 0 in the declaration
- makes their class abstract
- only non-abstract classes can be instantiated




Pure Virtual methods

C++ 11

Concept

- unimplemented methods that must be overridden
- marked by = 0 in the declaration
- makes their class abstract
- only non-abstract classes can be instantiated

```
1 // Error : abstract class
2 Shape s;
3
4 // ok, draw has been implemented
5 Polygon p;
6
7 // Shape type still usable
8 Shape & s = p;
9 s.draw();
```



Polymorphism and destruction

C++ 98

Owning base pointers

We sometimes need to maintain owning pointers to base classes:

```
1 struct Drawable {  
2     virtual void draw() = 0;  
3 };  
4 Drawable* getImpl();  
5  
6 Drawable* p = getImpl();  
7 p->draw();  
8 delete p;
```

- What happens when `p` is deleted?
- What if a class deriving from `Drawable` has a destructor?



Polymorphism and destruction

C++ 98

Owning base pointers

We sometimes need to maintain owning pointers to base classes:

```
1 struct Drawable {  
2     virtual void draw() = 0;  
3 };  
4 std::unique_ptr<Drawable> getImpl(); // better API  
5  
6 auto p = getImpl();  
7 p->draw();
```

- What happens when `p` is deleted?
- What if a class deriving from `Drawable` has a destructor?



Polymorphism and destruction

C++ 11

Virtual destructors

- We can mark a destructor as **virtual**
- This selects the right destructor based on the runtime type

```
1 struct Drawable {  
2     virtual ~Drawable() = default;  
3     virtual void draw() = 0;  
4 };  
5 Drawable* p = getImpl(); // returns derived obj.  
6 p->draw();  
7 delete p; // dynamic dispatch to right destructor
```

Good practice: Virtual destructors

If you expect users to inherit from your class and override methods (i.e. use your class polymorphically), declare its destructor **virtual**



Pure Abstract Class aka Interface

C++ 11

Definition of pure abstract class

- a class that has
 - no data members
 - all its methods pure virtual
 - a **virtual** destructor
- the equivalent of an Interface in Java

```
1 struct Drawable {  
2     virtual ~Drawable()  
3             = default;  
4     virtual void draw() = 0;  
5 }
```

Drawable
virtual void draw() = 0;



Overriding overloaded methods

C++ 98

Concept

- overriding an overloaded method will hide the others
- unless you inherit them using **using**

```
1 struct BaseClass {  
2     virtual int foo(std::string);  
3     virtual int foo(int);  
4 };  
5 struct DerivedClass : BaseClass {  
6     using BaseClass::foo;  
7     int foo(std::string) override;  
8 };  
9 DerivedClass dc;  
10 dc.foo(4);      // error if no using
```



Polymorphism

C++ 98

Exercise: Polymorphism

- go to exercises/polymorphism
- look at the code
- open trypoly.cpp
- create a Pentagon, call its perimeter method
- create a Hexagon, call its perimeter method
- create a Hexagon, call its parent's perimeter method
- retry with virtual methods




Multiple Inheritance

C++ 98

Concept

- one class can inherit from multiple parents



```
1 class TextBox :  
2     public Rectangle, Text {  
3         // inherits from both  
4         // publicly from Rectangle  
5         // privately from Text  
6     }
```



The diamond shape


C++ 98

Definition

- situation when one class inherits several times from a given grand parent

Problem

- are the members of the grand parent replicated?




Virtual inheritance

C++ 98

Solution

- inheritance can be **virtual** or not
 - **virtual** inheritance will “share” parents
 - standard inheritance will replicate them
- most derived class will call the virtual base class’s constructor

```
1 class Text : public virtual Drawable {...};  
2 class Rectangle : public virtual Drawable {...};
```



Multiple inheritance advice

C++ 98

Good practice: Avoid multiple inheritance

- Except for inheriting from interfaces (=no data members)
- And for rare special cases

Absolutely avoid diamond-shape inheritance

- This is a sign that your architecture is not correct
- In case you are tempted, think twice and change your mind



Virtual inheritance

C++ 98

Exercise: Virtual inheritance

- go to `exercisescode/virtual_inheritance`
- look at the code
- open `trymultiherit.cpp`
- create a `TextBox` and call `draw`
- Fix the code to call both `draws` by using types
- retry with virtual inheritance



Type casting

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups

Type casting

C++ 98 Adv

5 types of casts in C++ - These 2 should be used

- **static_cast<Target>(arg)**
 - converts type if the static types allow it
 - including using implicit conversion
 - essentially compile time
- **dynamic_cast<Target>(arg)**
 - checks if object at address of “arg” is convertible to Target
 - throws std::bad_cast or returns **nullptr** if it's not
 - essentially run time



Type casting example

C++ 98 Adv

Practically - godbolt

```
1 struct A{ virtual ~A()=default; } a;
2 struct B : A {} b;
3
4 A& c = static_cast<A&>(b); // OK. b is also an A
5 B& d = static_cast<B&>(a); // UB: a is not a B
6 B& e = static_cast<B&>(c); // OK. c is a B
7
8 B& f = dynamic_cast<B&>(c); // OK, c is a B
9 B& g = dynamic_cast<B&>(a); // throws std::bad_cast
10 B& foo(A& h) { return dynamic_cast<B&>(h); }
11 B& i = foo(c); // OK, c is a B
12
13 B* j = dynamic_cast<B*>(&a); // nullptr. a not a B.
14 if (j != nullptr) {
15     // Will not reach this
16 }
```



Type casting

C++ 98 Adv

5 types of casts in C++- These 3 should be avoided

- **const_cast<Target>(arg)**
 - removes (or adds) constness from a type
 - if you think you need this, rather improve your design!
- **reinterpret_cast<Target>(arg)**
 - changes type irrespective of what 'arg' is
 - almost never a good idea!
- C-style: (Target)arg
 - Force-changes type in C-style. No checks. Don't use.

Casts to avoid

```
1 void func(A const & a) {  
2     A& ra = a;                      // Error: not const  
3     A& ra = const_cast<A&>(a); // Compiles. Bad design!  
4     // Evil! Don't do this:  
5     B* b = reinterpret_cast<B*>(&a);  
6     B* b = (B*)&a;  
7 }
```



Operator overloading

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- **Operator overloading**
- Function objects
- Name Lookups



Operator overloading example

C++ 98

```
1 struct Complex {  
2     float m_real, m_imaginary;  
3     Complex(float real, float imaginary);  
4     Complex operator+(const Complex& other) {  
5         return Complex{m_real + other.m_real,  
6                         m_imaginary + other.m_imaginary};  
7     }  
8 };  
9  
10    Complex c1{2, 3}, c2{4, 5};  
11    Complex c3 = c1 + c2; // (6, 8)
```



Operator overloading

C++ 98

Defining operators for a class

- implemented as a regular method
 - either inside the class, as a member function
 - or outside the class (not all)
- with a special name (replace @ by operators from below)

Expression	As member	As non-member
@a	(a).operator@()	operator@(a)
a@b	(a).operator@(b)	operator@(a,b)
a=b	(a).operator=(b)	cannot be non-member
a(b...)	(a).operator()(b...)	cannot be non-member
a[b]	(a).operator[](b)	cannot be non-member
a->	(a).operator->()	cannot be non-member
a@	(a).operator@(0)	operator@(a,0)

possible operators: + - * / % ^ & | ~ ! = < >

+ = - = * = / = % = ^ = & = | = << >> >>= <<=

== != <= >= <=> && || ++ -- , ->* -> () []



Why have non-member operators?

C++ 98

Symmetry

```
1 struct Complex {  
2     float m_real, m_imaginary;  
3     Complex operator+(float other) {  
4         return Complex{m_real + other, m_imaginary};  
5     }  
6 };  
7 Complex c1{2.f, 3.f};  
8 Complex c2 = c1 + 4.f;    // ok  
9 Complex c3 = 4.f + c1;    // not ok !!
```



Why have non-member operators?

C++ 98

Symmetry

```
1 struct Complex {  
2     float m_real, m_imaginary;  
3     Complex operator+(float other) {  
4         return Complex{m_real + other, m_imaginary};  
5     }  
6 };  
7 Complex c1{2.f, 3.f};  
8 Complex c2 = c1 + 4.f;    // ok  
9 Complex c3 = 4.f + c1;    // not ok !!  
10    Complex operator+(float a, const Complex& obj) {  
11        return Complex{a + obj.m_real, obj.m_imaginary};  
12    }
```

Other reason to have non-member operators?

C++ 98

Extending existing classes

```
1 struct Complex {  
2     float m_real, m_imaginary;  
3     Complex(float real, float imaginary);  
4 };  
5  
6 std::ostream& operator<<(std::ostream& os,  
7                               const Complex& obj) {  
8     os << "(" << obj.m_real << ", "  
9          << obj.m_imaginary << ")";  
10    return os;  
11 }  
12 Complex c1{2.f, 3.f};  
13 std::cout << c1 << std::endl; // Prints '(2, 3)'
```



Chaining operators

C++ 98

In general, return a reference to the left value

```
1 struct Complex {
2     float m_real, m_imaginary;
3     Complex& operator=( const Complex& other ) {
4         m_real = other.m_real;
5         m_imaginary = other.m_imaginary;
6         return *this;
7     }
8 };
9 Complex c1{2.f, 3.f};
10 Complex c2, c3;
11 // right to left associativity
12 c3 = c2 = c1;
13 // left to right associativity
14 std::cout << c1 << c2 << c3 << std::endl;
```



Friend declarations

C++ 98

Concept

- Functions/classes can be declared **friend** within a class scope
- They gain access to all private/protected members
- Useful for operators such as `a + b`
- Don't abuse friends to go around a wrongly designed interface
- Avoid unexpected modifications of class state in a friend!

operator+ as a friend

```
1 class Complex {
2     float m_r, m_i;
3     friend Complex operator+(Complex const & a, Complex const & b);
4 public:
5     Complex ( float r, float i ) : m_r(r), m_i(i) {}
6 };
7 Complex operator+(Complex const & a, Complex const & b) {
8     return Complex{ a.m_r+b.m_r, a.m_i+b.m_i };
9 }
```



Exercise: Operators

Write a simple class representing a fraction and pass all tests

- go to exercises/operators
- look at operators.cpp
- inspect main and complete the implementation of
`class Fraction` step by step
- you can comment out parts of main to test in between



Function objects

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- **Function objects**
- Name Lookups



Function objects

C++ 98

Concept

- also known as functors (no relation to functors in math)
- a class that implements `operator()`
- allows to use objects in place of functions
- with constructors and data members

```
1 struct Adder {  
2     int m_increment;  
3     Adder(int increment) : m_increment(increment) {}  
4     int operator()(int a) { return a + m_increment; }  
5 };  
6 Adder inc1{1}, inc10{10};  
7 int i = 3;  
8 int j = inc1(i); // 4  
9 int k = inc10(i); // 13  
10 int l = Adder{25}(i); // 28
```



Function objects

C++ 20

Function objects as function arguments - godbolt

```
1 int count_if(const auto& range, auto predicate) {
2     int count = 0; // ↑ template (later)
3     for (const auto& e : range)
4         if (predicate(e)) count++;
5     return count;
6 }
7 struct IsBetween {
8     int lower, upper;
9     bool operator()(int value) const {
10         return lower < value && value < upper;
11     }
12 };
13 int arr[] {1, 2, 3, 4, 5, 6, 7};
14 std::cout << count_if(arr, IsBetween{2, 6}); // 3
15 // prefer: std::ranges::count_if
```



Name Lookups

3 Object orientation (OO)

- Objects and Classes
- Inheritance
- Constructors/destructors
- Static members
- Allocating objects
- Advanced OO
- Type casting
- Operator overloading
- Function objects
- Name Lookups



Basics of name lookup

C++ 98 Adv

Example code

```
1 std::cout << std::endl;
```

How to find the declaration of a name?

Mainly 2 cases :

- qualified lookup, for names preceded by '::'
 - here cout and endl
 - name is only looked for in given class/namespace/enum class
- unqualified lookup
 - here for std and **operator<<**
 - name is looked for in a sequence of scopes until found
 - remaining scopes are not examined
 - in parallel Argument Dependent Lookup (ADL) may happen



Unqualified name lookup and ADL

C++ 98 Adv

Ordered list of scopes (simplified)

- file (only for global level usage)
- current namespace/block, enclosing namespaces/blocks, etc...
- current class if any, base classes if any, etc...

Argument Dependent Lookup (simplified)

Only for call expression

- e.g. `f(a, b)` or `a + b`

The compiler also examines arguments one by one and searches:

- class, if any
- direct and indirect base classes, if any
- enclosing namespace

ADL consequences (1)

C++ 98 Adv

Use standalone/non-member functions

When a method is not accessing the private part of a class, make it a function in the same namespace

Don't write :

```
1  namespace MyNS {  
2      struct A {  
3          T func(...);  
4      };  
5  }
```

Prefer :

```
6  namespace MyNS {  
7      struct A { ... };  
8      T func(const A&, ...);  
9  }
```

Advantages :

- minimal change in user code, func still feels part of class A
- makes sure func does not touch internal state of A

Notes :

- non-member func has to be in same namespace as A
- please avoid global namespace



ADL consequences (2)

C++ 98 Adv

Customization points and using

Don't write :

```
1 N::A a,b;  
2 std::swap(a, b);
```

Prefer :

```
3 N::A a,b;  
4 using std::swap;  
5 swap(a, b);
```

Advantages :

- allows to use std::swap by default
- but benefits from any dedicated overload

```
6 namespace N {  
7     class A { ... };  
8     // optimized swap for A  
9     void swap(A&, A&);  
10 }
```

Hidden Friends

C++ 11 Adv

Hidden friends

- Friend functions *defined* in class scope are “hidden friends”
 - Without any declaration outside the class
- They are *not* member functions (no access to `this`)

operator<< as hidden friend

```
1 class Complex {
2     float m_real, m_imag;
3 public:
4     Complex(float, float) { ... }
5     friend
6     std::ostream & operator<<(std::ostream & os,
7                                     Complex const & c) {
8         return os << c.m_real << " + " << c.m_imag << "i";
9     }
10 };
11 std::cout << Complex{2.f, 3.f}; // Prints 2 + 3i
```

Hidden Friends and ADL

C++ 11 Adv

Advantages of hidden friends

- Are *only* found via ADL, e.g. `std::cout << complex;`
- Compiler needs to consider less functions during ordinary name lookup (faster compilation)
- Avoid accidental implicit conversions

Accidental conversions - two counterexamples

```
1 std::ostream & operator<< // out-of-class definition
2   (std::ostream & os, Complex const & c) { ... }
3 struct Fraction {
4   int num, denom;
5   operator Complex() const { ... } // conversion op
6 };
7
8 std::cout << Fraction{2,4}; // Prints 0.5 + 0i
9 //calls operator<<(std::cout, Fraction{2,4}.operator Complex());
```



Hidden Friends and ADL

C++ 11 Adv

Advantages of hidden friends

- Are *only* found via ADL, e.g. `std::cout << complex;`
- Compiler needs to consider less functions during ordinary name lookup (faster compilation)
- Avoid accidental implicit conversions

Accidental conversions - two counterexamples

```
1 std::ostream & operator<< // out-of-class definition
2   (std::ostream & os, Complex const & c) { ... }
3 struct Fraction {
4   int num, denom;
5 };
6
7 Complex::Complex(Fraction f); // converting ctor
8
9 std::cout << Fraction{2,4}; // Prints 0.5 + 0i
10 //calls operator<<(std::cout, Complex(Fraction{2, 3}));
```



Good practice for operator overloading

C++ 98 Adv

Good practice: Operator overloading

- Must declare in-class, as member operators:
 - `operator=`, `operator()`, `operator[]`, `operator->`
- Prefer in-class declaration, as member operators:
 - compound assignment operators, e.g. `operator+=`
 - unary operators
 - e.g. `operator++`, `operator--`, `operator!`, unary `operator-` (negate), dereference operator `operator*`
- Prefer definition in-class as hidden friends:
 - binary arithmetic operators, e.g. `operator+`, `operator|`
 - binary logical operators, e.g. `operator<`
 - stream operators, e.g. `operator<<`
- Avoid overloading:
 - Address-of `operator&`
 - Boolean logic operators, e.g. `operator||`
 - Comma operator `operator,`
 - Member dereference operators `operator->*`



Core modern C++

1 History and goals

- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization

2 Language basics

3 Object orientation (OO)

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates

5 Expert C++

6 Useful tools

7 Concurrency

8 C++ and python



Constness

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI^I and smart pointers
- Initialization



The const keyword

- indicates that the element to the left is constant
 - when nothing on the left, applies to the right
- this element won't be modifiable in the future
- this is all checked at compile time

```
1 int const i = 6;
2 const int i = 6; // equivalent
3
4 // error: i is constant
5 i = 5;
6
7 auto const j = i; // works with auto
```



Constness and pointers

C++ 98

```
1 int a = 1, b = 2;
2
3 int const *i = &a; // pointer to const int
4 *i = 5; // error, int is const
5 i = &b; // ok, pointer is not const
6
7 int * const j = &a; // const pointer to int
8 *j = 5; // ok, value can be changed
9 j = &b; // error, pointer is const
10
11 int const * const k = &a; // const pointer to const int
12 *k = 5; // error, value is const
13 k = &b; // error, pointer is const
14
15 int const & l = a; // reference to const int
16 l = b; // error, reference is const
17
18 int const & const l = a; // compile error
```

Member function constness

C++ 98

The `const` keyword for member functions

- indicates that the function does not modify the object
- in other words, `this` is a pointer to a constant object

```
1 struct Example {  
2     void foo() const {  
3         // type of 'this' is 'Example const*'  
4         data = 0; // Error: member function is const  
5     }  
6     void foo() { // ok, overload  
7         data = 1; // ok, 'this' is 'Example*'  
8     }  
9     int data;  
10};  
11 Example const e1; e1.foo(); // calls const foo  
12 Example e2; e2.foo(); // calls non-const foo
```

Member function constness

C++ 98

Constness is part of the type

- T **const** and T are different types
- but: T is automatically converted to T **const** when needed

```
1 void change(int & a);
2 void read(int const & a);
3
4 int a = 0;
5 int const b = 0;
6
7 change(a); // ok
8 change(b); // error
9 read(a);   // ok
10 read(b);  // ok
```

Exercise: Constness

- go to exercises/constness
- open constplay.cpp
- try to find out which lines won't compile
- check your guesses by compiling for real



Constant Expressions

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI^I and smart pointers
- Initialization



Generalized Constant Expressions

C++ 14 Adv

Reason of being

- use functions to compute constant expressions at compile time



Generalized Constant Expressions

C++ 14 Adv

Reason of being

- use functions to compute constant expressions at compile time

Example

```
1 constexpr int f(int x) {
2     if (x > 1) return x * f(x - 1);
3     return 1;
4 }
5 constexpr int a = f(5); // computed at compile-time
6 int b = f(5); // maybe computed at compile-time
7 int n = ...; // runtime value
8 int c = f(n); // computed at runtime
```



Generalized Constant Expressions(2)

C++ 14 Adv

Notes

- A `constexpr` function *may* be executed at compile time.
 - Arguments must be `constexpr` or literals in order to allow compile time evaluation
 - Function body must be visible to the compiler
- A `constexpr` function can also be used at runtime
- A `constexpr` variable must be initialized at compile time
- Classes can have `constexpr` member functions
- Objects used in constant expressions must be of *literal type*:
 - integral, floating-point, enum, reference, pointer type
 - union (of at least one) or array of literal types
 - class type with a `constexpr` constructor and the destructor is trivial (or `constexpr` since C++20)
- A `constexpr` function is implicitly `inline` (header files)



Limitations of constexpr functions

C++ 11 Adv

C++11

- Non-virtual function with a single return statement

C++14/C++17

- no try-catch, goto or asm statements
- no uninitialized/static/thread_local/non-literal-type variables

C++20

- no coroutines or static/thread_local/non-literal-type variables
- throw and asm statements allowed, but may not be executed
- transient memory allocation (memory allocated at compile-time must be freed again at compile-time)
- virtual functions and uninitialized variables allowed

C++23

- no coroutines, or execution of throw and asm statements
- transient memory allocation
- everything else allowed

Further details on [cppreference](#)



Real life example

C++ 11 Adv

```
1 constexpr float toSI(float v, char unit) {
2     switch (unit) {
3         case 'k': return 1000.0f*v;
4         case 'm': return 0.001f*v;
5         case 'y': return 0.9144f*v;
6         case 'i': return 0.0254f*v;
7         ...
8         default: return v;
9     }
10 }
11 constexpr float fromSI(float v, char unit) {
12     switch (unit) {
13         case 'k': return 0.001f*v;
14         case 'y': return 1.093f*v;
15         ...
16     }
17 }
```



Real life example(2)

C++ 11 Adv

```
1 class DimLength {
2     float m_value;
3 public:
4     constexpr DimLength(float v, char unit):
5         m_value(toSI(v, unit)) {
6     }
7     constexpr float get(char unit) const {
8         return fromSI(m_value, unit);
9     }
10 };
11 constexpr DimLength km(1, 'k');
12 constexpr float km_y = km.get('y');
13 constexpr float km_i = km.get('i');
14 static_assert(km_y == 1093, "expected km == 1093 yards!");
```



compile-time branches

C++ 17 Adv

if constexpr

- takes a `constexpr` expression as condition
- evaluates at compile time
- key benefit: the discarded branch can contain invalid code

Example code

```
1 template <typename T>
2 auto remove_ptr(T t) {
3     if constexpr (std::is_pointer_v<T>) {
4         return *t;
5     } else {
6         return t;
7     }
8 }
9 int i = ...; int *j = ...;
10 int r = remove_ptr(i); // equivalent to i
11 int q = remove_ptr(j); // equivalent to *j
```



Immediate functions

C++ 20 Adv

Motivation

- Force a function to be executed at compile-time
- Runtime evaluation is forbidden
- Same restrictions as `constexpr` functions
- New keyword: `consteval`

Example

```
1 consteval int f(int x) {
2     if (x > 1) return x * f(x - 1);
3     return 1;
4 }
5 constexpr int a = f(5); // computed at compile-time
6 int b = f(5); // computed at compile-time
7 int n = ...; // runtime value
8 int c = f(n); // compilation error
```

Motivation

- Like a `constexpr` variable, a `constinit` variable guarantees compile-time initialization, but can be modified afterwards
- Only allowed for `static`/global and `thread_local` variables
- Initializer must be a constant expression, but `constexpr` destruction is not required

Example

```
1 constexpr int f(int x) {  
2     if (x > 1) return x * f(x - 1);  
3     return 1;  
4 }  
5 constexpr int a = f(5); // CT init, not modifiable  
6 int b = f(5); // maybe CT init, modifiable  
7 constinit int c = f(5); // CT init, modifiable
```

Exceptions

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions**
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI^I and smart pointers
- Initialization



Exceptions

C++ 98

Purpose

- to handle *exceptional* events that happen rarely
- and cleanly jump to a place where the error can be handled

In practice

- add an exception handling block with `try ... catch`
 - when exceptions are possible *and can be handled*
- throw an exception using `throw`
 - when a function cannot proceed or recover internally

```

1  try {
2      process_data(f);
3  } catch (const
4      std::out_of_range& e) {
5      std::cerr << e.what();
6  }
7  void process_data(file &f) {
8      ...
9      if (i >= buffer.size())
10         throw std::out_of_range{
11             "buf overflow"};
12 }
```

Exceptions

C++ 98

Throwing exceptions

- objects of any type can be thrown (even e.g. `int`)

Good practice: Throwing exceptions

- prefer throwing standard exception classes
- throw objects by value

```
1 #include <stdexcept>
2 void process_data(file& f) {
3     if (!f.open())
4         throw std::invalid_argument{"stream is not open"};
5     auto header = read_line(f); // may throw an IO error
6     if (!header.starts_with("BEGIN"))
7         throw std::runtime_error{"invalid file content"};
8     std::string body(f.size()); // may throw std::bad_alloc
9     ...
10 }
```

Exceptions

C++ 98

Standard exceptions

- `std::exception`, defined in header `<exception>`
 - Base class of all standard exceptions
 - Get error message: `virtual const char* what() const;`
 - Please derive your own exception classes from this one
- From `<stdexcept>`:
 - `std::runtime_error`, `std::logic_error`,
`std::out_of_range`, `std::invalid_argument`, ...
 - Store a string: `throw std::runtime_error{"msg"}`
 - You should use these the most
- `std::bad_alloc`, defined in header `<new>`
 - Thrown by standard allocation functions (e.g. `new`)
 - Signals failure to allocate
 - Carries no message
- ...

Exceptions

C++ 98

Catching exceptions

- a catch clause catches exceptions of the same or derived type
- multiple catch clauses will be matched in order
- if no catch clause matches, the exception propagates
- if the exception is never caught, std::terminate is called

```
1 try {
2     process_data(f);
3 } catch (const std::invalid_argument& e) {
4     bad_files.push_back(f);
5 } catch (const std::exception& e) {
6     std::cerr << "Failed to process file: " << e.what();
7 }
```

Good practice: Catching exceptions

- Catch exceptions by const reference

Exceptions

C++ 98

Rethrowing exceptions

- a caught exception can be rethrown inside the catch handler
- useful when we want to act on an error, but cannot handle and want to propagate it

```
1 try {
2     process_data(f);
3 } catch (const std::bad_alloc& e) {
4     std::cerr << "Insufficient memory for " << f.name();
5     throw; // rethrow
6 }
```



Exceptions

C++ 98

Catching everything

- sometimes we need to catch all possible exceptions
- e.g. in main, a thread, a destructor, interfacing with C, ...

```
1
2 try {
3     callUnknownFramework();
4 } catch(const std::exception& e) {
5     // catches std::exception and all derived types
6     std::cerr << "Exception: " << e.what() << std::endl;
7 } catch(...) {
8     // catches everything else
9     std::cerr << "Unknown exception type" << std::endl;
10 }
```



Exceptions

C++ 98

Stack unwinding

- all objects on the stack between a `throw` and the matching `catch` are destructed automatically
- this should cleanly release intermediate resources
- make sure you are using the RAII idiom for your own classes

```

1  class C { ... };
2  void f() {
3      C c1;
4      throw exception{};
5      // start unwinding
6      C c2; // not run
7  }
8  void g() {
9      C c3; f();
10 }

11 int main() {
12     try {
13         C c4;
14         g();
15         cout << "done"; // not run
16     } catch(const exception&) {
17         // c1, c3 and c4 have been
18         // destructed
19     }
20 }
```

Exceptions

C++ 17

Good practice: Exceptions

- use exceptions for *unlikely* runtime errors outside the program's control
 - bad inputs, files unexpectedly not found, DB connection, ...
- *don't* use exceptions for logic errors in your code
 - use assert and tests
- *don't* use exceptions to provide alternative/skip return values
 - you can use std::optional, std::variant or std::expected (C++ 23)
 - avoid using the global C-style errno
- never throw in destructors
- see also the [C++ core guidelines](#) and the [ISO C++ FAQ](#)



Exceptions

C++ 98

A more illustrative example

- exceptions are very powerful when there is much code between the error and where the error is handled
- they can also rather cleanly handle different types of errors
- **try/catch** statements can also be nested

```
1  try {                                1  void process_file(File const & file) {
2    for (File const &f : files) {          2    ...
3      try {                            3    if (handle = open_file(file))
4        process_file(f);              4      throw bad_file(file.status());
5      }                                5    while (!handle) {
6      catch (bad_file const & e) {     6      line = read_line(handle);
7        ... // loop continues       7      database.insert(line); // can throw
8      }                                8      // bad_db
9    }                                9    }
10 } catch (bad_db const & e) {           10  }
11 ... // loop aborted
```



Exceptions

C++ 98

Cost

- exceptions have little cost if no exception is thrown
 - they are recommended to report *exceptional* errors
- for performance, when error raising and handling are close, or errors occur often, prefer error codes or a dedicated class
- when in doubt about which error strategy is better, profile!

Avoid

```
for (string const &num: nums) {  
    try {  
        int i = convert(num); // can  
                           // throw  
        process(i);  
    } catch (not_an_int const &e) {  
        ... // log and continue  
    }  
}
```

Prefer

```
for (string const &num: nums) {  
    optional<int> i = convert(num);  
    if (i) {  
        process(*i);  
    } else {  
        ... // log and continue  
    }  
}
```



noexcept specifier

C++ 11

noexcept

- a function with the **noexcept** specifier states that it guarantees to not throw an exception

```
int f() noexcept;
```

- either no exceptions are thrown or they are handled internally
- if one is thrown, 'std::terminate' is called
- allows the compiler to optimize around that knowledge

- a function with **noexcept(expression)** is only **noexcept** when expression evaluates to **true** at compile-time

```
int safe_if_8B() noexcept(sizeof(long)==8);
```

Good practice: noexcept

- Use **noexcept** on leaf functions where you know the behavior
- C++11 destructors are **noexcept** - never throw from them



noexcept operator

C++ 11 Adv

- the `noexcept(expression)` operator checks at compile-time whether an expression can throw exceptions
- returns a `bool`, which is `true` if no exceptions can be thrown

```
1 constexpr bool callCannotThrow = noexcept(f());  
2 if constexpr (callCannotThrow) { ... }
```

```
3 template <typename Function>  
4 void g(Function f) noexcept(noexcept(f())) {  
5     ...  
6     f();  
7 }
```

Exceptions

C++ 98 Adv

Exception Safety Guarantees

A function can offer the following guarantees:

- No guarantee
 - An exception will lead to undefined program state
- Basic exception safety guarantee
 - An exception will leave the program in a defined state
 - At least destructors must be able to run
 - This is similar to the moved-from state
 - No resources are leaked
- Strong exception safety guarantee
 - The operation either succeeds or has no effect
 - When an exception occurs, the original state must be restored
 - This is hard to do, possibly expensive, maybe impossible
- No-throw guarantee
 - The function never throws (e.g. is marked `noexcept`)
 - Errors are handled internally or lead to termination

Exceptions

Alternatives to exceptions

- The global variable `errno` (avoid)
- Return an error code (e.g. an `int` or an `enum`)
- Return a `std::error_code` (C++ 11)
- If failing to produce a value is not a hard error, consider returning `std::optional<T>` (C++ 17)
- Return `std::expected<T, E>` (C++ 23), where T is the return type on success and E is the type on failure
- Terminate the program
 - e.g. by calling `std::terminate()` or `std::abort()`
- Contracts: Specify function pre- and postconditions. Planned for C++ 20, but removed. Will come back in the future



Further resources

- [Exceptions and Error Handling - ISO C++ FAQ](#)
- [Back to Basics: Exceptions - Klaus Iglberger - CppCon 2020](#)
- [The Unexceptional Exceptions - Fedor Pikus - CppCon 2015](#)
- [The Dawn of a New Error, \(C++ error-handling and exceptions\) - Phil Nash - CppCon 2019 \(alternatives\)](#)
- [Exceptions the Other Way Round - Sean Parent - CppNow 2022 \(theoretical approach to errors\)](#)
- [P0709: Zero-overhead deterministic exceptions: Throwing values - Herb Sutter \(proposal - distant future\)](#)
 - [De-fragmenting C++: Making Exceptions and RTTI More Affordable and Usable - Herb Sutter CppCon 2019](#)

Move semantics

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI^I and smart pointers
- Initialization



Move semantics: the problem

C++ 11 Adv

Inefficient code

```
1 void swap(std::vector<int> &a,
2             std::vector<int> &b) {
3     std::vector<int> c = a;
4     a = b;
5     b = c;
6 }
7 std::vector<int> v(10000), w(10000);
8 ...
9 swap(v, w);
```



Move semantics: the problem

C++ 11 Adv

Inefficient code

```
1 void swap(std::vector<int> &a,
2             std::vector<int> &b) {
3     std::vector<int> c = a;
4     a = b;
5     b = c;
6 }
7 std::vector<int> v(10000), w(10000);
8 ...
9 swap(v, w);
```

What happens during swap

- one allocation and one release for 10k **ints**
- a copy of 30k **ints**

Move semantics: the problem

C++ 11 Adv

Vector's built-in efficient swap

```
1 std::vector<int> v(10'000), w(10'000);  
2 ...  
3 v.swap(w);
```



Move semantics: the problem

C++ 11 Adv

Vector's built-in efficient swap

```
1 std::vector<int> v(10'000), w(10'000);  
2 ...  
3 v.swap(w);
```

What happens during swap

- 3 swaps of `int*` (9 copies)
- only some pointers to the underlying storage are swapped



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> input) {  
2     // ... change elements, compute result  
3     return result;  
4 }  
5  
6 std::vector<int> values{...};  
7 consumeVector(values); // being copied now
```



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> input) {  
2     // ... change elements, compute result  
3     return result;  
4 }  
5  
6 std::vector<int> values{...};  
7 consumeVector(values); // being copied now
```

Pass by copy

- One unnecessary (large?) allocation and release
- Unnecessary copy of the **ints**



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> const & input) {  
2     std::vector tmp{input};  
3     // ... change elements, compute result  
4     return result;  
5 }  
6 std::vector<int> values{...};  
7 consumeVector(values); // maybe copied internally?
```



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> const & input) {  
2     std::vector tmp{input};  
3     // ... change elements, compute result  
4     return result;  
5 }  
6 std::vector<int> values{...};  
7 consumeVector(values); // maybe copied internally?
```

Use references to avoid copies?

- Working with pass by reference only moves allocation and copy to a different place



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> & input) {  
2     // ... change elements, compute result  
3     return result;  
4 }  
5  
6 std::vector<int> values{...};  
7 consumeVector(values); // maybe modified?  
8 somethingElse(values); // safe to use values now???
```

So non-const references?

- Non-**const** references could work, but does the outside know that we're changing the vector?
- Could lead to bugs



Move semantics: the problem

C++ 98 Adv

Another potentially inefficient code

```
1 T consumeVector(std::vector<int> & input) {  
2     // ... change elements, compute result  
3     return result;  
4 }  
5  
6 std::vector<int> values{...};  
7 consumeVector(values); // maybe modified?  
8 somethingElse(values); // safe to use values now???
```

The ideal situation

Have a way to express that we move the vector's content



Move semantics: value categories

C++ 11 Adv

Lvalue (left value / locator value)

- An expression is an lvalue
 - If its evaluation determines identity of object/function
 - If the object is not expiring (“xvalue”)
- Lvalues have a persistent address in memory
- Lvalues can be assigned to unless **const**

Lvalue examples

```
1 int i;
2 int & j = i;
3 T t;
4 t.m = 5; // lvalue, because t is lvalue
5 ++i;      // returns underlying object
```



Move semantics: value categories

C++ 11 Adv

Rvalue

- An expression is an rvalue when from these primary categories
 - “pure rvalue”: Result of built-in operator; initialiser of object
 - “xvalue”: Expiring value, resources can “moved”
- Has no persistent address in memory
 - Cannot use the addressof operator
- Cannot be used as left-hand operand of built-in assignments

Examples for rvalues

```
1 i = 42;      // 42 is pure rv
2 i = a + b;  // a+b is pure rv
3 T t = T{};  // T{} is pure rv
4 i++;        // Pure rv (post-increment returns temp.)
5 T{}.m;       // m is expiring because T{} is pure rv
6 std::move(x); // Converts lvalue into xvalue
```



Move semantics: value categories

C++ 11 Adv

Rvalue

- An expression is an rvalue when from these primary categories
 - “pure rvalue”: Result of built-in operator; initialiser of object
 - “xvalue”: Expiring value, resources can “moved”
- Has no persistent address in memory
 - Cannot use the addressof operator
- Cannot be used as left-hand operand of built-in assignments

Things that don't compile with rvalues

- 1 `a + b = 2; // Cannot assign to pure rv`
- 2 `&(a+b); // Cannot take address of pure rv`
- 3 `&i++; // Same, post-increment returns pure rv`
- 4 `T{}.m = 2; // cannot assign if m is built-in type`



Move semantics: Terminology

C++ 11 Adv

Source of confusion

- The name “rvalue” originates from them *mostly* being on right-hand side
- But it denotes value categories of *expressions*
- Rvalues can be assigned to for non-builtin types

Example: rvalue on left-hand side

```
1 struct T{  
2     int m;  
3     T& operator=(int i);  
4 };  
5 std::cout << (T{} = 5).m << "\n";
```

Not possible with built-in types

```
1 std::cout << (int{} = 5) << "\n"; // Error
```



Move semantics: rvalue references

C++ 11 Adv

New reference type: rvalue reference

- Declared with `T&&`
- Does not bind to lvalues

Binding rules

Lvalues

```

T lv;
T &      ref = lv; //OK
T &&    rvref = lv; //Error
T const& cr = lv; //OK

```

Rvalues

```

T &      ref = T{}; //Error
T &&    rvref = T{}; //OK
T const& cr = T{}; //OK

```

Overload resolution prefers rvalue reference

```

1 void f(T const &);
2 void f(T &&);      // Selected. Enables move
3 f(T{});

```



Move semantics

C++ 11 Adv

The idea

- use rvalue references to reuse resources
- 2 new special member functions in every class:
 - a **move constructor** similar to copy constructor
 - a **move assignment operator** similar to assignment operator
(now called copy assignment operator)



Move semantics

C++ 11 Adv

The idea

- use rvalue references to reuse resources
- 2 new special member functions in every class:
 - a **move constructor** similar to copy constructor
 - a **move assignment operator** similar to assignment operator
(now called copy assignment operator)

Practically

```
1 T(T const & other); // copy construction
2 T(      T&& other); // move construction
3 T& operator=(T const & other); // copy assignment
4 T& operator=(      T&& other); // move assignment
```



Move semantics

C++ 11 Adv

A few points

- move constructor and assignment operator are allowed to leave the source object "empty"
 - so do not use the source object afterward
 - leave the source in a valid state (for its destructor)
- if no move semantic is implemented, copies will be performed
- the language and STL understand move semantic
- the compiler uses copy elision or moves whenever possible
 - e.g. when passing temporaries or returning from a function



Move semantics

C++ 11 Adv

A few points

- move constructor and assignment operator are allowed to leave the source object "empty"
 - so do not use the source object afterward
 - leave the source in a valid state (for its destructor)
- if no move semantic is implemented, copies will be performed
- the language and STL understand move semantic
- the compiler uses copy elision or moves whenever possible
 - e.g. when passing temporaries or returning from a function

Practically (in C++ 11)

```
1 T f() { T r; return r; } // move r out of f
2 T v = f(); // move returned (temporary) T into v
3 void g(T a, T b, T c);
4 g(f(), T{}, v); // move, move, copy
```



Move semantics

C++ 11 Adv

In some cases, you want to force a move

```
1 void swap(T &a, T &b) {  
2     T c = a;    // copy construct  
3     a = b;    // copy assign  
4     b = c;    // copy assign  
5 }
```



Move semantics

C++ 11 Adv

In some cases, you want to force a move

```
1 void swap(T &a, T &b) {  
2     T c = a;    // copy construct  
3     a = b;    // copy assign  
4     b = c;    // copy assign  
5 }
```

Explicitly request moving

- using the `std::move` function
- which is basically a cast to an rvalue reference

```
6 void swap(T &a, T &b) {  
7     T c = std::move(a);        // move construct  
8     a = std::move(b);        // move assign  
9     b = static_cast<T&>(c); // move assign (don't)  
10 }
```



Move semantics: recommended implementation C++ 11 Adv

Use copy and swap idiom

- implement an efficient swap function for your class
 - preferably hidden friend and symmetric
- move constructor
 - consider delegating to default constructor
 - swap `*this` with parameter (source)
- move assignment as `operator=(T source)`
 - parameter passed by value; caller can move or copy into it
 - swap parameter with `*this`
 - end of scope: parameter destroys former content of `*this`
- alternative: move assignment as `operator=(T&& source)`
 - swap parameter with `*this`
 - 1 swap less, separate copy assignment operator needed
 - former content of `*this` destroyed with caller argument
- swap, move constructor/assignment must be `noexcept!`



Move semantics: recommended implementation C++ 11 Adv

Practically

```
1 class Movable {
2     Movable();
3     Movable(const Movable &other);
4     Movable(Movable &&other) noexcept :
5         Movable() {           // constructor delegation
6             swap(*this, other);
7         }
8     Movable& operator=(Movable other) noexcept { // by value
9         swap(*this, other);
10        return *this;
11    }
12    friend void swap(Movable &a, Movable &b) noexcept {...}
13};
14 Movable a, b;
15 a = b;           // operator= copies b into "other"
16 a = std::move(b); // operator= moves b into "other"
```



Move semantics: alternative implementation

C++ 11 Adv

Practically

```
1 class Movable {
2     Movable();
3     Movable(const Movable &other);
4     Movable(Movable &&other) noexcept :
5         Movable() {           // constructor delegation
6             swap(*this, other);
7         }
8     Movable& operator=(const Movable& other);
9     Movable& operator=(Movable&& other) noexcept {
10         swap(*this, other);
11         return *this;
12     }
13     friend void swap(Movable &a, Movable &b) noexcept { ... }
14 };
```



Move semantics: Don't forget noexcept!

C++ 11 Adv

Vector pessimisation when move constructor can throw

- When a vector reallocates, it must copy/move all elements
 - An exception during move aborts the relocation of elements
 - Move might not be revertible, so vector falls back to copy
 - This can significantly slow down the resize operation
- With **noexcept**, the much faster move is used

Vector resize with copy/move

```
1 struct Movable1 { Movable1(Movable1 &&other); };
2 struct Movable2 { Movable2(Movable2 &&other) noexcept; };
3 while (vector1.size() < 10000) {
4     vector1.push_back(Movable1{}); // Copies
5     vector2.push_back(Movable2{}); // Moves
6 }
```



Move Semantic

C++ 11 Adv

Exercise: Move semantics

- go to exercises/move
- look at the code and run it with callgrind
- understand how inefficient it is
- implement move semantic the easy way in NVector
- run with callgrind and see no improvement
- understand why and fix test.cpp
- see efficiency improvements

prerequisite: be able to use simple templated code



Copy elision

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization

Copy elision

C++ 17 Adv

Where does it take place ?

```
1 struct Foo { ... };
2 Foo f() {
3     return Foo(); // return-value opt. (RVO)
4 }
5 Foo g() {
6     Foo foo;
7     return foo; // named return-value opt. (NRVO)
8 }
9 int main() {
10    // compiler must avoid these copies:
11    Foo a = f();
12    Foo b = Foo();
13    Foo c = Foo(Foo());
14    // copy elision allowed, but not guaranteed:
15    Foo d = g();
16 }
```



Guaranteed copy elision

C++ 17 Adv

Allows to write code not allowed in C++ 14 (would not compile)

Guaranteed: Return value optimization (RVO)

```
1 struct Foo {  
2     Foo() { ... }  
3     Foo(const Foo &) = delete;  
4     Foo(Foo &&) = delete;  
5     Foo& operator=(const Foo &) = delete;  
6     Foo& operator=(Foo &&) = delete;  
7 };  
8 Foo f() {  
9     return Foo(); // ok, no move  
10 }  
11 int main() {  
12     Foo foo = f(); // ok, no move  
13 }
```

Templates

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization



Templates

C++ 17

Concept

- The C++ way to write reusable code
 - like macros, but fully integrated into the type system
- Applicable to functions, classes and variables

```
1 template<typename T>
2 const T & max(const T &a, const T &b) {
3     return b < a ? a : b;
4 }
5 template<typename T>
6 struct Vector {
7     int m_len;
8     T* m_data;
9 };
10 template <typename T>
11 std::size_t size = sizeof(T);
```




Templates

C++ 98

Notes on templates

- they are compiled for each instantiation
- they need to be defined before used
 - so all template code must typically be in headers
 - or declared to be available externally (`extern template`)
- this may lead to longer compilation times and bigger binaries



Templates

C++ 98

Template parameters

- can be types, values or other templates
- you can have several
- default values allowed starting at the last parameter

```
1 template<typename KeyType=int, typename ValueType=KeyType>
2 struct Map {
3     void set(const KeyType &key, ValueType value);
4     ValueType get(const KeyType &key);
5     ...
6 };
7
8 Map<std::string, int> m1;
9 Map<float> m2;    // Map<float, float>
10 Map<> m3;        // Map<int, int>
11 Map m4;           // Map<int, int>, C++17
```



Template parameters

C++ 98

typename vs. class keyword

- for declaring a template type parameter, the **typename** and **class** keyword are semantically equivalent
- template template parameters require C++ 17 for **typename**

```
1 template<typename T>
2 T func(T a); // equivalent to:
3 template<class T>
4 T func(T a);
5
6 template<template<class> class C>
7 C<int> func(C<int> a); // equivalent to:
8 template<template<typename> class C>
9 C<int> func(C<int> a); // equivalent to:
10 template<template<typename> typename C> // C++17
11 C<int> func(C<int> a);
```

Template implementation

C++ 98

```
1 template<typename KeyType=int, typename ValueType=KeyType>
2 struct Map {
3     // declaration and inline definition
4     void set(const KeyType &key, ValueType value) {
5         ...
6     }
7     // just declaration
8     ValueType get(const KeyType &key);
9 };
10
11 // out-of-line definition
12 template<typename KeyType, typename ValueType>
13 ValueType Map<KeyType, ValueType>::get
14     (const KeyType &key) {
15     ...
16 }
```

Nested templates

C++ 98 Adv

Nested templates - godbolt

```
1 template<typename KeyType=int, typename ValueType=KeyType>
2 struct Map {
3     template<typename OtherValueType>
4     void set(const KeyType &key, OtherValueType value) {
5         ...
6     }
7     template<typename OtherKeyType>
8     ValueType get(const OtherKeyType &key);
9 };
10
11 template<typename KeyType, typename ValueType> //for class
12 template<typename OtherKeyType> //for member function
13 ValueType Map<KeyType, ValueType>::get
14     (const OtherKeyType &key) { ... }
```



Non-type template parameter C++ 98 / C++ 17 / C++ 20

template parameters can also be values

- integral types, pointer, enums in C++ 98
- `auto` in C++ 17
- literal types (includes floating points) in C++ 20

```
1 template<unsigned int N>
2 struct Polygon {
3     float perimeter() {
4         return 2 * N * std::sin(PI / N) * radius;
5     }
6     float radius;
7 };
8
9 Polygon<19> nonadecagon{3.3f};
```



Template specialization

C++ 98

Specialization

Templates can be specialized for given values of their parameter

```
1 template<typename F, unsigned int N>
2 struct Polygon { ... }; // primary template
3
4 template<typename F> // partial specialization
5 struct Polygon<F, 6> {
6     F perimeter() { return 6 * radius; }
7     F radius;
8 };
9 template<> // full specialization
10 struct Polygon<int, 6> {
11     int perimeter() { return 6 * radius; }
12     int radius;
13 };
```

Template argument deduction

C++ 98 Adv

Template argument deduction

- Template arguments deduced from (function) arguments
- Template arguments can always be specified explicitly
- Only for function templates (Before C++ 14)

```
1 template <typename T, typename U>
2 void f(T t, U u) { ... }
3
4 f(42, true);           // deduces T=int, U=bool
5                         // calls f<int, bool>(42, true);
6 f<float>(42, true);  // sets T=float, deduces U=bool
7                         // calls f<float, bool>(42, true);
8                         // 42 converted to float before call
```



Template argument deduction

C++ 11 Adv

Deduced contexts

- Compiler can even deduce template arguments inside certain expressions (pattern matching)
- See [cppreference](#) for details

```
1 template <typename T>
2 void f(T* p) { ... }
3
4 const int * ip = ...;
5 f(ip); // deduces T=const int
6
7 template <typename T, std::size_t N>
8 void g(std::array<T*, N> a) { ... }
9
10 std::array<int*, 3> aip = ...;
11 g(aip); // deduces T=int, N=3
```

Template argument deduction

C++ 11 Adv

Non-deduced contexts

- Deduction from certain expressions is impossible/forbidden

```
1 template <typename C>
2 void f(typename C::value_type v) { ... }
3 f(std::vector<int>{...}); // cannot deduce C
4 // from a dependent type
5 template <typename T, std::size_t N>
6 void g(std::array<T, N * 2> a) { ... }
7 g(std::array<int, 4>{...}); // deduces T=int,
8 // cannot deduce N from expression
9 template <typename T>
10 void h(std::vector<T> v) { ... }
11 h({1, 2, 3}); // error, braced-initializer list has no type
12 h(std::vector<int>{1, 2, 3}); // ok, T=int
```



Template argument deduction

C++ 11 Adv

Non-deduced contexts

- Deduction from certain expressions is impossible/forbidden

```
1 template <typename C, typename F>
2 void reduce(const C& cont, const F& f = std::plus{});  
3 reduce(std::vector<int>{...}); // error: cannot deduce F  
4 // would need: <typename C, typename F = decltype(std::plus{})>  
5  
6 template<typename T>
7 const T & max(const T &a, const T &b) { ... }  
8 int i = 3;  
9 max(i, 3.14f); // deduces T=int and T=float, error  
10 // either: max<float>(i, 3.14f);  
11 // or:     max(static_cast<float>(i), 3.14f);
```



Template argument deduction

C++ 11 Adv

Deduction in partial specializations

- Partial specializations also deduce template arguments
 - with similar rules as for function templates
 - but some more restrictions (cf. [cppreference](#))

```
1 template <typename T>
2 struct S { ... }; // primary template
3
4 template <typename T>
5 struct S<T*> { ... }; // specialization 1
6
7 template <typename U, std::size_t N>
8 struct S<std::array<U*, N>> { ... }; // specialization 2
9
10 S<int> s1; // prim. tmpl. (T=int) ok, spec. 1/2 invalid
11 S<int*> s2; // prim. tmpl. (T=int*) and spec. 1 (T=int) ok
12 // spec. 1 is more specialized, will be used
```



Template argument deduction

C++ 98 Adv

Specialization vs. overloading

- For function templates we can choose between specialization and overloading
- Partial specialization of function templates is forbidden

```
1 template <typename T>
2 void f(T t) { ... }
3
4
5 void f(int* t) { ... }
6
7
8 template <typename T>
9 void f(T* t) { ... }
```

```
10 template <typename T>
11 void f(T t) { ... }
12
13 template <>
14 void f<int*>(int* t) { ... }
15
16 // part. spec. forbidden:
17 template <typename T>
18 void f<T*>(T* t) {...}
```

Template argument deduction

C++ 11 Adv

Disadvantages of specialization vs. overloading

- Specialization always needs a primary template
 - Sometimes this does not make sense
- Partial specializations of function templates is forbidden
 - So we need SFINAE workarounds or concepts

Could you express this with specializations?

```
1 template <typename T>
2 void f(T* p) { ... }
3
4 template <typename T>
5 void f(std::unique_ptr<T> p) { ... }
```

Good practice: Specialization vs. overloading

Prefer overloading function templates over template specialization



Class Template Argument Deduction (CTAD) C++ 17 Adv

CTAD

- Deduce the template arguments for a class template
- Based on construction arguments
- Only when no template arguments provided
- Since C++ 20: CTAD for aggregates (no constructor needed)

Practically - godbolt

```
1 template<typename A, typename B, typename C = double>
2 struct Triple {
3     Triple(A a, B b, C c) : a(a), b(b), c(c) {}
4     A a; B b; C c;
5 };
6 Triple t{42, true, 3.14}; // Triple<int, bool, double>
7 Triple<int> t2{42, true, 3.14}; // compilation error
8 Triple<int, bool> t3{42, true, 3.14}; // not CTAD
```

Class Template Argument Deduction (CTAD) C++ 17 Adv

Deduction guides

- Describe how constructor argument types are mapped to class template arguments

```
1 template<typename A, typename B>
2 struct Pair {
3     Pair(A a, B b) : a(a), b(b) {}
4     A a; B b;
5 };
6
7
8
9 Pair p{42, "hello"}; // Pair<int, const char*>
```



Class Template Argument Deduction (CTAD) C++ 17 Adv

Deduction guides

- Describe how constructor argument types are mapped to class template arguments

```
1 template<typename A, typename B>
2 struct Pair {
3     Pair(A a, B b) : a(a), b(b) {}
4     A a; B b;
5 };
6 template<typename A>
7 Pair(A, const char*) -> Pair<A, std::string>;
8
9 Pair p{42, "hello"}; // Pair<int, std::string>
```



Class Template Argument Deduction (CTAD) C++ 17 Adv

Standard library examples

```
1 std::pair p{1.2, true}; // std::pair<double, bool>
2 std::tuple t{1.2, true, 32};
3 // std::tuple<double, bool, int>
4 std::vector v{1, 2, 3}; // std::vector<int>
5 std::list l{v.begin(), v.end()}; // std::list<int>
6 std::array a{1, 2, 3}; // std::array<int, 3>
7
8 std::mutex m;
9 std::lock_guard l{m}; // std::lock_guard<std::mutex>
```



The full power of templates

C++ 98

Exercise: Templates

- go to exercises/templates
- look at the OrderedVector code
- compile and run playwithsort.cpp. See the ordering
- modify playwithsort.cpp and reuse OrderedVector with Complex
- improve OrderedVector to template the ordering
- test reverse ordering of strings (from the last letter)
- test order based on Manhattan distance with complex type
- check the implementation of Complex
- try ordering complex of complex



Lambdas

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- **Lambdas**
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization



Lambda expressions

C++ 11

Definition

A lambda expression is a function with no name



Lambda expressions

C++ 11

Definition

A lambda expression is a function with no name

Python example

```
1 data = [1,9,3,8,3,7,4,6,5]
2
3 # without lambdas
4 def isOdd(n):
5     return n%2 == 1
6 print(filter(isOdd, data))
7
8 # with lambdas
9 print(filter(lambda n:n%2==1, data))
```



C++ Lambdas

C++ 11

Simplified syntax

```
1 auto f = [] (arguments) -> return_type {  
2     statements;  
3 };
```

- The return type specification is optional
- `f` is an instance of a functor type, generated by the compiler

Usage example

```
4 int data[] {1,2,3,4,5};  
5 auto f = [](int i) {  
6     std::cout << i << " squared is " << i*i << '\n';  
7 };  
8 for (int i : data) f(i);
```



Trailing function return type

C++ 11

An alternate way to specify a function's return type

```
int f(float a); // classic
auto f(float a) -> int; // trailing
auto f(float a) { return 42; } // deduced, C++14
```



Trailing function return type

C++ 11

An alternate way to specify a function's return type

```
int f(float a); // classic
auto f(float a) -> int; // trailing
auto f(float a) { return 42; } // deduced, C++14
```

When to use trailing return type

- Only way to specify return type for lambdas
- Allows to simplify inner type definition

```
1 class Equation {
2     using ResultType = double;
3     ResultType evaluate();
4 }
5 Equation::ResultType Equation::evaluate() {...}
6 auto Equation::evaluate() -> ResultType {...}
```



Capturing variables

C++ 11

Adaptable lambdas

- Adapt lambda's behaviour by accessing variables outside of it
- This is called “capture”



Capturing variables

C++ 11

Adaptable lambdas

- Adapt lambda's behaviour by accessing variables outside of it
- This is called “capture”

First attempt in C++

```
1 int increment = 3;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [] (int x) { return x+increment; };
4 for(int& i : data) i = f(i);
```



Capturing variables

C++ 11

Adaptable lambdas

- Adapt lambda's behaviour by accessing variables outside of it
- This is called “capture”

First attempt in C++

```
1 int increment = 3;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [] (int x) { return x+increment; };
4 for(int& i : data) i = f(i);
```

Error

```
error: 'increment' is not captured
[] (int x) { return x+increment; );
```



Capturing variables

C++ 11

The capture list

- local variables outside the lambda must be explicitly captured
 - unlike in Python, Java, C#, Rust, ...
- captured variables are listed within initial []



Capturing variables

C++ 11

The capture list

- local variables outside the lambda must be explicitly captured
 - unlike in Python, Java, C#, Rust, ...
- captured variables are listed within initial []

Example

```
1 int increment = 3;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [increment](int x) { return x+increment; };
4 for(int& i : data) i = f(i);
```



Default capture is by value

C++ 11

Code example

```
1 int sum = 0;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [sum](int x) { sum += x; };
4 for (int i : data) f(i);
```



Default capture is by value

C++ 11

Code example

```
1 int sum = 0;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [sum](int x) { sum += x; };
4 for (int i : data) f(i);
```

Error

```
error: assignment of read-only variable 'sum'
        [sum](int x) { sum += x; });
```



Default capture is by value

C++ 11

Code example

```
1 int sum = 0;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [sum](int x) { sum += x; };
4 for (int i : data) f(i);
```

Error

```
error: assignment of read-only variable 'sum'
        [sum](int x) { sum += x; });
```

Explanation

- By default, variables are captured by value
- The lambda's **operator()** is **const inline**



Capture by reference

C++ 11

Simple example

In order to capture by reference, add '&' before the variable

```
1 int sum = 0;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [&sum](int x) { sum += x; };
4 for (int i : data) f(i);
```



Capture by reference

C++ 11

Simple example

In order to capture by reference, add '&' before the variable

```
1 int sum = 0;
2 int data[] {1,9,3,8,3,7,4,6,5};
3 auto f = [&sum](int x) { sum += x; };
4 for (int i : data) f(i);
```

Mixed case

One can of course mix values and references

```
5 int sum = 0, off = 1;
6 int data[] {1,9,3,8,3,7,4,6,5};
7 auto f = [&sum, off](int x) { sum += x + off; };
8 for (int i : data) f(i);
```



Capture by value vs. by reference

C++ 11

See the difference between val and ref

```
1 int data[] {1,9,3,8,3,7,4,6,5};  
2 int increment = 3;  
3 auto val = [ inc](int x) { return x+inc; };  
4 auto ref = [&inc](int x) { return x+inc; };  
5  
6 increment = 4;  
7  
8 for(int& i : data) i = val(i); // increments by 3  
9 for(int& i : data) i = ref(i); // increments by 4
```



Capture with an initializer

In C++ 14, can declare captures with initializers

```
1 auto f = [inc = 1+2](int x) { return x+inc; };
2 auto g = [inc = getInc()] (int x) { return x+inc; };
3 for(int& i : data) i = f(i); // increments by 3
4 for(int& i : data) i = g(i); // unknown increment
```



Anatomy of a lambda

C++ 11

Lambdas are pure syntactic sugar - `cppinsight`

- They are replaced by a functor during compilation

```
1 int sum = 0, off = 1;
2 auto l =
3 [&sum, off]
4
5 (int x) {
6     sum += x + off;
7 }
8
9 l(42);
```

```
13 int sum = 0, off = 1;
14 struct __lambda4 {
15     int& sum; int off;
16     __lambda4(int& s, int o)
17         : sum(s), off(o) {}
18
19     auto operator()(int x) const {
20         sum += x + off;
21     }
22 };
23 auto l = __lambda4{sum, off};
24 l(42);
```

Some nice consequences

- Lambda expressions create ordinary objects
- They can be copied, moved, or inherited from



Capture list

C++ 11

all by value

```
[=] (...) { ... };
```



Capture list

C++ 11

all by value

```
[=] (...) { ... };
```

all by reference

```
[&] (...) { ... };
```



Capture list

C++ 11

all by value

```
[=] (...) { ... };
```

all by reference

```
[&] (...) { ... };
```

mix

```
[&, b] (...) { ... };  
[=, &b] (...) { ... };
```

Capture list - this

C++ 11 Adv

Inside a (non-static) member function, we can capture **this**.

```
1 [    this](...) { use(*this); };
2 [      &] (...) { use(*this); };
3 [&, this] (...) { use(*this); };
4 [=] (...) { use(*this); } // deprecated in C++20
5 [=, this] (...) { use(*this); } // allowed in C++20
```

Since the captured **this** is a pointer, ***this** refers to the object by reference.



Capture list - this

C++ 11 Adv

Inside a (non-static) member function, we can capture **this**.

```
1 [    this](...) { use(*this); };
2 [      &] (...) { use(*this); };
3 [&, this] (...) { use(*this); };
4 [=] (...) { use(*this); } // deprecated in C++20
5 [=, this] (...) { use(*this); } // allowed in C++20
```

Since the captured **this** is a pointer, ***this** refers to the object by reference.

```
1 [*this] (...) { use(*this); } // C++17
2 [&, *this] (...) { use(*this); } // C++17
3 [=, *this] (...) { use(*this); } // C++17
```

The object at ***this** is captured by value (the lambda gets a copy).

Details in [this blog post](#).

Generic lambdas

C++ 14 Adv

Generic lambdas (aka. polymorphic lambdas)

- The type of lambda parameters may be `auto`.

```
auto add = [](auto a, auto b) { return a + b; };
```

- The generated `operator()` becomes a template function:

```
template <typename T, typename U>
auto operator()(T a, U b) const { return a + b; }
```

- The types of `a` and `b` may be different.

Explicit template parameters (C++ 20)

```
auto add = []<typename T>(T a, T b)
{ return a + b; };
```

- The types of `a` and `b` must be the same.



Higher-order lambdas

C++ 11 Adv

Example - godbolt

```
1 auto build_incremoter = [](int inc) {
2     return [inc](int value) { return value + inc; };
3 };
4 auto inc1 = build_incremoter(1);
5 auto inc10 = build_incremoter(10);
6 int i = 0;
7 i = inc1(i);    // i = 1
8 i = inc10(i);   // i = 11
```

How it works

- `build_incremoter` returns a function object
- this function's behavior depends on a parameter
- note how `auto` is useful here!

Lambda improvements

C++ 20 Adv

Lambda improvements in C++ 20

- Allowed in unevaluated contexts,
e.g. within `decltype`, `sizeof`, `typeid`, etc.
- Without captures, are default-constructible and assignable

Example - godbolt

```
1 struct S {  
2     decltype([](int i) { std::cout << i; }) f;  
3 } s;  
4 s.f(42); // prints "42"  
5  
6 std::set<T, decltype([](T a, T b) { ... })> s2;
```



The STL

4

Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization



The Standard Template Library

C++ 98

What it is

- A library of standard templates
- Has almost everything you need
 - strings, containers, iterators
 - algorithms, functions, sorters
 - functors, allocators
 - ...
- Portable
- Reusable
- Efficient



The Standard Template Library

C++ 98

What it is

- A library of standard templates
- Has almost everything you need
 - strings, containers, iterators
 - algorithms, functions, sorters
 - functors, allocators
 - ...
- Portable
- Reusable
- Efficient

Use it

and adapt it to your needs, thanks to templates



STL in practice

C++ 14

STL example - godbolt

```
1 #include <vector>
2 #include <algorithm>
3 #include <functional>      // `import std;` in C++23
4 #include <iterator>
5 #include <iostream>
6
7 std::vector<int> in{5, 3, 4};    // initializer list
8 std::vector<int> out(3);        // constructor taking size
9 std::transform(in.begin(), in.end(),           // input range
10               out.begin(),             // start result
11               std::negate{});          // function obj
12 std::copy(out.begin(), out.end(),           // -5 -3 -4
13           std::ostream_iterator<int>{std::cout, " "});
```



STL's concepts

C++ 98

containers

- data structures for managing a range of elements, irrespective of:
 - the data itself (templated)
 - the memory allocation of the structure (templated)
 - the algorithms that may use the structure (iterators)

Examples (→ string and container library on cppreference)

- string, string_view (C++ 17)
- list, forward_list (C++ 11), vector, deque, array (C++ 11)
- [multi]map, [multi]set (C++ 23: flat_[multi]map, flat_[multi]set)
- unordered_[multi]map (C++ 11), unordered_[multi]set (C++ 11)
- stack, queue, priority_queue
- span (C++ 20)
- non-containers: bitset, pair, tuple (C++ 11), optional (C++ 17), variant (C++ 17), any (C++ 17), expected (C++ 23)



Containers: std::vector

C++ 11

```

1 #include <vector>
2 std::vector<T> v{5, 3, 4}; // 3 Ts, 5, 3, 4
3 std::vector<T> v(100);    // 100 default constr. Ts
4 std::vector<T> v(100, 42); // 100 Ts with value 42
5 std::vector<T> v2 = v;           // copy
6 std::vector<T> v2 = std::move(v); // move, v is empty
7
8 std::size_t s = v.size();
9 bool empty = v.empty();

10
11 v[2] = 17;           // write element 2
12 T& t = v[1000];    // access element 1000, bug!
13 T& t = v.at(1000); // throws std::out_of_range
14 T& f = v.front();   // access first element
15 v.back() = 0;        // write to last element
16 T* p = v.data();    // pointer to underlying storage

```



Containers: std::vector

C++ 11

```
1 std::vector<T> v = ...;
2 auto b = v.begin(); // iterator to first element
3 auto e = v.end();   // iterator to one past last element
4 // all following operations, except reserve, invalidate
5 // all iterators (b and e) and references to elements
6
7 v.resize(100); // size changes, grows: new T{}s appended
8                   //           shrinks: Ts at end destroyed
9 v.reserve(1000); // size remains, memory increased
10 for (T i = 0; i < 900; i++)
11     v.push_back(i); // add to the end
12 v.insert(v.begin() + 3, T{});
13
14 v.pop_back(); // removes last element
15 v.erase(v.end() - 3); // removes 3rd-last element
16 v.clear(); // removes all elements
```



Containers: std::unordered_map

C++ 11

Conceptually a container of std::pair<Key const, Value>

```

1 #include <unordered_map>
2 std::unordered_map<std::string, int> m;
3 m["hello"] = 1; // inserts new key, def. constr. value
4 m["hello"] = 2; // finds existing key
5 auto [it, isNewKey] = m.insert({"hello", 0}); // no effect
6 // ^ C++17: "Structured binding"
7 int val = m["world"]; // inserts new key (val == 0)
8 int val = m.at("monde"); // throws std::out_of_range
9
10 if (auto it = m.find("hello"); it != m.end()) // C++17
11     m.erase(it); // remove by iterator (fast)
12 if (m.contains("hello")) // C++20
13     m.erase("hello"); // remove by key, 2. lookup, bad
14 for (auto const& [k, v] : m) // iterate k/v pairs (C++17)
15     std::cout << k << ":" << v << '\n';

```



- The standard utility to create hash codes
- Used by std::unordered_map and others
- Can be customized for your types via template specialization

```
1 #include <functional>
2 std::hash<std::string> h;
3 std::cout << h("hello"); // 2762169579135187400
4 std::cout << h("world"); // 8751027807033337960
5
6 class MyClass { int a, b; ... };
7 template<> struct std::hash<MyClass> {
8     std::size_t operator()(MyClass const& c) {
9         std::hash<int> h;
10        return h(c.a) ^ h(c.b); // xor to combine hashes
11    }
12};
```

STL's concepts

C++ 11

iterators

- generalization of pointers
- allow iteration over some data, irrespective of:
 - the container used (templated)
 - the data itself (container is templated)
 - the consumer of the data (templated algorithm)
- examples
 - `std::reverse_iterator`, `std::back_insert_iterator`, ...

Iterator example - godbolt

```
1 std::vector<int> const v = {1,2,3,4,5,6,7,8,9};  
2 auto const end = v.rend() - 3; // arithmetic  
3 for (auto it = v.rbegin();  
4     it != end;      // compare positions  
5     it += 2)        // jump 2 positions  
6     std::cout << *it; // dereference, prints: 975
```



STL's concepts

C++ 98

algorithms

- implementation of an algorithm working on data
- with a well defined behavior (defined complexity)
- irrespective of
 - the data handled
 - the container where the data live
 - the iterator used to go through data (almost)
- examples
 - `for_each`, `find`, `find_if`, `count`, `count_if`, `search`
 - `copy`, `swap`, `transform`, `replace`, `fill`, `generate`
 - `remove`, `remove_if`
 - `unique`, `reverse`, `rotate`, `shuffle`, `partition`
 - `sort`, `partial_sort`, `merge`, `make_heap`, `min`, `max`
 - `lexicographical_compare`, `iota`, `reduce`, `partial_sum`
- see also [105 STL Algorithms in Less Than an Hour](#) and the [algorithms library](#) on cppreference



STL's concepts

C++ 98

functors / function objects

- generic utility functions
- as structs with `operator()`
- mostly useful to be passed to STL algorithms
- implemented independently of
 - the data handled (templated)
 - the context (algorithm) calling it
- examples
 - plus, minus, multiplies, divides, modulus, negate
 - equal_to, less, greater, less_equal, ...
 - logical_and, logical_or, logical_not
 - bit_and, bit_or, bit_xor, bit_not
 - identity, not_fn
 - bind, bind_front
- see also documentation on [cppreference](#)



Functors / function objects

C++ 11

Example

```
1 struct Incrementer {
2     int m_inc;
3     Incrementer(int inc) : m_inc(inc) {}
4
5     int operator()(int value) const {
6         return value + m_inc;
7     }
8 };
9 std::vector<int> v{1, 2, 3};
10 const auto inc = 42;
11 std::transform(v.begin(), v.end(), v.begin(),
12                 Incrementer{inc});
```



Prefer lambdas over functors

C++ 11

With lambdas

```
1 std::vector<int> v{1, 2, 3};  
2 const auto inc = 42;  
3 std::transform(begin(v), end(v), begin(v),  
4 [inc](int value) {  
5     return value + inc;  
6 });
```



Prefer lambdas over functors

C++ 11

With lambdas

```
1 std::vector<int> v{1, 2, 3};  
2 const auto inc = 42;  
3 std::transform(begin(v), end(v), begin(v),  
4                 [inc](int value) {  
5                     return value + inc;  
6                 });
```

Good practice: Use STL algorithms with lambdas

- Prefer lambdas over functors when using the STL
- Avoid binders like `std::bind2nd`, `std::ptr_fun`, etc.



Range-based for loops with STL containers

C++ 11

Iterator-based loop (since C++ 98)

```
1 std::vector<int> v = ...;
2 int sum = 0;
3 for (std::vector<int>::iterator it = v.begin();
4       it != v.end(); it++)
5     sum += *it;
```



Range-based for loops with STL containers

C++ 11

Iterator-based loop (since C++ 98)

```
1 std::vector<int> v = ...;
2 int sum = 0;
3 for (std::vector<int>::iterator it = v.begin();
4       it != v.end(); it++)
5     sum += *it;
```

Range-based for loop (since C++ 11)

```
6 std::vector<int> v = ...;
7 int sum = 0;
8 for (auto a : v) { sum += a; }
```



Range-based for loops with STL containers

C++ 11

Iterator-based loop (since C++ 98)

```
1 std::vector<int> v = ...;
2 int sum = 0;
3 for (std::vector<int>::iterator it = v.begin();
4       it != v.end(); it++)
5     sum += *it;
```

Range-based for loop (since C++ 11)

```
6 std::vector<int> v = ...;
7 int sum = 0;
8 for (auto a : v) { sum += a; }
```

STL way (since C++ 98)

```
9 std::vector<int> v = ...;
10 int sum = std::accumulate(v.begin(), v.end(), 0);
11 // std::reduce(v.begin(), v.end()); // C++17
```



More examples

C++ 17

```
1 std::list<int> l = ...;
2
3 // Finds the first element in a list between 1 and 10.
4 const auto it = std::find_if(l.begin(), l.end(),
5     [] (int i) { return i >= 1 && i <= 10; });
6 if (it != l.end()) {
7     int element = *it; ...
8 }
9
10 // Computes sin(x)/(x + DBL_MIN) for elements of a range.
11 std::vector<double> r(l.size());
12 std::transform(l.begin(), l.end(), r.begin(),
13     [] (auto x) { return std::sin(x)/(x + DBL_MIN); });
14
15 // reduce/fold (using addition)
16 const auto sum = std::reduce(v.begin(), v.end());
```



More examples

C++ 11

```
1 std::vector<int> v = ...;
2
3 // remove duplicates
4 std::sort(v.begin(), v.end());
5 auto newEndIt = std::unique(v.begin(), v.end());
6 v.erase(newEndIt, v.end());
7
8 // remove by predicate
9 auto p = [](int i) { return i > 42; };
10 auto newEndIt = std::remove_if(v.begin(), v.end(), p);
11 v.erase(newEndIt, v.end());
12
13 // remove by predicate (C++20)
14 std::erase_if(v, p);
```



Welcome to lego programming!

C++ 98



Using the STL

C++ 98

Exercise: STL

- go to exercises/stl
- look at the non STL code in randomize.nostl.cpp
 - it creates a vector of ints at regular intervals
 - it randomizes them
 - it computes differences between consecutive ints
 - and the mean and variance of it
- open randomize.cpp and complete the “translation” to STL
- see how easy it is to reuse the code with complex numbers



Using the STL

C++ 98

Be brave and persistent!

- you may find the STL quite difficult to use
- template syntax is really tough
- it is hard to get right, compilers spit out long error novels
 - but, compilers are getting better with error messages
- C++ 20 will help with concepts and ranges
- the STL is extremely powerful and flexible
- it will be worth your time!



Standard library

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization

std::string_view cppreference

C++ 17 Adv

Non owning read-only view of a contiguous char sequence

- Doesn't allocate memory
- Similar interface to std::string, but read-only
- Easy to copy, faster for some calls eg. substr ($O(1)$ vs $O(n)$)
- The data pointed to has to outlive the string_view

Some example uses

```
1 constexpr std::string_view sv {"Some example"};
2 auto first = sv.substr(0, sv.find_first_of(" "));
3 std::string some_string {"foo bar"};
4 std::string_view sv_str(some_string);
5 char foo[3] = {'f', 'o', 'o'};
6 std::string_view carr(foo, std::size(foo));
```



std::span<T> - Concept cppreference

C++ 20 Adv

Non owning view of a contiguous sequence of items

- Doesn't allocate memory
- Allows to modify items, but not the container itself
 - no reallocation, push or pop
- Provides container like interface
 - iterators, size, front/back, subspan, ...
- Standard contiguous containers convert to it automatically
 - std::vector, std::array, C arrays, ...
 - not std::list or std::deque (no contiguous storage)
 - span should thus be preferred as function argument
- Simply a pair (pointer, size), so cheap to copy
 - and thus to be passed by value
- span can also have a static extent
 - meaning the size is determined at compile time
 - and thus only a pointer is stored by span



std::span<T> - Usage

Some example - godbolt

```
1 void print(std::span<const int> c) {
2     for (auto a : c) { std::cout << a << " "; }
3     std::cout << "\n";
4 }
5 void times2(std::span<int> c) {
6     for(auto &e : c) { e *= 2; }
7 }
8 int a[]{23, 45, 67, 89};
9 print(a);                      // 23 45 67 89
10
11 std::vector v{1, 2, 3, 4, 5};
12 std::span<const int> sv = v;
13 print(sv.subspan(2, 2)); // 3 4
14
15 std::array a2{-14, 55, 24};
16 times2(a2);
17 print(a2);                  // -28 110 48
18
19 std::span<const int, 3> sa2 = a2;
20 std::cout << sizeof(sv) << " " << sizeof(sa2) << "\n"; // 16 8
21 std::span<const int, 3> s2a2 = a; // compilation failure, invalid conversion
```



std::optional cppreference

C++ 17 Adv

Manages an optionally contained value

- Contextually converts to bool, telling if it contains something
- Has value semantics (Copy, Move, Compare, stack alloc.)
- Useful for the return value of a function that may fail
- Useful in place of pointers where value semantics are intuitive

```
1 std::optional<Phone> parse_phone(std::string_view in) {
2     if (is_valid_phone(in))
3         return in; // equiv. to optional<Phone>{in};
4     return {};   // or: return std::nullopt; (empty opt.)
5 }
6 auto v = parse_phone(...);
7 if (v) {                                // or: v.has_value()
8     process_phone(v.value()); // or: *v (unchecked)
9     v->call(); // calls Phone::call() (unchecked)
10 }
11 v.reset(); assert(!v.has_value()); // or: v = {};
```



Code example

```
1 // address may be given for a person or not
2 void processPerson(std::string_view name,
3                     std::optional<Address> address);
4 Address addr = wonderland();
5 processPerson("Alice", std::move(addr));
6
7 // Every person has a name, but not always an address
8 struct Person {
9     std::string name;
10    std::optional<Address> address;
11 };
12 std::vector<Person> ps = ...;
13 std::sort(ps.begin(), ps.end(),
14           [](const Person& a, const Person& b) {
15     return a.address < b.address;
16 }); // sorts by address, persons w/o address at front
```



Exercise: Handling optional results

- go to `exercises/optional/optional.cpp`
- modify `mysqrt` so to return an `std::optional<double>`
- guess the consequences on `square`



std::expected cppreference

C++ 23 Adv

Manages either of 2 values (expected or not)

- templated by the 2 value types
- Useful for the return value of a function that may fail
 - and would then return another type (error type)

Practically - godbolt

```
1 enum class Error {...};  
2 std::expected<int, Error> parse(std::string_view in) {  
3     if (is_valid(in)) return convert_to_int(in);  
4     return std::unexpected(Error{...});  
5 }  
6 auto v = parse(...);  
7 if (v) { // or v.has_value()  
8     std::cout << *v; // (unchecked)  
9     foo(v.value()); // may throw bad_expected_access  
10 } else  
11     log(v.error());
```



std::variant cppreference

C++ 17 Adv

A type-safe union

- Allows the variable to hold any of the given types
- std::get reads the value of the variant
- and throws std::bad_variant_access for bad accesses
- Currently held alternative can be probed

Code example

```
1 std::variant<int, float, string> opt{100}; // holding int
2 int ival = std::get<int>(opt); // or std::get<0>(opt)
3 try {
4     float val = std::get<float>(opt) // will throw
5 } catch (std::bad_variant_access const& ex) {...}
6 if (std::holds_alternative<float>(opt))
7     std::cout << std::get<float>(opt);
8 if (const float* v = std::get_if<float>(&opt)) { ... }
9 const int current_alternative = opt.index(); // == 0
```

std::variant and the visitor pattern

C++ 17 Adv

std::visit cppreference

- Applies a “visitor” to a given variant
 - A visitor is a callable able to handle all different types
 - Must return **void** or the same type for all overloads

Practically - godbolt

```
1 struct Visitor {  
2     auto operator()(int i)    {return "i:"+std::to_string(i);}  
3     auto operator()(float f) {return "f:"+std::to_string(f);}  
4     auto operator()(const std::string& s) { return "s:"+s;}  
5     template <typename T>  
6     auto operator()(const T& t) { return std::string{"?"}; }  
7 };  
8 void print(std::variant<int, float, std::string, char> v) {  
9     std::cout << std::visit(Visitor{}, v) << '\n';  
10 }  
11 print(100); print(42.0f); print("example"); print('A');
```



std::variant and the lambda visitor pattern

C++ 17 Adv

Idea

- use inheritance to group a set of lambdas

Practically - godbolt

```
1 template<typename... Ts> // covered in expert part
2 struct Overload : Ts ... { using Ts::operator()...; };
3 template<typename... Ts>
4 Overload(Ts...) -> Overload<Ts...>; // obsolete in C++20
5 int main(){
6     std::variant<int, float, std::string, char> v{ ... };
7     auto overloadSet = Overload {
8         [](int i) { std::cout << "i32:" << i; },
9         [](std::string s) { std::cout << s; },
10        [](auto a) { std::cout << "other"; },
11    };
12    std::visit(overloadSet, v);
13 }
```

std::variant exercise

C++ 17 Adv

Exercise: An alternative to oo polymorphism

- go to `exercises/variant/variant.cpp`
- replace inheritance from a common base class with an `std::variant` on the three kinds of particle.
- two solutions given : with `std::get_if` and with `std::visit`.



std::any cppreference

C++ 17 Adv

a type-safe container for single values of any type

- Allows a variable to hold any type (say bye to `void*`)
- `std::any_cast` reads the internal value
- and throws `std::bad_any_cast` for bad accesses
- `std::any_cast` only matches types 1:1, ignoring inheritance

Code example

```
1 std::any val{100};           // holding int
2 val = std::string("hello"); // holding string
3 std::string s = std::any_cast<std::string>(val);
4 try {
5     int val = std::any_cast<int>(val); // will throw
6 } catch (std::bad_any_cast const& ex) {...}
7 if (val.type() == typeid(int)) // requires RTTI
8     std::cout << std::any_cast<int>(val);
9 val.reset(); assert(!val.has_value());
```



non-member begin and end

C++ 11 Adv

The problem in C++ 98

STL containers and arrays have different syntax for loop

```
1 std::vector<int> v;
2 int a[] = {1,2,3};
3 for(auto it = v.begin(); it != v.end(); it++) {...}
4 for(int i = 0; i < 3; i++) {...}
```



non-member begin and end

C++ 11 Adv

The problem in C++ 98

STL containers and arrays have different syntax for loop

```
1 std::vector<int> v;
2 int a[] = {1,2,3};
3 for(auto it = v.begin(); it != v.end(); it++) {...}
4 for(int i = 0; i < 3; i++) {...}
```

A new syntax

```
5 for(auto it = begin(v); it != end(v); it++) {...}
6 for(auto i = begin(a); i != end(a); i++) {...}
```



non-member begin and end

C++ 11 Adv

The problem in C++ 98

STL containers and arrays have different syntax for loop

```
1 std::vector<int> v;
2 int a[] = {1,2,3};
3 for(auto it = v.begin(); it != v.end(); it++) {...}
4 for(int i = 0; i < 3; i++) {...}
```

A new syntax

```
5 for(auto it = begin(v); it != end(v); it++) {...}
6 for(auto i = begin(a); i != end(a); i++) {...}
```

Allowing the best syntax

```
7 for(auto & element : v) {...}
8 for(auto & element : a) {...}
```

std::tuple cppreference

C++ 11 Adv

```
1 #include <tuple>
2
3 std::tuple<int, float, bool> t{42, 3.14f, false};
4 auto t = std::make_tuple(42, 3.14f, false);
5 std::tuple t{42, 3.14f, false}; // C++17 CTAD
6
7 float& f = std::get<1>(t);    // get 1. member
8 bool& b = std::get<bool>(t); // get bool member
9
10 using Tup = decltype(t);
11 constexpr auto s = std::tuple_size_v<Tup>; // 3
12 using E2Type = std::tuple_element<2, Tup>; // bool
13
14 auto t2 = std::tuple_cat(t, std::tuple{'A'}, t);
15 // <int, float, bool, char, int, float, bool>
16 bool& b2 = std::get<bool>(t2); // error
```

std::tuple

C++ 11 Adv

```
1 void f(std::tuple<int, float, bool> tuple) {
2     int a; float b; bool c;
3     std::tie(a, b, c) // tuple<int&, float&, bool&>
4         = tuple;
5     // Use structured bindings in C++17
6 }
7
8 int& i = ...;
9 std::tuple{i};                      // tuple<int>
10 std::make_tuple(i);                 // tuple<int>
11 std::make_tuple(std::ref(i));       // tuple<int&>
12 std::tuple{std::ref(i)};           // C++17 CTAD, tuple<std::reference_wrapper<int>>
```

Structured Binding Declarations

C++ 17 Adv

- Helps with `std::tuple`, tuple-like, array or aggregate types
- Automatically creates variables and ties them
- CV qualifiers and references allowed

Practically - godbolt

```
1 std::tuple<int, double, long> tuple = ...;
2 auto [ a, b, c ] = tuple;
3
4 struct S { int i; float f; bool b; } s = ...;
5 auto&& [ a, b, c ] = std::move(s);
6
7 int arr[] = {1, 3, 4};
8 const auto& [ a, b, c ] = arr;
9
10 std::unordered_map<K, V> map = ...;
11 for (const auto& [key, value] : map) { ... }
```



Structured Binding Declarations

C++ 17 Adv

“Tuple-like binding protocol”

- For non-aggregates, you can opt-in to structured bindings
- The compiler requires these hooks:
 - Specialize `std::tuple_size`
 - Specialize `std::tuple_element`
 - Provide either `T::get<I>()` or `get<I>(T)`

```
1 class C { ... }; // complex, cannot change
2 template <>
3 struct std::tuple_size<C> {
4     static constexpr std::size_t value = ...;
5 };
6 template <std::size_t I>
7 struct std::tuple_element<I, C> { using type = ...; };
8 template <std::size_t I>
9 auto& get(C& c) { ... }; // -> std::tuple_element<I, C>&
```

C++ date and time utilities

C++ 11 / C++ 20

`std::chrono` library present in `<chrono>` header

Clocks

- consists of a starting point (or epoch) and a tick rate
 - E.g. January 1, 1970 and every second
- C++ defines several clock type
 - `system_clock` system time, aka wall clock time, or C time
 - `steady_clock` monotonic but unrelated to wall clock time
 - `high_resolution_clock` clock with the smallest tick period

`time_point` and `duration`

- provide easy manipulation of times and duration
- clock dependent
- `duration_cast` allows conversions between duration types
 - available helper types : nanoseconds, microseconds, milliseconds, seconds, minutes, hours, ...



Practical usage / timing some C++ code

C++ 11 Adv

How to measure the time spent in some code - godbolt

```
1 #include <chrono>
2
3 std::chrono::high_resolution_clock clock;
4 auto start = clock::now();
5 ... // code to be timed
6 std::chrono::duration<float> ticks = clock.now() - start;
7
8 auto millis = std::chrono::duration_cast
9                 <std::chrono::milliseconds>(ticks);
10 std::cout << "it took " << ticks.count() << " ticks"
11           << ", that is " << millis.count() << " ms\n";
```



Practical usage / timing some C++ code

C++ 11 Adv

How to measure the time spent in some code - godbolt

```
1 #include <chrono>
2
3 std::chrono::high_resolution_clock clock;
4 auto start = clock::now();
5 ... // code to be timed
6 std::chrono::duration<float> ticks = clock.now() - start;
7
8 auto millis = std::chrono::duration_cast
9                 <std::chrono::milliseconds>(ticks);
10 std::cout << "it took " << ticks.count() << " ticks"
11             << ", that is " << millis.count() << " ms\n";
```

Warning

- this does not measure the amount of CPU used !
- neither the time spent on a CPU (think suspended threads)



Pseudo-random number generators (PRNG)

C++ 11 Adv

Concept

- For generating pseudo-random numbers the STL offers:
 - engines and engine adaptors to generate pseudo random bits
 - distributions to shape these into numbers
 - access to (potential) hardware entropy
- All found in header `<random>`

Example

```
1 #include <random>
2 std::random_device rd;
3 std::default_random_engine engine{rd()};
4 std::normal_distribution dist(5.0, 1.5); // μ, σ
5 double r = dist(engine);
```



Engines

- Generate random bits (as integers), depending on algorithm
- Can be seeded via constructor or seed()
- Algorithms: `linear_congruential_engine`,
`mersenne_twister_engine`, `subtract_with_carry_engine`
- Adaptors: `discard_block_engine`, `independent_bits_engine`,
`shuffle_order_engine`
- Aliases: `minstd_rand0`, `minstd_rand`, `mt19937`, `mt19937_64`,
`ranlux24`, `ranlux48`, `knuth_b` (use these)
- `std::default_random_engine` aliases one of the above

Engine creation with seed

```
1 std::default_random_engine engine{42}; // or empty
2 auto i = engine(); // operator(): random integer
```



Determinism and random device

C++ 11 Adv

Engines / PRNGs

- Are *deterministic*
- Will produce the same number sequence for the same seed
- May be useful for reproducibility

std::random_device

- Provides *non-deterministic* random numbers
- Uses hardware provided entropy, if available
- May become slow when called too often,
e.g. when hardware entropy pool is exhausted
- Use only to seed engine, if *non-deterministic* numbers needed

Non-deterministic engine seeding

```
1 std::random_device rd;  
2 std::default_random_engine engine{rd()};
```



Distributions

C++ 11 Adv

Distributions

- Take random bits from engines and shape them into numbers
- Standard library provides a lot of them
 - Uniform, bernoulli, poisson, normal and sampling distributions
 - See [cppreference](#) for the full list

Distributions sharing a non-deterministic engine

```
1 std::random_device rd;
2 std::default_random_engine engine{rd()};
3 std::uniform_int_distribution<int> dist(2,7); //min,max
4 const int size = dist(engine); // gen. random number
5 std::vector<int> v(size);
6 std::normal_distribution<double> norm(5.0, 1.5); //μ,σ
7 std::generate(begin(v), end(v),
8                 [&]{ return norm(engine); });
```



C random library

C++ 98 Adv

C random library

- The C library (`<cstdlib>`) offers a primitive PRNG:
- `rand()` returns random int between 0 and `RAND_MAX` (incl.)
- `srand(seed)` sets the global seed

Disadvantages

- Seeding is not thread-safe; `rand()` may be.
- Returned numbers have small, implementation-defined range
- No guarantee on quality
- Mapping a random number to a custom range is hard
- E.g.: `rand() % 10` yields biased numbers and is wrong!

Good practice: Strongly avoid the C random library

Use the C++ 11 facilities from the `<random>` header



Ranges

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI and smart pointers
- Initialization



The range concept

- A range is a sequence of items that you can iterate over
- It must provide iterators to its beginning and end, obtainable via `std::range::begin` and `std::range::end`
- Range concepts are located in `std::ranges` namespace
- All STL containers are ranges

Variations on ranges

There are actually different types of ranges

- `input_range`, `output_range`, `forward_range`,
`bidirectional_range`, `random_access_range`,
`contiguous_range`
- corresponding to the different iterator categories



The view concept

- A view is a non-owning object built from a range or another view after applying some range adaptor(s)
- The time complexity to copy, move, assign is constant
- Range adaptors can be applied/chained with `|` and are lazy
- They live in the `std::views` namespace
- They introduce easy functional programming to the STL

Example - godbolt

```
1 auto const numbers = std::views::iota(0, 6);
2 auto even = [](int i) { return 0 == i % 2; };
3 auto square = [](int i) { return i * i; };
4 auto results = numbers | std::views::filter(even)
5                         | std::views::transform(square);
6 for (auto a : results) { ... } // 0, 4, 16
```



Many range adaptors are provided by the STL C++ 20 Adv

Useful subset

all a view on all elements

filter applies a filter

transform applies a transformation

take takes only first n elements

drop skips first n elements

join/split joins ranges or splits into subranges

reverse reverses view order

elements for tuple views, extract nth elements of each tuple

keys/values for pair like views, takes 1st/2nd element of each pair

Remember you can combine them via (|)



Back to laziness

C++ 23 Adv

Views are lazy

- Computation is only triggered when iterating
- And can be stopped by e.g. take
- Here, the minimal number of iterations is performed

Example - godbolt

```
1 // print first 20 prime numbers above 1000000
2 for (int i: std::views::iota(1000000)
3     | std::views::filter(odd)
4     | std::views::filter(isPrime)
5     | std::views::take(20)) {
6     std::cout << i << " ";
7 }
```



RAII and smart pointers

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAII and smart pointers
- Initialization



Pointers: why are they error prone?

C++ 98

They need initialization

```
1 char *s;
2 try {
3     foo(); // may throw
4     s = new char[100];
5     read_line(s);
6 } catch (...) { ... }
7 process_line(s);
```



Pointers: why are they error prone?

C++ 98

They need initialization

Seg Fault

```
1 char *s;
2 try {
3     foo(); // may throw
4     s = new char[100];
5     read_line(s);
6 } catch (...) { ... }
7 process_line(s);
```



Pointers: why are they error prone?

C++ 98

They need initialization

Seg Fault

```
1 char *s;
2 try {
3     foo(); // may throw
4     s = new char[100];
```

They need to be released

```
1 char *s = new char[100];
2 read_line(s);
3 if (s[0] == '#') return;
4 process_line(s);
5 delete[] s;
```



Pointers: why are they error prone?

C++ 98

They need initialization

Seg Fault

```
1 char *s;
2 try {
3     foo(); // may throw
4     s = new char[100];
```

They need to be released

Memory leak

```
1 char *s = new char[100];
2 read_line(s);
3 if (s[0] == '#') return;
4 process_line(s);
5 delete[] s;
```



Pointers: why are they error prone?

C++ 98

They need initialization

Seg Fault

```

1  char *s;
2  try {
3      foo(); // may throw
4      s = new char[100];

```

They need to be released

Memory leak

```

1  char *s = new char[100];
2  read_line(s);

```

They need clear ownership

```

1  char *s = new char[100];
2  read_line(s);
3  vec.push_back(s);
4  set.add(s);
5  std::thread t1{func1, vec};
6  std::thread t2{func2, set};

```



Pointers: why are they error prone?

C++ 98

They need initialization

Seg Fault

```

1  char *s;
2  try {
3      foo(); // may throw
4      s = new char[100];

```

They need to be released

Memory leak

```

1  char *s = new char[100];
2  read_line(s);

```

They need clear ownership

Who should release ?

```

1  char *s = new char[100];
2  read_line(s);
3  vec.push_back(s);
4  set.add(s);
5  std::thread t1{func1, vec};
6  std::thread t2{func2, set};

```



This problem exists for any resource

C++ 11

For example with a file

```
1 std::FILE *handle = std::fopen(path, "w+");
2 if (nullptr == handle) { throw ... }
3 std::vector v(100, 42);
4 write(handle, v);
5 if (std::fputs("end", handle) == EOF) {
6     return;
7 }
8 std::fclose(handle);
```

Which problems do you spot in the above snippet?



Resource Acquisition Is Initialization (RAII)

C++ 98

Practically

Use variable construction/destruction and scope semantics:

- wrap the resource inside a class
- acquire resource in constructor
- release resource in destructor
- create an instance on the stack
 - automatically destructed when leaving the scope
 - including in case of exception
- use move semantics to pass the resource around



RAII in practice

C++ 98

An RAII File class

```
1  class File {
2  public:
3      // constructor: acquire resource
4      File(const char* filename)
5          : m_handle(std::fopen(filename, "w+")) {
6          // abort constructor on error
7          if (m_handle == nullptr) { throw ...; }
8      }
9      // destructor: release resource
10     ~File() { std::fclose(m_handle); }
11     void write (const char* str) {
12         ...
13     }
14 private:
15     std::FILE* m_handle; // wrapped resource
16 };
```



Usage of File class

```
1 void log_function() {  
2     // file opening, aka resource acquisition  
3     File logfile("logfile.txt");  
4  
5     // file usage  
6     logfile.write("hello logfile!"); // may throw  
7  
8     // file is automatically closed by the call to  
9     // its destructor, even in case of exception!  
10 }
```

Good practice: Use `std::fstream` for file handling

The standard library provides `std::fstream` to handle files, use it!



std::unique_ptr

C++ 11

A RAII pointer

- wraps and behaves like a regular pointer
- get underlying pointer using `get()`
- when destroyed, deletes the object pointed to
- has move-only semantic
 - the pointer has unique ownership
 - copying will result in a compile error



std::unique_ptr

C++ 11

A RAII pointer

- wraps and behaves like a regular pointer
- get underlying pointer using `get()`
- when destroyed, deletes the object pointed to
- has move-only semantic
 - the pointer has unique ownership
 - copying will result in a compile error

```
1 #include <memory>
2 void f(std::unique_ptr<Foo> ptr) {
3     ptr->bar();
4 } // deallocation when f exits
5
6 std::unique_ptr<Foo> p{ new Foo{} }; // allocation
7 f(std::move(p)); // transfer ownership
8 assert(p.get() == nullptr);
```



Quiz

C++ 11

What do you expect?

```
1 void f(std::unique_ptr<Foo> ptr);  
2 std::unique_ptr<Foo> uptr(new Foo{});  
3 f(uptr); // transfer of ownership
```



Quiz

C++ 11

What do you expect?

```
1 void f(std::unique_ptr<Foo> ptr);  
2 std::unique_ptr<Foo> uptr(new Foo{});  
3 f(uptr); // transfer of ownership
```

Compilation Error - godbolt

```
test.cpp:15:5: error: call to deleted constructor  
of 'std::unique_ptr<Foo>'  
    f(uptr);  
           ^~~~  
/usr/include/c++/4.9/bits/unique_ptr.h:356:7: note:  
  'unique_ptr' has been explicitly marked deleted here  
  unique_ptr(const unique_ptr&) = delete;  
^
```



std::make_unique

C++ 14

std::make_unique

- allocates and constructs an object with arguments and wraps it with `std::unique_ptr` in one step
- no `new` or `delete` calls anymore!
- no memory leaks if used consistently



std::make_unique

C++ 14

std::make_unique

- allocates and constructs an object with arguments and wraps it with std::unique_ptr in one step
- no `new` or `delete` calls anymore!
- no memory leaks if used consistently

std::make_unique usage

```
1 {  
2     // calls new File("logfile.txt") internally  
3     auto f = std::make_unique<File>("logfile.txt");  
4     f->write("hello logfile!");  
5 } // deallocation at end of scope
```



Arrays and unique_ptr

C++ 11 Adv

Dynamic arrays

- `unique_ptr` can wrap arrays
- Use `T[]` as template parameter
- This will enable the subscript operator
- The default constructor is called for each element
- If size known at compile time, prefer `std::array`
- If size might change, prefer `std::vector`

```
1 auto b = std::make_unique<Foo[]>(10);
2 b[3] = ...;
3 b[4].someFunction();
4
5 // deallocations at end of scope
```



RAII or raw pointers

C++ 11

When to use what?

- Always use RAII for resources, in particular allocations
 - You thus never have to release / deallocate yourself
- Use raw pointers as non-owning, re-bindable observers
- Remember that `std::unique_ptr` is move only



RAII or raw pointers

C++ 11

When to use what?

- Always use RAII for resources, in particular allocations
 - You thus never have to release / deallocate yourself
- Use raw pointers as non-owning, re-bindable observers
- Remember that std::unique_ptr is move only

A question of ownership

```
1 std::unique_ptr<T> produce();
2 void observe(const T&);
3 void modifyRef(T&);
4 void modifyPtr(T*);
5 void consume(std::unique_ptr<T>);
6 std::unique_ptr<T> pt{produce()}; // Receive ownership
7 observe(*pt); // Keep ownership
8 modifyRef(*pt); // Keep ownership
9 modifyPtr(pt.get()); // Keep ownership
10 consume(std::move(pt)); // Transfer ownership
```



std::unique_ptr usage summary

C++ 11

Good practice: std::unique_ptr

- `std::unique_ptr` is about lifetime management
 - use it to tie the lifetime of an object to a unique RAII owner
 - use raw pointers/references to refer to another object without owning it or managing its lifetime
- use `std::make_unique` for creation
- strive for having no `new/delete` in your code
- for dynamic arrays, `std::vector` may be more useful



std::shared_ptr

C++ 11

`std::shared_ptr` : a reference-counting pointer

- wraps a regular pointer similar to `unique_ptr`
- has move and copy semantic
- uses reference counting internally
 - "Would the last person out, please turn off the lights?"
- reference counting is thread-safe, therefore a bit costly

`std::make_shared` : creates a `std::shared_ptr`

```
1 {  
2     auto sp = std::make_shared<Foo>(); // #ref = 1  
3     vector.push_back(sp);           // #ref = 2  
4     set.insert(sp);               // #ref = 3  
5 } // #ref 2
```



weak_ptr

weak_ptr: a non-owning observer

- Sometimes want to observe resources without keeping them alive
 - `shared_ptr`? Resource stays alive
 - Raw pointer? Risk of dangling pointer
- The solution is to construct a `weak_ptr` from a shared pointer
- To access the resource, convert the weak into a `shared_ptr`

```
1 std::shared_ptr<Cache> getSharedCache();
2 std::weak_ptr<Cache> weakPtr{ getSharedCache() };
3 // ... shared cache may be invalidated here
4 if (std::shared_ptr<Cache> cache = weakPtr.lock()) {
5     // Cache is alive, we actively extend its lifetime
6     return cache->findItem(...);
7 } else {
8     // Cache is nullptr, we need to do something
9     weakPtr = recomputeCache(...);
```

Quiz: std::shared_ptr in use

C++ 11

What is the output of this code? - godbolt

```
1 auto shared = std::make_shared<int>(100);
2 auto print = [shared](){
3     std::cout << "Use: " << shared.use_count() << " "
4                     << "value: " << *shared << "\n";
5 };
6 print();
7 {
8     auto ptr{ shared };
9     (*ptr)++;
10    print();
11 }
12 print();
```

Quiz: std::shared_ptr in use

C++ 11

What is the output of this code? - godbolt

```
1 auto shared = std::make_shared<int>(100);
2 auto print = [shared](){
3     std::cout << "Use: " << shared.use_count() << " "
4                     << "value: " << *shared << "\n";
5 };
6 print();
7 {
8     auto ptr{ shared };
9     (*ptr)++;
10    print();
11 }
12 print();
```

Use: 2 value: 100
Use: 3 value: 101
Use: 2 value: 101

Quiz: std::shared_ptr in use

C++ 11

What is the output of this code?

```
1 auto shared = std::make_shared<int>(100);
2 auto print = [&shared](){
3     std::cout << "Use: " << shared.use_count() << " "
4                     << "value: " << *shared << "\n";
5 };
6 print();
7 {
8     auto ptr{ shared };
9     (*ptr)++;
10    print();
11 }
12 print();
```

Use: 1 value: 100
Use: 2 value: 101
Use: 1 value: 101

Quiz: shared_ptr and weak_ptr in use

C++ 11 Adv

What is the output of this code? - godbolt

```
1 auto shared = std::make_shared<int>(100);
2 std::weak_ptr<int> weak{ shared };
3 print(); // with print as before
4
5 auto ptr = weak.lock();
6 (*ptr)++; print();
7
8 ptr = nullptr; print();
9
10 function(weak); print();
```



Quiz: shared_ptr and weak_ptr in use

C++ 11 Adv

What is the output of this code? - godbolt

```
1 auto shared = std::make_shared<int>(100);
2 std::weak_ptr<int> weak{ shared };
3 print(); // with print as before
4
5 auto ptr = weak.lock();
6 (*ptr)++; print();
7
8 ptr = nullptr; print();
9
10 function(weak); print();
```

Use: 1 value: 100

Use: 2 value: 101

Use: 1 value: 101

Use: 1 (or more) value: ???



Good practice: Single responsibility principle (SRP)

Every class should have only one responsibility.

Good practice: Rule of zero

- If your class has any special member functions (except ctor.)
 - Your class probably deals with a resource, use RAI
 - Your class should only deal with this resource (SRP)
 - Apply rule of 3/5: write/default/delete all special members
- Otherwise: do not declare any special members (rule of zero)
 - A constructor is fine, if you need some setup
 - If your class holds a resource as data member:
wrap it in a smart pointer, container, or any other RAI class



smart pointers

C++ 98

Exercise: Smart pointers

- go to [exercises/smartPointers](#)
- compile and run the program. It doesn't generate any output.
- Run with valgrind if possible to check for leaks

```
$ valgrind --leak-check=full --track-origins=yes ./smartPointers
```

- In the *essentials course*, go through `problem1()` and `problem2()` and fix the leaks using smart pointers.
- In the *advanced course*, go through `problem1()` to `problem4()` and fix the leaks using smart pointers.
- `problem4()` is the most difficult. Skip if not enough time.



Initialization

4 Core modern C++

- Constness
- Constant Expressions
- Exceptions
- Move semantics
- Copy elision
- Templates
- Lambdas
- The STL
- Standard library
- Ranges
- RAI! and smart pointers
- Initialization



Initialization

Initializing scalars

```

1 int i(42);      // direct initialization
2 int i{42};      // direct list initialization (C++11)
3 int i = 42;     // copy initialization
4 int i = {42};   // copy list initialization (C++11)

```

- All of the above have the same effect: $i == 42$

Narrowing conversions

```

1 int i(42.3);    // i == 42
2 int i{42.3};    // compilation error
3 int i = 42.3;   // i == 42
4 int i = {42.3}; // compilation error

```

- Braced initialization prevents narrowing conversions



Initialization

Initializing class types with constructors

```
1 struct C {  
2     C(int i);  
3     C(int i, int j);  
4 };  
5  
6 C c(42);      // calls C(42)  
7 C c{42};       // same  
8 C c = 42;      // calls C(42) in C++17, before C(C(42))  
9 C c = {42};    // same  
10  
11 C c(1, 2);    // calls C(1, 2)  
12 C c{1, 2};    // calls C(1, 2)  
13 C c = (1, 2); // calls C(2), comma operator  
14 C c = {1, 2}; // calls C(1, 2) in C++17, before C(C(1, 2))  
15  
16 C c(1.1, 2.2); // calls C(1, 2)  
17 C c{1.1, 2.2}; // compilation error
```



Initialization

std::initializer_list

```

1 struct C {
2     C(int i, int j);
3     C(std::initializer_list<int> l);
4 };
5
6 C c(1, 2);      // calls C(1, 2)
7 C c{1, 2};      // calls C(std::initializer_list<int>{1, 2})
8 C c = (1, 2);  // calls C(2), comma operator, error
9 C c = {1, 2};  // calls C(std::initializer_list<int>{1, 2})

```

- A constructor with a std::initializer_list parameter takes precedence over other overloads.

Beware

```

1 std::vector<int> v1(4, 5); // {5, 5, 5, 5}
2 std::vector<int> v2{4, 5}; // {4, 5}

```

Initialization

Type deduction

```
1 auto i(42);      // int
2 auto i{42};      // C++11: std::initializer_list<int>
                   // since C++17: int
3
4 auto i = 42;     // int
5 auto i = {42};   // std::initializer_list<int>
```



Initialization

C++ 11 Adv

Default initialization

```
T t;
```

- If T has a default constructor (including compiler generated), calls it
- If T is a fundamental type, t is left uninitialized (undefined value)
- If T is an array, the same rules are applied to each item



Initialization

C++ 11 Adv

Value initialization

```
1 T t{};      // value initialization
2 T t = {};   // same
3 f(T());    // passes value-initialized temporary
4 f(T{});    // same
5 T t();     // function declaration
```

- Basically zero-initializes `t` (or the temporary), unless `T` has a user defined constructor, which is run in this case
- Details on [cppreference](#)



Initialization

C++ 11 Adv

Aggregate initialization

```
1 struct S { int i; double d; };
2
3 S s{1, 2.3};           // s.i == 1, s.d == 2.3
4 S s = {1, 2.3};        // same
5 S s{.i = 1, .d = 2.3} // same, designated init. (C++20)
6 S s = {.i = 1, .d = 2.3} // same
7 S s{1};               // s.i == 1, s.d == 0.0
8 S s = {1};             // same
9 S s{.i = 1}            // same (C++20)
10 S s = {.i = 1}          // same (C++20)
11 S s{}; // value init, s.i == 0, s.d == 0.0
12 S s; // default init, undefined values
```

- An aggregate is a class with no constructors, only public base classes (C++ 17) and members, no virtual functions, no default member initializers (until C++ 14)
- Details on [cppreference](#)



Initialization

C++ 11 Adv

Good practice: Initialization

- In generic code, for a generic type T:
 - `T t;` performs the minimally necessary initialization
 - `T t{};` performs full and deterministic initialization
- Prefer `T t{};` over `T t; memset(&t, 0, sizeof(t));`
- Prefer braces over parentheses to avoid accidental narrowing conversions. E.g. `T t{1, 2};`
 - Unless T has a `std::initializer_list` constructor that you do not want to call!
 - Be wary of `std::vector`!
- The STL uses value initialisation when resizing containers
 - E.g. `vec.resize(vec.size() + 10);`
- Aggregates are very flexible. If your class does not need special initialization, make it an aggregate (rule of zero)



Initialization

Further resources

- The Nightmare of Initialization in C++ - Nicolai Josuttis - CppCon 2018
- Initialization in modern C++ - Timur Doumler - Meeting C++ 2018



Expert C++

1 History and goals

- Variadic templates
- Perfect forwarding
- SFINAE
- Concepts
- Modules
- Coroutines

2 Language basics

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

- The `<=>` operator

6 Useful tools

7 Concurrency

8 C++ and python

The <=> operator

5 Expert C++

- The <=> operator
- Variadic templates
- Perfect forwarding
- SFINAE
- Concepts
- Modules
- Coroutines



The burden of comparison operators

Motivation

- One often needs `operator<` for a user-defined class
 - e.g. when sorting a `std::vector`
 - e.g. when using it as a key for `std::set` or `std::map`
- Many other operators are also desirable for completeness
 - `operator>`, `operator>=`, `operator<=`, ...
 - often implemented reusing e.g. `operator<` and `operator==`
- Should be defined as hidden friend functions
- Much boilerplate code to write. Too much...



The three-way comparison operator

C++ 20 Adv

Idea

- C++20 introduces `operator<=>`
- named the **three-way comparison operator**.
 - unofficially called **spaceship operator**
- allowing to implement all comparisons in one go

How it works

- it returns *something* which can be compared to 0
 - similar to `std::strcmp`
- lower, greater, or equal to 0 means respectively *lower than*, *greater than* and *equivalent to*
- It is provided by default for all built-in types and many types in the standard library

The three-way comparison operator practically C++ 20 Adv

Output of operator<=> - godbolt

```
1 template <typename T>
2 void three_way_compare( T lhs, T rhs ) {
3     auto res = lhs <=> rhs;
4     std::cout << lhs << "<=>" << rhs << ":" "
5             << (res<0) << (res==0) << (res>0)
6             << '\n';
7 }
8 int main() {
9     three_way_compare(1, 2); // 1<=>2: 100
10    three_way_compare(2, 2); // 2<=>2: 010
11    three_way_compare(2, 1); // 2<=>1: 001
12 }
```

Different kinds of ordering

C++ 20 Adv

The return type of operator`<=>`

- for integers `operator<=>` returns `std::strong_ordering`
- `weak_ordering` and `partial_ordering` also exist

3 types of ordering

`strong` exactly one test among `<0`, `==0`, and `>0` will return `true`

`weak` like `strong` but two *equivalent* values may differ

- they are however *equivalent* for ranking
- e.g. rational numbers $2/3$ and $4/6$

`partial` like `weak` but some values are incomparable

- for some values all tests may return `false`
- e.g. compare a floating point to `Nan`



Exercising different orderings

C++ 20 Adv

Example - godbolt

```
1 struct Ratio {
2     unsigned n, d ;
3     friend std::weak_ordering operator<=>( Ratio const & a,
4                                                 Ratio const & b ) {
5         return (a.n * b.d) <=> (a.d * b.n);
6     }
7     friend std::ostream & operator<<( std::ostream & os, Ratio const & r ) {
8         return (os << r.n << '/' << r.d);
9     }
10 };
11 int main() {
12     // Ratio uses weak_ordering
13     three_way_compare(Ratio{1, 2}, Ratio{2, 3}); // 1/2<=>2/3 : 100
14     three_way_compare(Ratio{2, 3}, Ratio{4, 6}); // 2/3<=>4/6 : 010
15     three_way_compare(Ratio{2, 3}, Ratio{1, 2}); // 2/3<=>1/2 : 001
16
17     // floats use partial_ordering
18     three_way_compare(+0., -0.); // 0<=>-0 : 010
19     three_way_compare(0., 1./0.); // 0<=>inf : 100
20     three_way_compare(0., 0./0.); // 0<=>-nan : 000
21 }
```



Compiler-generated comparison operators

C++ 20 Adv

For a given user-defined class

- defining `operator<=` allows the compiler to use it when encountering the comparison operators `<`, `<=`, `>` and `>=`
- of course, one can still provide a custom implementation
- the compiler will *NOT* add a default implementation for `operator==` and `operator!=`
 - as those operators mean **equal**, rather than **equivalent**
 - if `operator<=` does not provide a strong order, it is advised not to define `operator==`



Compiler-generated operator<=>

C++ 20 Adv

Default operator<=> implementation

- One can ask the compiler to provide a default implementation for `operator<=>` and/or `operator==`
 - declaring them with “`= default`”
 - it will compare the member variables, one by one
- Can be useful e.g. for tuples
- Can be wrong, e.g. for the Ratio class
- If `operator<=>` is defaulted and no `operator==` is defined, then the compiler also provides `operator==`



Summary

- Defining `operator<=>` allows you to use `operator<`, `operator>`, `operator<=`, and `operator>=` for free
- The standard library defines a few kinds of orderings
 - strong, weak and partial
- If `operator<=>` does not define a strong ordering, avoid defining `operator==`



Variadic templates

5 Expert C++

- The `<=>` operator
- **Variadic templates**
- Perfect forwarding
- SFINAE
- Concepts
- Modules
- Coroutines



Basic variadic template

C++ 11 Adv

The idea

- a parameter accepting arbitrarily many arguments (i.e. pack)
- template parameter pack for e.g. types, function parameter packs for values, and expansions, details on [cppreference](#)

Recursive example `cppinsight`

```
1 template<typename T>
2 T sum(T v) { return v; }
3
4 template<typename T,
5           typename... Args>      // temp. param. pack
6 T sum(T first, Args... args) { // func. param. pack
7     return first + sum(args...); // pack expansion
8 }
9 int s = sum(1, 2, 3, 8, 7);
```

A couple of remarks

C++ 11 Adv

About performance

- do not be afraid of recursion
- everything is at compile time!
- unlike C-style variadic functions
 - e.g. `printf(const char* fmt, ...);`

Why it is better than variadic functions

- it's more performant
- type safety is included
- it applies to everything, including user-defined types



Parameter packs

C++ 11 Adv

Parameter packs

- only a few operations are supported
 - query the number of elements in a pack with `sizeof...`
 - expand it with `x...`
- can be empty
- can hold different types

```
1 template<typename... Args>
2 void f(Args... args) {
3     constexpr std::size_t N = sizeof...(Args);
4     std::tuple<Args...> t{args...};
5 }
6 f(0, 1, 2);           // Args = [int, int, int]
7 f();                 // Args = empty
8 f(0, true, 3.14);   // Args = [int, bool, double]
```



Fold expressions

C++ 17 Adv

The idea

- reduces a parameter pack over a binary operator
- details on [cppreference](#)

Example

```
1 template<typename... Args>
2 auto sum1(Args... args) {
3     return (args + ...);      // unary fold over +
4 }                           // parentheses mandatory
5 template<typename... Args>
6 auto sum2(Args... args) {
7     return (args + ... + 0); // binary fold over +
8 }                           // parentheses mandatory
9 int sum = sum1(); // error
10 int sum = sum2(); // ok
```

Fold expressions

C++ 17 Adv

Example: call a function on all arguments - [cppinsight](#)

```
1 template<typename T>
2 void print(const T& t) {
3     std::cout << t;
4 }
5 template<typename... Ts>
6 void printAll(const Ts&... ts) {
7     (print(ts),...); // fold over comma operator
8 }
```

C++ 11 workaround (don't use anymore)

```
1 template<typename... Ts>
2 void printAll(const Ts&... ts) {
3     int dummy[] {(print(ts),0)...};
4 }
```

Variadic class template

C++ 11 Adv

The tuple example, simplified - godbolt

```
1 template <typename... Ts>
2 struct tuple {};
3
4 template <typename T, typename... Ts>
5 struct tuple<T, Ts...> : tuple<Ts...> {
6     tuple(T head, Ts... tail) :
7         tuple<Ts...>(tail...), m_head(head) {}
8     T m_head;
9 }
10
11 tuple<double, uint64_t, const char*>
12     t1(12.2, 42, "big");
```



std::integer_sequence

C++ 14 Adv

Packs of values

- We can also have non-type template parameter packs
- Useful to pass lists of compile-time constants around
- The standard library includes a few helpers

```
1 template<typename T, T... Is>
2 struct integer_sequence { ... };
3
4 template<size_t... Is>
5 using index_sequence =
6     integer_sequence<std::size_t, Is...>;
7
8 template<size_t N>
9 using make_index_sequence =
10    index_sequence</*0, 1, ..., N-1*/>;
```

std::integer_sequence

C++ 14 Adv

Example - make_from_tuple with helper

```
1 template<typename T,
2         typename... Args, std::size_t... Is>
3 T helper(std::tuple<Args...> args,
4          std::index_sequence<Is...>) {
5     return T(std::get<Is>(args)...);
6 }
7 template<typename T, typename... Args>
8 T make_from_tuple(std::tuple<Args...> args) {
9     return helper<T>(args,
10                      std::make_index_sequence<sizeof...(Args)>{});
11 }
12
13 struct S { S(int, float, bool) { ... } };
14 std::tuple t{42, 3.14, false};
15 S s = make_from_tuple<S>(t);
```



std::integer_sequence

C++ 20 Adv

Example - make_from_tuple with lambda

```
1 template<typename T, typename... Args>
2 T make_from_tuple(std::tuple<Args...> args) {
3     return [&]<std::size_t... Is>(
4         std::index_sequence<Is...>){
5             return T(std::get<Is>(args)...);
6         }(std::make_index_sequence<sizeof...(Args)>{});
7     }
8
9     struct S { S(int, float, bool) { ... } };
10    std::tuple t{42, 3.14, false};
11    S s = make_from_tuple<S>(t);
```



Variadic templates

C++ 11 Adv

Exercise: Variadic templates

- go to exercises/variadic
- you will extend the tuple from the last slides
- follow the instructions in the source code



Perfect forwarding

5 Expert C++

- The `<=>` operator
- Variadic templates
- **Perfect forwarding**
- SFINAE
- Concepts
- Modules
- Coroutines



The problem

C++ 11 Adv

How to write a generic wrapper function?

```
1 template <typename T>
2 void wrapper(T arg) {
3     // code before
4     func(arg);
5     // code after
6 }
```

Example usage :

- emplace_back
- make_unique



Why is it not so simple?

C++ 11 Adv

```
1 template <typename T>
2 void wrapper(T arg) {
3     func(arg);
4 }
```

What about references?

- what if func takes a reference to avoid copies?
- wrapper would force a copy and we fail to use references



Second try, second failure ?

C++ 11 Adv

```
1 template <typename T>
2 void wrapper(T& arg) {
3     func(arg);
4 }
5 wrapper(42);
// invalid initialization of
// non-const reference from
// an rvalue
```

and **const** T& won't work when passing something non const



Solution: cover all cases

C++ 11 Adv

```
1 template <typename T>
2 void wrapper(T& arg) { func(arg); }
3
4 template <typename T>
5 void wrapper(const T& arg) { func(arg); }
6
7 template <typename T>
8 void wrapper(T&& arg) { func(std::move(arg)); }
```



The new problem: scaling to more arguments C++ 11 Adv

```
1 template <typename T1, typename T2>
2 void wrapper(T1& arg1, T2& arg2)
3 { func(arg1, arg2); }

4

5 template <typename T1, typename T2>
6 void wrapper(const T1& arg1, T2& arg2)
7 { func(arg1, arg2); }

8

9 template <typename T1, typename T2>
10 void wrapper(T1& arg1, const T2& arg2)
11 { func(arg1, arg2); }

12 ...
```

Exploding complexity

- for n arguments, 3^n overloads
- you do not want to try $n = 5\dots$

Reference collapsing

C++ 11 Adv

Reference to references

- Are formally forbidden, but can occur sometimes

```
1 template <typename T>
2 void foo(T t) { T& k = t; } //int& &, error in C++98
3 int ii = 4;
4 foo<int&>(ii); // want to pass by reference
```

C++ 11 added rvalue-references

- More combinations possible, like: `int&&` &, or `int&&` &&

Reference collapsing

- Rule: When multiple references are involved, & always wins
- $T\&\& \&$, $T\& \&\&$, $T\& \&$ → $T\&$
- $T\&\& \&\&$ → $T\&\&$



Forwarding references

C++ 11 Adv

Rvalue reference in type-deducing context

```
1 template <typename T>
2 void f(T&& t) { ... }
```

Forwarding reference

- Next to a template parameter, that can be deduced, `&&` is not an rvalue reference, but a “forwarding reference”
 - aka. “universal reference”
 - So, this applies only to functions
- The template parameter is additionally allowed to be deduced as an lvalue reference, and reference collapsing proceeds:
 - if an lvalue of type U is given, T is deduced as `U&` and the parameter type `U&&` collapses to `U&`
 - otherwise (rvalue), T is deduced as U
 - and forms the parameter type `U&&`



Forwarding references

C++ 11 Adv

Examples

```
1 template <typename T>
2 void f(T&& t) { ... }
3
4 f(4);           // rvalue -> T is int&
5 double d = 3.14;
6 f(d);           // lvalue -> T is double&
7 float g() {...}
8 f(g());         // rvalue -> T is float&
9 std::string s = "hello";
10 f(s);          // lvalue -> T is std::string&
11 f(std::move(s)); // rvalue -> T is std::string&
12 f(std::string{"hello"}); // rvalue -> std::string&
```



Forwarding references

C++ 20 Adv

Careful!

```
1 template <typename T> struct S {  
2     S(T&& t) { ... }           // rvalue references  
3     void f(T&& t) { ... }    // NOT forwarding references  
4};
```

Half-way correct version

```
1 template <typename T> struct S {  
2  
3     template <typename U>  
4         S(U&& t) { ... } // deducing context -> fwd. ref.  
5     template <typename U>  
6         void f(U&& t)    // deducing context -> fwd. ref.  
7         // ... but now U can be a different type than T  
8};
```

Forwarding references

Careful!

```
1 template <typename T> struct S {  
2     S(T&& t) { ... }           // rvalue references  
3     void f(T&& t) { ... }    // NOT forwarding references  
4};
```

Correct version

```
1 template <typename T> struct S {  
2     template <typename U, std::enable_if_t<  
3         std::is_same_v<std::decay_t<U>, T>, int> = 0>  
4         S(U&& t) { ... } // deducing context -> fwd. ref.  
5     template <typename U>  
6         void f(U&& t)    // deducing context -> fwd. ref.  
7         requires std::same_as<std::decay_t<U>, T> { ... }  
8};
```

Perfect forwarding

C++ 11 Adv

std::remove_reference

- Type trait to remove reference from a type
- If T is a reference type, `remove_reference_t<T>` is the type referred to by T, otherwise it is T.

```
1 template <typename T>
2 struct remove_reference { using type = T; };
3 template <typename T>
4 struct remove_reference<T&> { using type = T; };
5 template <typename T>
6 struct remove_reference<T&&> { using type = T; };
7
8 template <typename T>
9 using remove_reference_t =
10   typename remove_reference<T>::type; // C++14
```



Perfect forwarding

C++ 11 Adv

std::forward

- Forward lvalue-/rvalueness
- Keeps references and maps non-references to rvalue references

```
1 template<typename T>
2 T&& forward(remove_reference_t<T>& t) noexcept { // 1.
3     return static_cast<T&&>(t);
4 }
5 template<typename T>
6 T&& forward(remove_reference_t<T>&& t) noexcept { // 2.
7     return static_cast<T&&>(t);
8 }
```

- if T is `int`, selects 2., returns `int&&`
- if T is `int&`, selects 1., returns `int& &&`, i.e. `int&`
- if T is `int&&`, selects 2., returns `int&& &&`, i.e. `int&&`



Perfect forwarding

C++ 11 Adv

Example - putting it all together

```
1 template <typename... T>
2 void wrapper(T&&... args) {
3     func(std::forward<T>(args)...);
4 }
```

- if we pass an rvalue reference $U\&&$ to wrapper
 - $T=U$, arg is of type $U\&&$
 - func will be called with a $U\&&$
- if we pass an lvalue reference $U\&$ to wrapper
 - $T=U\&$, arg is of type $U\&$ (reference collapsing)
 - func will be called with a $U\&$
- if we pass a plain U (rvalue) to wrapper
 - $T=U$, arg is of type $U\&&$ (no copy in wrapper)
 - func will be called with a $U\&&$
 - if func takes a plain U , copy happens there, as expected



Real life example

C++ 11 Adv

```
1 template<typename T, typename... Args>
2 unique_ptr<T> make_unique(Args&&... args) {
3     return unique_ptr<T>
4         (new T(std::forward<Args>(args)...));
5 }
```



SFINAE

5 Expert C++

- The `<=>` operator
- Variadic templates
- Perfect forwarding
- **SFINAE**
- Concepts
- Modules
- Coroutines



Substitution Failure Is Not An Error (SFINAE) C++ 11 Adv

The main idea

- substitution replaces template parameters with the provided arguments (types or values)
- if it leads to invalid code, do not report an error but try other overloads/specializations

Example

```
1 template <typename T>
2 void f(typename T::type arg) { ... }
3 void f(int a) { ... }
4
5 f(1); // Calls void f(int)
```

Note: SFNAE is largely superseded by concepts in C++ 20



decltype

C++ 11 Adv

The main idea

- gives the type of the result of an expression
- the expression is not evaluated (i.e. unevaluated context)
- at compile time

Example

```
1 struct A { double x; };
2 A a;
3 decltype(a.x) y;           // double
4 decltype((a.x)) z = y;    // double& (lvalue)
5 decltype(1 + 2u) i = 4;   // unsigned int
6
7 template<typename T, typename U>
8 auto add(T t, U u) -> decltype(t + u);
9 // return type depends on template parameters
```

The main idea

- function to create reference to a “fake” object at compile time
- useful for types that cannot easily be constructed
- use only in unevaluated contexts, e.g. inside `decltype`

```
1 struct Default {
2     int foo() const;
3 };
4 class NonDefault {
5     private: NonDefault();
6     public: int foo() const;
7 };
8 decltype(Default().foo()) n1 = 1;      // int
9 decltype(NonDefault().foo()) n2 = 2; // error
10 decltype(std::declval<NonDefault>().foo()) n3 = 3;
```



true_type and false_type

C++ 11 Adv

The main idea

- encapsulate a compile-time boolean as type
- can be inherited

Example

```
1 struct truth : std::true_type { };
2
3 constexpr bool test = truth::value; // true
4 constexpr truth t;
5 constexpr bool test = t(); // true
6 constexpr bool test = t; // true
```

Possible implementation

```
1 using true_type = integral_constant<bool, true>;
2 using false_type = integral_constant<bool, false>;
```

Using SFINAE for introspection

C++ 11 Adv

The main idea

- use a primary template, inheriting from `false_type`
- use a template specialization, inheriting from `true_type`
 - which depends on the feature you want to introspect, as part of the context where template argument deduction happens
- let SFINAE choose between the two templates

Example

```
1 template <typename T, typename = void>
2 struct hasFoo : std::false_type {};
3 // primary template
4 template <typename T>
5 struct hasFoo<T, decltype(std::declval<T>().foo())>
6     : std::true_type {};
7 // template special.
8 struct A{};
9 struct B{ void foo(); };
10 static_assert(!hasFoo<A>::value, "A has no foo()");
11 static_assert(hasFoo<B>::value, "B has foo()");
```

Not so easy actually...

C++ 11 Adv

Example - godbolt

```
1 template <typename T, typename = void>
2 struct hasFoo : std::false_type {};
3 template <typename T>
4 struct hasFoo<T, decltype(std::declval<T>().foo())>
5   : std::true_type {};
6
7 struct A{};
8 struct B{ void foo(); };
9 struct C{ int foo(); };
10
11 static_assert(!hasFoo<A>::value, "A has no foo()");
12 static_assert(hasFoo<B>::value, "B has foo()");
13 static_assert(!hasFoo<C>::value, "C has no foo()");
14 static_assert(hasFoo<C,int>::value, "C has foo()");
```



Using void_t

C++ 17 Adv

Concept

- Maps a sequence of given types to `void`
- Introduced in C++ 17, though trivial to implement in C++ 11
- Can be used in specializations to check the validity of an expression

Implementation in header type_traits

```
1 template <typename...>
2 using void_t = void;
```



Previous introspection example using void_t C++ 17 Adv

Example - godbolt

```
1 template <typename T, typename = void>
2 struct hasFoo : std::false_type {};
3
4 template <typename T>
5 struct hasFoo<T,
6     std::void_t<decltype(std::declval<T>().foo())>>
7 : std::true_type {};
8
9 struct A{}; struct B{ void foo(); };
10 struct C{ int foo(); };
11
12 static_assert(!hasFoo<A>::value, "unexpected foo()");
13 static_assert(hasFoo<B>::value, "expected foo()");
14 static_assert(hasFoo<C>::value, "expected foo()");
```



Standard type traits

C++ 11 Adv

Standard type traits

- Don't implement SFINAE introspection for everything, use the standard library!
- Standard type traits in header `<type_traits>`
- They look like `std::is_*<T>`
- Checks at compile time whether T is
 - `float`, `signed`, `final`, `abstract`, `default_constructible`, ...
- Result in nested boolean constant value
- Since C++ 17, variable template versions: `std::is_*_v<T>`, giving the bool directly

```
1 constexpr bool b = std::is_pointer<int*>::value;
2 constexpr bool b = std::is_pointer_v<int*>; // C++17
```

SFINAE and the STL

C++ 11/C++ 14

```
enable_if cppreference
```

```
template<bool B, typename T=void>
struct enable_if;
```

- If B is true, has a nested alias type to type T
- otherwise, has no such alias

Example

```
1 template<bool B, typename T=void>
2 struct enable_if {};
3 template<typename T>
4 struct enable_if<true, T> { using type = T; };
5
6 template<bool B, typename T = void> // C++14
7 using enable_if_t = typename enable_if<B, T>::type;
```



Usage example - godbolt

```
1 template <typename T> constexpr bool is_iterable ... ;  
2  
3 template<typename T,  
4           typename std::enable_if_t  
5           <!is_iterable<T>, bool> = true>  
6 void print(T const& t) {  
7     std::cout << t << "\n";  
8 }  
9 template<typename T,  
10           typename std::enable_if_t  
11           <is_iterable<T>, bool> = true>  
12 void print(T const& t) {  
13     for (auto const& item : t) {  
14         std::cout << item << "\n";  
15     }  
16 }
```

Note : using **if constexpr** would be simpler



Back to variadic class templates

C++ 11 Adv

The tuple_element trait

```
1 template <size_t I, typename Tuple>
2 struct tuple_element;
3
4 template <typename T, typename... Ts>
5 struct tuple_element<0, tuple<T, Ts...>> {
6     using type = T;
7 };
8
9 template <size_t I, typename T, typename... Ts>
10 struct tuple_element<I, tuple<T, Ts...>> {
11     using type = typename
12         tuple_element<I - 1, tuple<Ts...>>::type;
13 };
```



Back to variadic class templates

C++ 14 Adv

The tuple get function

```
1 template <size_t I, typename... Ts>
2 std::enable_if_t<I == 0,
3     typename tuple_element<0, tuple<Ts...>>::type&>
4 get(tuple<Ts...>& t) {
5     return t.m_head;
6 }
7 template <size_t I, typename T, typename... Ts>
8 std::enable_if_t<I != 0,
9     typename tuple_element<I - 1, tuple<Ts...>>::type&>
10 get(tuple<T, Ts...>& t) {
11     return get<I - 1>(static_cast<tuple<Ts...>&>(t));
12 }
```



with if constexpr and return type deduction

C++ 17 Adv

The tuple get function

```
1 template <size_t I, typename T, typename... Ts>
2 auto& get(tuple<T, Ts...>& t) {
3     if constexpr(I == 0)
4         return t.m_head;
5     else
6         return get<I - 1>(static_cast<tuple<Ts...>&>(t));
7 }
```

Good practice: SFNAE vs. if constexpr

- `if constexpr` can replace SFNAE in many places.
- It is usually more readable as well. Use it if you can.



Concepts

5

Expert C++

- The `<=>` operator
- Variadic templates
- Perfect forwarding
- SFINAE
- Concepts
- Modules
- Coroutines



Requirements and concepts

C++ 20 Adv

Motivation

- Generic programming is made of variable, function and class templates which can be instantiated with different types.
- It is frequent to instantiate them with **unsuited types**, and the resulting compilation errors are cryptic.
- As a last resort, authors provide **documentation**, and practice tricky **template meta-programming**.
- C++20 brings **simpler ways to define constraints** on template parameters.



The world before concepts

C++ 17 Adv

C++ 17 work around: SFINAE

- Unsuited arguments can be avoided by inserting fake template arguments, leading to a substitution failure

Practical code - godbolt

```
1 template
2 <typename T,
3   typename = std::enable_if_t<std::is_floating_point_v<T>>>
4 bool equal( T e1, T e2 ) {
5   return std::abs(e1-e2) < std::numeric_limits<T>::epsilon();
6 }
7 ... equal(10,5+5) ...
```



The world before concepts

C++ 17 Adv

C++ 17 work around: SFINAE

- Unsuited arguments can be avoided by inserting fake template arguments, leading to a substitution failure

Practical code - godbolt

```
1 template
2 <typename T,
3   typename = std::enable_if_t<std::is_floating_point_v<T>>>
4 bool equal( T e1, T e2 ) {
5   return std::abs(e1-e2) < std::numeric_limits<T>::epsilon();
6 }
7 ... equal(10,5+5) ...
<source>:11:12: error: no matching function for call to 'equal(int, int)'
11   if (equal(10,5+5)) { std::cout << "FAILURE\n"; }
<source>:7:6: note: candidate: 'template<class T, class> bool equal(T, T)'
    bool equal( T e1, T e2 )
<source>:7:6: note:  template argument deduction/substitution failed:
In file included from <source>:1:
.../type_traits: In substitution of 'template<bool _Cond, class _Tp>
using enable_if_t = typename std::enable_if::type with bool _Cond = false; _Tp = void':
<source>:6:14:  required from here
.../type_traits:2514:11: error: no type named 'type' in 'struct std::enable_if<false, void>'
2514     using enable_if_t = typename enable_if<_Cond, _Tp>::type;
```



Basic requirements

C++ 20 Adv

A new keyword

- The keyword **requires** lets us define various constraints.

With concepts - godbolt

```
1 template<typename T>
2     requires std::is_floating_point_v<T>
3     bool equal( T e1, T e2 ) {
4         return std::abs(e1-e2) < std::numeric_limits<T>::epsilon();
5     }
6     ... equal(10,5+5) ...
```



Basic requirements

C++ 20 Adv

A new keyword

- The keyword **requires** lets us define various constraints.

With concepts - godbolt

```
1 template<typename T>
2     requires std::is_floating_point_v<T>
3     bool equal( T e1, T e2 ) {
4         return std::abs(e1-e2) < std::numeric_limits<T>::epsilon();
5     }
6     ... equal(10,5+5) ...
7
8 <source>:11:12: error: no matching function for call to 'equal(int, int)'
9     11     if (equal(10,5+5)) { std::cout << "FAILURE\n"; }
10    <source>:7:6: note: candidate: 'template<class T, class> bool equal(T, T)'
11    7     bool equal( T e1, T e2 )
12    <source>:7:6: note:   template argument deduction/substitution failed:
13    <source>:7:6: note: constraints not satisfied
14    <source>: In substitution of 'template<class T> ... bool equal(T, T) with T = int':
15    <source>:11:12:   required from here
16    <source>:7:6:   required by the constraints of 'template<class T> ... bool equal(T, T)'
17    <source>:6:15: note: the expression 'is_floating_point_v<T>' with T = int' evaluated to 'false'
18    6     requires std::is_floating_point_v<T>
```



Requirements and overloads

C++ 20 Adv

Example of several competing templates

```
1 template<typename T>
2 bool equal( T e1, T e2 ) { return (e1==e2); }
3
4 template< typename T>
5 requires std::is_floating_point_v<T>
6 bool equal( T e1, T e2 )
7 {return std::abs(e1-e2)<std::numeric_limits<T>::epsilon();}
```

Requirements affect overload resolution

- Overload resolution considers the second function as a better match when the requirements are fulfilled.



Concepts

C++ 20 Adv

Definition

- a **concept** gives a name to a given set of requirements
- useful when the same requirements are reused frequently

A new keyword: concept

```
1 template< typename T>
2 concept MyFloatingPoint =
3     std::is_floating_point_v<T> &&
4     std::numeric_limits<T>::epsilon() > 0;
5
6 template<typename T>
7 requires MyFloatingPoint<T>
8 bool equal( T e1, T e2 )
9 {return std::abs(e1-e2)<std::numeric_limits<T>::epsilon();}
```



Some usages of concepts

C++ 20 Adv

Constrained template parameters

- a concept can replace **typename** in a template parameter list

```
1 template<MyFloatingPoint T>
2 bool equal( T e1, T e2 ) {
3     return std::abs(e1-e2)<std::numeric_limits<T>::epsilon();
4 }
```

Constrained variables

- a concept can be used together with **auto** to impose requirements on the type of a variable

Example code

```
4 MyFloatingPoint auto f = 3.14f;
5 MyFloatingPoint auto d = 3.14;
6 MyFloatingPoint auto i = 3; // compile error
```



Some usages of concepts

C++ 20 Adv

Abbreviated function templates

- function parameters can be constrained as well
 - similar to variables
- the template head becomes obsolete
- the function is still a template though!
- unconstrained `auto` parameters are allowed as well

```
1 // no template <...> here!
2 bool equal( MyFloatingPoint auto e1,
3             MyFloatingPoint auto e2 ) {
4     return std::abs(e1-e2) <
5             std::numeric_limits<decltype(e1)>::epsilon();
6 }
7
8 // unconstrained abbreviated function template:
9 void equal(auto e1, auto e2) { ... }
```



Standard concepts

C++ 20 Adv

Never reinvent the wheel

- Writing a bug-proof concept is an expert task
- Prefer the ones provided by the standard library

E.g.: the floating point concept

```
1 #include <concepts>
2 bool equal( std::floating_point auto e1,
3             std::floating_point auto e2 ) {
4     return std::abs(e1-e2) <
5             std::numeric_limits<decltype(e1)>::epsilon();
6 }
```



Checking concepts

C++ 20 Adv

Concepts as Boolean operators

- Concepts can be used wherever a `bool` is expected
- they can appear in `if constexpr` conditions
 - since they are evaluated at compile-time

Using concepts with `if constexpr`

```
1 template<typename T>
2     bool equal( T e1, T e2 )
3 {
4     if constexpr (std::floating_point<T>) {
5         return std::abs(e1-e2)
6             < std::numeric_limits<T>::epsilon();
7     } else {
8         return (e1==e2);
9     }
10 }
```



Advanced requirements overview

C++ 20 Adv

requires as an expression

requires can express more than basic requirements. It can

- include other basic requirements
- list expressions that must be valid
- check the return type of some expressions

Practically

```
1 template<typename T>
2 concept StreamableAndComparableNumber =
3     requires( T v1, T v2, std::ostream os ) {
4         requires std::integral<T> || std::floating_point<T>;
5         os << v1 << v2;
6         { equal(v1,v2) } -> std::convertible_to<bool>;
7     };
```

Remember: use standard concepts first



Requirements and concepts

C++ 20 Adv

Summary

- A template can now *require* properties of its parameters
- Compiler error messages clearly state which argument does not fulfill which requirement
- A set of requirements can be gathered in a *concept*
- Overload resolution takes requirements into account
- The standard library provides many ready-to-use concepts
- Writing a new good concept is an expert task



Exercise time

C++ 20 Adv

Exercise: Concepts

- go to exercises/concepts
- you will use concepts to optimize some loop on iterators
- follow the instructions in the source code



Modules

5 Expert C++

- The `<=>` operator
- Variadic templates
- Perfect forwarding
- SFINAE
- Concepts
- **Modules**
- Coroutines



Modules

C++ 20 Adv

Motivation

- Modules allow for sharing declarations and definitions across translation units
- Header files do the same, but modules have benefits:
 - Better isolation. No cross-talk between multiple pieces of included code, or between included code and your code.
 - Better control over public interface of your library. Avoid leaking library dependencies into user code.
 - Faster compilation (`import fmt; >1500x faster`)

Textual includes

```
1 #include "math.h"
```

Module importation

```
1 import math;
```

Writing modules

C++ 20 Adv

ratio.cpp

```
1 export module math;
2 import <string>;
3 namespace utils { // ns. and module names orthogonal
4     export struct Ratio { float num, denom; };
5     export std::string toString(const Ratio& r) { ... }
6     void normalize(Ratio& r) { ... } // not exported
7 }
```

main.cpp

```
1 import <iostream>;
2 import math;
3 int main() {
4     auto r = utils::Ratio{1, 3};
5     std::cout << utils::toString(r);
6 }
```

Writing modules

C++ 20 Adv

Exporting declarations

- By default, everything inside a module is internal
- Declarations become visible when declared with `export` in the primary module interface
- You can only export entities at namespace scope (including the global scope)
- You cannot export entities with internal linkage (declared static or inside anonymous namespace)
- You cannot export aliases to non-exported types



Writing modules

C++ 20 Adv

What to export

```
1  export module math;  
2  
3  export void f();           // function  
4  export float pi = 3.14f;   // variable  
5  export namespace utils {   // namespace  
6      enum Status { Yes, No }; // enum (exported)  
7  }  
8  namespace utils {  
9      export struct Ratio { ... }; // class  
10     export using R = Ratio;    // alias  
11     float e = 2.71f;          // (not exported)  
12  }  
13  export { // content of export block  
14      void h();  
15  }
```

Writing modules

C++ 20 Adv

What not to export

```
1  export module math;  
2  
3  namespace {  
4      export void f();           // error  
5      export float pi = 3.14f;    // error  
6      export struct Ratio { ... }; // error  
7  }  
8  
9  export static void f();           // error  
10 export static float pi = 3.14f; // error  
11  
12 class C { ... };  
13 export using D = C;           // error
```



Visibility and reachability

C++ 20 Adv

ratio.cpp

```
1 export module math;  
2 struct Ratio { float num, denom; };  
3 export Ratio golden() { ... } // golden is visible
```

main.cpp

```
1 import math;  
2 import <iostream>;  
3 int main() {  
4     Ratio r = golden(); // error: Ratio not visible  
5     auto r = golden(); // OK: Ratio is reachable  
6     std::cout << r.num; // OK: members of reachable  
7                                     // types are visible  
8     using R = decltype(r);  
9     R r2{1.f, 3.f};           // OK  
10 }
```



Module structure

C++ 20 Adv

```
1 module;                                // global module fragment (opt.)
2 #include <cassert>                      // only preprocessor
3 // ...
4
5 export module math;                     // module preamble starts with
6                               // exported module name.
7 import <string>;                      // only imports allowed
8 // ...
9
10 export struct S { ... }; // module purview started
11 export S f();           // at first non-import
12 std::string h() { ... }
13
14 module :private               // private module fragment (opt.)
15 S f() { ... }                // changes here trigger
16                               // no recompilation of BMI
```



Module structure

Pitfall

```
1 export module math;  
2 #include <header.h>  
3 import <string>;  
4 // ...
```

- Headers included after `export module` add to the module's preamble and purview
- This is usually not what you want

Put includes into global module fragment

```
1 module;  
2 #include <vector>  
3 export module math;  
4 import <string>;  
5 // ...
```



Module structure

C++ 20 Adv

Like with header files, we can split interface and implementation:

Interface

```
1 export module math;  
2  
3 export struct S { ... };  
4 export S f();
```

Implementation

```
1 module;  
2 #include <cassert>  
3 module math;  
4 import <string>;  
5  
6 S f() { ... }  
7 std::string b() { ... }
```



Submodules and re-exporting

C++ 20 Adv

Modules can (re-)export other modules using `export import`:

Module one

```
1 export module one;  
2  
3 export struct S { ... };
```

Module two

```
1 export module two;  
2 import <string>;  
3  
4 export std::string b()  
5 { ... }
```

Module three

```
1 export module three;  
2 import one;  
3 export import two;  
4  
5 export S f() { ... }
```

Module four

```
1 export module four;  
2 import three;  
3 // b and f visible  
4 // S and std::string only  
5 // reachable
```



Module partitions

C++ 20 Adv

We can split a module's *interface* into multiple files:

Primary module interface

```
1 export module math;  
2 export import :structs;  
3  
4 export S f();
```

Module interface partition

```
1  
2 export module math:structs;  
3  
4 export struct S { ... };
```

Implementation

```
1 module math;  
2 import <string>;  
3  
4 S f() { ... }  
5 std::string b() { ... }
```



Module implementation units

C++ 20 Adv

We can split a module's *implementation* into multiple files:

Interface

```
1 export module math;  
2 import :foo; // no export  
3 import :bar; // no export  
4 import <string>;  
5  
6 export struct S { ... };  
7 export S f();  
8 std::string b();
```

Implementation unit foo

```
1 module math:foo;  
2  
3 S f() { ... }
```

Implementation unit bar

```
1 module math:bar;  
2  
3 std::string b() { ... }
```

Standard library modules

C++ 23 Adv

Standard library modules

- Use `import std;` for the entire C++ standard library.
Basically everything in namespace `std`.
 - E.g. `memcpy(a, b, n)`, `sqrt(val)`, `uint32_t`
 - Note that you need to prefix with `std::`:
- Use `import std.compat;` for the entire C and C++ standard library
- No macros are exported. For these, you still need to include the corresponding headers.

```
1 import std;          // everything C++
2 #include <cassert> // for assert()
3 #include <climits> // for e.g. CHAR_BIT
```

- Definition of further, smaller modules of the standard library is in progress (e.g. fundamentals, containers, networking, etc.)



Purpose

- For easier transition from header-based projects to modules
- Allows to import any header file as if it was a module
- Can be mixed with regular header inclusion

```
1 #include <set>
2 import <vector>;
3
4 #include "header.hpp"
5 import "otherheader.hpp";
```

Effect

- All declarations in the imported header will be exported
- An imported header is not affected by the preprocessor state. Thus, you cannot “configure” the header by defining a macro before importing. This allows the header to be precompiled.
- An imported header may still provide preprocessor macros

```
1 #define MY_MACRO 1
2 #include "header.hpp" // may be affected by MY_MACRO
3 import "otherheader.hpp"; // ignores MY_MACRO
```



Automatic include to import translation

C++ 20 Adv

Importable headers

- A header file which can successfully be either included or imported is called “importable”
 - I.e. the header does not require setup of preprocessor state before inclusion
- All C++ standard library headers are importable
- C wrapper headers (`<c*>`) are not

Include translation

- Allows the compiler to automatically turn includes into imports
- Controlled via compiler flags and the build system
- Basically a standard replacement for precompiled headers

Build system

h.hpp

```
#include <vector>
```

...

a.cpp

```
#include "h.hpp"
```

...

b.cpp

```
#include <vector>
```

```
#include "h.hpp"
```


...

c.cpp

```
#include <vector>
```

...

Traditional workflow



Build system

C++ 20 Adv

h.cpp

```
export module foo;
import <vector>;
...
```

a.cpp

```
import foo;
...
```

b.cpp


```
import <vector>;
import foo;
...
```

c.cpp

```
#include <vector>
...
```

New workflow

scan dependencies



Build system

New challenges

- Resolving a module name to the file name building it
- Translation units no longer independent
 - Extra tool needs to infer dependency graph before building
 - Called a “dependency scanner”
 - Needs to be communicated to the build system
 - Using a new standard dependency file format
 - Parallel and distributed builds need synchronization
- Compilation of module translation units produces a binary module interface (BMI) *and* an object file
 - Need to manage BMIs between multiple compiler invocations
- Tools beside the compiler need to build/read BMIs
 - E.g. static analysis, language servers for auto completion, etc.
- The C++ standard specifies very little on how this works
 - We may experience large implementation divergence



Status

- Dependency scanner since clang 16: `clang-scan-deps`
- Experimental module support in cmake since 3.26, uses `clang-scan-deps`, many fixes in later versions
- MSBuild (Windows) fully handles dependency scanning
- `g++ a.cpp b.cpp ...` “just works”, must be one line though
- Experimental extensions to GNU make exist to grow dependency graph on the fly while modules are discovered
- Header units unsupported in all build systems (May 2023)
- C++23: `import std;` supported in MSVC, partially in clang
 - GCC/clang/MSVC also plan support in C++20



Building modules (g++)

C++ 20 Adv

Case study: g++

- BMI is called CMI (compiled module interfaces)
- By default, g++ caches CMIs in subdirectory ./gcm.cache
- Uses a [module mapper](#) to resolve imports, which specifies a protocol and uses a central server for module maintenance
- Each module needs to be built before it can be imported
 - `g++ -std=c++20 -fmodules-ts -c ratio.cpp -o ratio.o`
 - Generates `ratio.o` and `./gcm.cache/ratio.gcm` (CMI)
- Each header unit needs to be built before it can be imported
 - `g++ -std=c++20 -fmodules-ts -x c++-system-header vector`
 - Generates e.g.
`./gcm.cache/usr/include/c++/11/vector.gcm` (CMI)
 - `g++ -std=c++20 -fmodules-ts -x c++-header ratio.h`
 - Generates e.g. `./gcm.cache/,/ratio.h.gcm` (CMI)

Modules

C++ 20 Adv

Guidance for today

- Start writing importable headers (no macro dependencies)
 - Dual use: `#include` before C++ 20 and `import` with C++ 20
- Watch progress on module support in your build system

Guidance for tomorrow

- Start writing modules when your users have C++ 20
- Maybe import standard headers when C++ 20 is available
- Start using `import std;` when available
- Future of header units unclear (implementation difficulties)



Resources

- A (Short) Tour of C++ Modules - Daniela Engert - CppCon 2021
(very technical)
- Understanding C++ Modules: [Part1](#), [Part2](#) and [Part3](#).
- So, You Want to Use C++ Modules ... Cross-Platform? - Daniela Engert - C++ on Sea 2023
- import CMake: 2023 State of C++20 modules in CMake - Bill Hoffman - CppNow 2023



Exercise: Modules

- go to exercises/modules
- convert the Complex.hpp header into a module named math

Exercise: Header units

- go to exercises/header_units
- convert all `#includes` into header unit `imports`



Coroutines

5 Expert C++

- The `<=>` operator
- Variadic templates
- Perfect forwarding
- SFINAE
- Concepts
- Modules
- Coroutines



Why do we need coroutines?

C++ 20 Adv

The generator use case (one of several)

In python one can write

```
1 for i in range(0, 10):
2     print(i)
```

What about it in C++?

```
3 someType range(int first, int last) { ... }
4 for (auto i : range(0, 10)) {
5     std::cout << i << "\n";
6 }
```

How can we implement range?

Requirements

- `someType` needs to be some iterable type
- `range` should execute lazily, only creating values on demand
- meaning `range` should be suspended and resumed

That's what coroutines are for!



What C++ 20 allows

C++ 20 Adv

Interesting part of the implementation (see [cppreference](#))

```
1 template<std::integral T>
2 Generator<T> range(T first, const T last) {
3     while (first < last) {
4         co_yield first++;
5     }
6 }
7 for (const char i : range(65, 91)) {
8     std::cout << i << ' ';
9 }
10 std::cout << '\n';
```

The bad news

- class **Generator** is not provided by the STL
- and is ~ 80 lines of boiler plate code...
- may be sorted out in C++ 23
- in the meantime, we'll have to dive into some details!



A few concepts first

C++ 20 Adv

Definition of a coroutine

- a function which can be suspended and resumed
 - potentially by another caller (even on a different thread)
- practically a function using one of:
 - `co_await` suspends execution
 - `co_yield` suspends execution and returns a value
 - `co_return` completes execution

Implications

- on suspension, the state of the function needs to be preserved
- the suspended function becomes pure data on the heap
- access to this data is given via the returned value
 - so the returned type must follow some convention
 - it needs to contain a `struct promise_type`

Minimal code

C++ 20 Adv

Definition of an empty coroutine

```
3 #include <coroutine>
4 struct Task {
5     struct promise_type {
6         Task get_return_object() { return {}; }
7         std::suspend_never initial_suspend() { return {}; }
8         std::suspend_never final_suspend() noexcept { return {}; }
9         void return_void(){}
10        void unhandled_exception() { throw; }
11    };
12 };
13 Task myCoroutine() { co_return; }
14 int main() { Task x = myCoroutine(); }
```

promise_type interface

`get_return_object` builds a `Task` from the given promise
`initial/final_suspend` called before/after coroutine starts/ends
`unhandled_exception` called in case of exception in coroutine
`return_void/value` returns final value of the coroutine




Minimal sequence of events (nothing suspended)

C++ 20

Adv

Compiler generated set of calls when calling myCoroutine

Credits: original idea for this graph comes from [this web site](#).

Suspending a coroutine

C++ 20 Adv

Main way to suspend a coroutine: `co_wait`

- it's an expression taking an "awaitable"/ "awaiter"
 - we'll ignore the distinction for now
- 2 trivial awaiters are provided by the STL:
`std::suspend_always` and `std::suspend_never`
- obviously only the first one will trigger a suspension
 - note the other one is used in the minimal code
- more details on awaiters later

Code

```
1 Task myCoroutine() {
2     std::cout << "Step 1 of coroutine\n";
3     co_await std::suspend_always{};
4     std::cout << "Step 2 of coroutine\n";
5     co_await std::suspend_always{};
6     std::cout << "final step\n";
7 }
```



Resuming a coroutine

C++ 20 Adv

Principles

- one needs to call `resume()` on the `coroutine_handle`
- BUT user code has no such handle
 - only the coroutine return object (here a `Task` instance)
- so one has to amend the `Task` class
 - and the `promise_type` as it's building the `Task` instance

Code

```
1 struct Task {
2     struct promise_type {
3         Task get_return_object() {
4             // build Task with the coroutine handle
5             return {std::coroutine_handle<promise_type>::from_promise(*this)};
6         }
7         ... // rest is unchanged
8     };
9     std::coroutine_handle<promise_type> h_;
10    void resume() { h_.resume(); }
11};
```



Resuming a coroutine

C++ 20 Adv

User code - godbolt

```
1 Task myCoroutine() {
2     std::cout << "Step 1 of coroutine\n";
3     co_await std::suspend_always{};
4     std::cout << "Step 2 of coroutine\n";
5     co_await std::suspend_always{};
6     std::cout << "final step\n";
7 }
8 int main() {
9     auto c = myCoroutine(); // Step 1 runs
10    std::cout << "In main, between Step 1 and 2\n";
11    c.resume();           // Step 2 runs
12    std::cout << "In main, between Step 2 and final step\n";
13    c.resume();           // final Step runs
14    // c.resume(); // would segfault!
15 }
```

```
Step 1 of coroutine
In main, between Step 1 and 2
Step 2 of coroutine
In main, between Step 2 and final step
final step
```

A few more details

C++ 20 Adv

About suspension of coroutines

- initial/final_suspend methods allow to suspend a coroutine at the very start/end (after `co_return`)
- a suspended coroutine is not deallocated
 - i.e. its state will stay on the heap
 - so you can leak coroutines!
- one can directly call `coroutine_handle::destroy` if the execution of a suspended coroutine should be aborted
 - which requires another piece of code in Task, e.g.
`~Task() { if (h_) h_.destroy(); }`
- the Task class should be move only

We need standard Task classes in the STL!

- `std::generator` may come in C++ 23

co_yield - returning a value upon suspension C++ 20 Adv

What co_yield really is

- an expression like `co_await`
- a shortcut for `co_await promise.yield_value(expr)`
- and we have to provide the `yield_value` method of `promise_type` and the Task accessor

Typical code

```
1 template<typename T> struct Task {
2     struct promise_type {
3         T value_;
4         std::suspend_always yield_value(T&& from) {
5             value_ = std::forward<T>(from); // caching the result
6             return {};
7         }
8         ... // rest is unchanged
9     };
10    T getValue() { return h_.promise().value_; } // accessor for user code
11};
```



co_yield - generator behavior

C++ 20 Adv

Generator behavior - more boiler plate code

```
1 template<typename T>
2 struct Task {
3     struct iterator {
4         std::coroutine_handle<promise_type> handle;
5         auto &operator++() {
6             handle.resume();
7             if (handle.done()) { handle = nullptr; } // == end iterator
8             return *this;
9         }
10        auto operator++(int) { ++*this; }
11        T const& operator*() const { return handle.promise().value_; }
12        ...
13    };
14    iterator begin() { resume(); return {h_}; }
15    iterator end() { return {nullptr}; }
16    ... // plus error handling is missing
17};
```



co_yield - final usage

C++ 20 Adv

Finally, we can use the generator nicely - godbolt

```
1 Task range(int first, int last) {
2     while (first != last) {
3         co_yield first++;
4     }
5 }
6 for (int i : range(0, 10)) {
7     std::cout << i << "";
8 } // 1 2 3 4 5 6 7 8 9
```



co_yield - final usage

C++ 20 Adv

Finally, we can use the generator nicely - godbolt

```
1 Task range(int first, int last) {
2     while (first != last) {
3         co_yield first++;
4     }
5 }
6 for (int i : range(0, 10)) {
7     std::cout << i << "";
8 } // 1 2 3 4 5 6 7 8 9
```

Wait a minute: 0 is missing in the output!

- resume called in `operator++` before 0 is printed
- we should pause in `initial_suspend`

```
10 struct Task {
11     struct promise_type {
12         ...
13         std::suspend_always initial_suspend() { return {}; }
14         ...
15     }
16 };
```



Laziness allows for infinite ranges

C++ 20 Adv

Generators do not need to be finite - godbolt

```
1 Task range(int first) {
2     while (true) {
3         co_yield first++;
4     }
5 }
6 for (int i : range(0) | std::views::take(10)) {
7     std::cout << i << "\n";
8 }
```



Real life example

C++ 20 Adv

A Fibonacci generator, from cppreference - godbolt

```
1 Task<long long int> fibonacci(unsigned n) {
2     if (n==0) co_return;
3     co_yield 0;
4     if (n==1) co_return;
5     co_yield 1;
6     if (n==2) co_return;
7     uint64_t a=0;
8     uint64_t b=1;
9     for (unsigned i = 2; i < n; i++) {
10         uint64_t s=a+b;
11         co_yield s;
12         a=b;
13         b=s;
14     }
15 }
16
17 int main() {
18     for (int i : fibonacci(10)) {
19         std::cout << i << "\n";
20     }
21 }
```

Coming back to awaitables/awaiters

C++ 20 Adv


awaitables vs awaiters

- awaitable is the type required by `co_await/co_yield`
 - or the expression given has to be convertible to awaitable
 - alternatively the `promise_type` can have an `await_transform` method returning an awaitable
- an awariter is retrieved from the awaitable by the compiler
 - either by calling operator `co_wait` on it
 - or via direct conversion if no such operator exists
- an awariter has 3 main methods
 - `await_ready`, `await_suspend` and `await_resume`



How awaiters work

Compiler generated set of calls



awaiter example

Switch to new thread, from cppreference - godbolt

```
1 auto switch_to_new_thread(std::jthread& out) {
2     struct awaiter {
3         std::jthread* th;
4         bool await_ready() { return false; }
5         void await_suspend(std::coroutine_handle<> h) {
6             *th = std::jthread([h] { h.resume(); });
7             std::cout << "New thread ID: " << th->get_id() << '\n';
8         }
9         void await_resume() {}
10    };
11    return awaiter{&out};
12 }
13 Task<int> resuming_on_new_thread(std::jthread& out) {
14     std::cout << "Started on " << std::this_thread::get_id() << '\n';
15     co_await switch_to_new_thread(out);
16     std::cout << "Resumed on " << std::this_thread::get_id() << std::endl;
17 }
18 int main() {
19     std::jthread out;
20     resuming_on_new_thread(out);
21 }
```

awaiter example

Switch to new thread, from cppreference - godbolt

```
1 auto switch_to_new_thread(std::jthread& out) {
2     struct awaiter {
3         std::jthread* th;
4         bool await_ready() { return false; }
5         void await_suspend(std::coroutine_handle<> h) {
6             *th = std::jthread([h] { h.resume(); });
7             std::cout << "New thread ID: " << th->get_id() << '\n';
8         }
9         void await_resume() {}
10    };
11    return awaiter{&out};
12 }
13 Task<int> resuming_on_new_thread(std::jthread& out) {
14     std::cout << "Started on " << std::this_thread::get_id() << '\n';
15     co_await switch_to_new_thread(out);
16     std::cout << "Resumed on " << std::this_thread::get_id() << std::endl;
17 }
18 int main() {
19     std::jthread out;
20     resuming_on_new_thread(out);
21 }
```

Output

```
Started on thread: 140144191063872
New thread ID: 140144191059712
Resumed on thread: 140144191059712
```



Summary on coroutines

C++ 20 Adv

Coroutines are usable in C++ 20

- allow to have nice generators
- also very useful for async/IO code
 - another big use case which we did not touch

But library support is poor for now

- standard Task/Generator classes are not provided
- leading to *lots of boiler plate code* for the end user
- std::generator is part of C++ 23's draft
- in the meantime, libraries exist, e.g. [CppCoro](#)



Useful tools

1 History and goals

- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen

2 Language basics

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

6 Useful tools

- C++ editor

7 Concurrency

8 C++ and python



C++ editor

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen

C++ editors and IDEs

Can dramatically improve your efficiency by

- Coloring the code for you to “see” the structure
- Helping with indenting and formatting properly
- Allowing you to easily navigate in the source tree
- Helping with compilation/debugging, profiling, static analysis
- Showing you errors and suggestions while typing

► Visual Studio Heavy, fully fledged IDE for Windows

► Visual Studio Code Editor, open source, portable, many plugins

► Eclipse IDE, open source, portable

► Emacs ► Vim Editors for experts, extremely powerful.

They are to IDEs what latex is to PowerPoint

CLion, Code::Blocks, Atom, NetBeans, Sublime Text, ...

Choosing one is mostly a matter of taste



Version control

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen

Version control

Please use one!

- Even locally
- Even on a single file
- Even if you are the only committer

It will soon save your day

A few tools

► git THE mainstream choice. Fast, light, easy to use

► mercurial The alternative to git

► Bazaar Another alternative

► Subversion Historical, not distributed - don't use

► CVS Archeological, not distributed - don't use



Git crash course

```
$ git init myProject
```

Initialized empty Git repository in myProject/.git/

```
$ vim file.cpp; vim file2.cpp
```

```
$ git add file.cpp file2.cpp
```

```
$ git commit -m "Committing first 2 files"
```

[master (root-commit) c481716] Committing first 2 files

...

```
$ git log --oneline
```

d725f2e Better STL test

f24a6ce Reworked examples + added stl one

bb54d15 implemented template part

...

```
$ git diff f24a6ce bb54d15
```

Code formatting

6 Useful tools

- C++ editor
- Version control
- **Code formatting**
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen

clang-format

.clang-format

- File describing your formatting preferences
- Should be checked-in at the repository root (project wide)
- `clang-format -style=LLVM -dump-config > .clang-format`
- Adapt style options with help from: <https://clang.llvm.org/docs/ClangFormatStyleOptions.html>

Run clang-format

- `clang-format --style=LLVM -i <file.cpp>`
- `clang-format -i <file.cpp>` (looks for .clang-format file)
- `git clang-format` (formats local changes)
- `git clang-format <ref>` (formats changes since git <ref>)
- Most editors/IDEs can find a .clang-format file and adapt

clang-format

Exercise: clang-format

- Go to any example
- Format code with:
`clang-format --style=GNU -i <file.cpp>`
- Inspect changes, try `git diff .`
- Revert changes using `git checkout -- <file.cpp>` or
`git checkout .`
- Go to exercises directory and create a `.clang-format` file
`clang-format -style=LLVM -dump-config > .clang-format`
- Run `clang-format -i <any_exercise>/*.cpp`
- Revert changes using `git checkout <any_exercise>`

The Compiling Chain


6 Useful tools

- C++ editor
- Version control
- Code formatting
- **The Compiling Chain**
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen



The compiling chain

C++ 17



The steps

cpp the preprocessor

handles the `#` directives (macros, includes)

creates “complete” source code (ie. translation unit)

g++ the compiler

creates machine code from C++ code

ld the linker

links several binary files into libraries and executables



Compilers

Available tools

▶ **gcc** the most common and most used
free and open source

▶ **clang** drop-in replacement of gcc
free and open source, based on LLVM

▶ **icc** ▶ **icx** Intel's compilers, proprietary but now free
optimized for Intel hardware
icc being replaced by icx, based on LLVM

▶ **Visual C++ / MSVC** Microsoft's C++ compiler on Windows

My preferred choice today

- **gcc** as the de facto standard in HEP
- **clang** in parallel to catch more bugs

Useful compiler options (gcc/clang)

Get more warnings

-Wall -Wextra get all warnings

-Werror force yourself to look at warnings

Optimization

-g add debug symbols (also: -g3 or -ggdb)

-Ox 0 = no opt., 1-2 = opt., 3 = highly opt. (maybe larger binary), g = opt. for debugging

Compilation environment

-I <path> where to find header files

-L <path> where to find libraries

-l <name> link with libname.so

-E / -c stop after preprocessing / compilation



How to inspect object files?

Listing symbols : nm

- gives list of symbols in a file
 - these are functions and constants
 - with their internal (mangled/encoded) naming
- also gives type and location in the file for each symbol
 - 'U' type means undefined
 - so a function used but not defined
 - linking will be needed to resolve it
- use -C option to demangle on the fly

```
> nm -C Struct.o
              U strlen
              U __Unwind_Resume
0000000000000008a T SlowToCopy::SlowToCopy(SlowToCopy const&)
00000000000000000000 T SlowToCopy::SlowToCopy()
0000000000000000064 T SlowToCopy::SlowToCopy(std::__cxx11::basic_string<char> const&)
```

How to inspect libraries/executables?

Listing dependencies : ldd

- gives (recursive) list of libraries required by the given argument
 - and if/where they are found in the current context
- use `-r` to list missing symbols (mangled)

```
> ldd -r trypoly
    linux-vdso.so.1 (0x00007f3938085000)
    libpoly.so => not found
    libstdc++.so.6 => /lib/x86_64-linux-gnu/libstdc++.so.6 (0x00
    [...]
    undefined symbol: _ZNK7Hexagon16computePerimeterEv      (./try
    undefined symbol: _ZNK7Polygon16computePerimeterEv      (./try
    undefined symbol: _ZN7HexagonC1Ef        (./trypoly.sol)
    undefined symbol: _ZN8PentagonC1Ef      (./trypoly.sol)
```



Makefiles

Why to use them

- an organized way of describing building steps
- avoids a lot of typing

Several implementations

- raw Makefiles: suitable for small projects
- cmake: portable, the current best choice
- automake: GNU project solution

```
test : test.cpp libpoly.so
       $(CXX) -Wall -Wextra -o $@ $^
libpoly.so: Polygons.cpp
       $(CXX) -Wall -Wextra -shared -fPIC -o $@ $^
clean:
       rm -f *o *so *~ test test.sol
```



CMake

- a cross-platform meta build system
- generates platform-specific build systems
- see also this [basic](#) and [detailed](#) talks

Example CMakeLists.txt

```
1 cmake_minimum_required(VERSION 3.18)
2 project(hello CXX)
3
4 find_package(ZLIB REQUIRED) # for external libs
5
6 add_executable(hello main.cpp util.h util.cpp)
7 set(CMAKE_CXX_STANDARD 20)
8 target_link_libraries(hello PUBLIC ZLIB::ZLIB)
```



CMake - Building

Building a CMake-based project

Start in the directory with the top-level CMakeLists.txt:

```
1 mkdir build # will contain all build-related files
2 cd build
3 cmake ..      # configures and generates a build system
4 cmake -DCMAKE_BUILD_TYPE=Release .. # pass arguments
5 ccmake .       # change configuration using terminal GUI
6 cmake-gui .   # change configuration using Qt GUI
7 cmake --build . -j8     # build project with 8 jobs
8 cmake --build . --target hello # build only hello
9 sudo cmake --install . # install project into system
10 cd ..
11 rm -r build # clean everything
```



Compiler chain

Exercise: Compiler chain

- go to exercises/polymorphism
- preprocess Polygons.cpp (`g++ -E -o output`)
- compile Polygons.o and trypoly.o (`g++ -c -o output`)
- use nm to check symbols in .o files
- look at the Makefile
- try make clean; make
- see linking stage of the final program using `g++ -v`
 - just add a `-v` in the Makefile command for trypoly target
 - run make clean; make
 - look at the collect 2 line, from the end up to “`-o trypoly`”
- see library dependencies of ‘trypoly’ using ‘`ldd`’



Web tools

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- **Web tools**
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen

Godbolt / Compiler Explorer

Concept

An online generic compiler with immediate feedback. Allows:

- trying various compilers in the browser
- inspecting the assembly generated
- use of external libraries (over 50 available !)
- running the produced code
- sharing small pieces of code via permanent short links

Typical usage

- check small pieces of code on different compilers
- check some new C++ functionality and its support
- optimize small pieces of code
- NOT relevant for large codes



Godbolt by example

Check effect of optimization flags

<https://godbolt.org/z/Pb8WsWjEx>

- Check generated code with -O0, -O1, -O2, -O3
- See how it gets shorter and simpler

The screenshot shows the Godbolt Compiler Explorer interface with three main panes:

- Left pane (C++ source code):**

```
1 #include <iostream>
2
3 constexpr int fact(int a) {
4     int n = 1;
5     for (int i = 1; i <=a; i++) {
6         n *= i;
7     }
8     return n;
9 }
10
11 int main() {
12     for (int i = 4; i < 8; i++) {
13         std::cout << fact(i);
14     }
15 }
```
- Middle pane (-O0 optimization level):**

```
1 .fact(int):
2     push rbp
3     mov rbp, rsp
4     mov DWORD PTR [rbp-20], edi
5     mov DWORD PTR [rbp-4], 1
6     mov DWORD PTR [rbp-8], 1
7     jmp .L2
8 .L3:
9     mov eax, DWORD PTR [rbp-4]
10    imul eax, DWORD PTR [rbp-8]
11    mov DWORD PTR [rbp-4], eax
12    add DWORD PTR [rbp-8], 1
13 .L2:
14    mov eax, DWORD PTR [rbp-8]
15    cmp eax, DWORD PTR [rbp-20]
16    jle .L3
17    mov eax, DWORD PTR [rbp-4]
18    pop rbp
19    ret
20 main:
21     push rbp
22     mov rbp, rsp
23     sub rsp, 16
24     mov DWORD PTR [rbp-4], 4
25     jmp .L6
26 .L7:
27     mov eax, DWORD PTR [rbp-4]
28     mov edi, eax
29     call fact(int)
30     mov esi, eax
31     mov edi, OFFSET FLAT:_ZSt4cout
32     call std::basic_ostream<char, std::char_traits<char>>::operator<<(void*)
```
- Right pane (-O3 optimization level):**

```
1 main:
2     sub    rsp, 8
3     mov    esi, 24
4     mov    edi, OFFSET FLAT:_ZSt4cout
5     call   std::basic_ostream<char, std::char_traits<char>>::operator<<(void*)
6     mov    esi, 120
7     mov    edi, OFFSET FLAT:_ZSt4cout
8     call   std::basic_ostream<char, std::char_traits<char>>::operator<<(void*)
9     mov    esi, 720
10    mov   edi, OFFSET FLAT:_ZSt4cout
11    call   std::basic_ostream<char, std::char_traits<char>>::operator<<(void*)
12    mov    esi, 5040
13    mov   edi, OFFSET FLAT:_ZSt4cout
14    call   std::basic_ostream<char, std::char_traits<char>>::operator<<(void*)
15    xor    eax, eax
16    add    rsp, 8
17    ret
18 _GLOBAL_sub_1_main:
19     sub    rsp, 8
20     mov    edi, OFFSET FLAT:_ZSt8__join
21     call   std::ios_base::Init::Init()
22     mov    edx, OFFSET FLAT:_ZSt8_dso_handle
23     mov    esi, OFFSET FLAT:_ZSt8__join
24
25
26
27
28
29
30
31
32
```

Below the panes, the status bar shows:

- Output (0): x86-64 gcc 12.1 | 600ms (98817B) - 5874 lines filtered
- Output of x86-64 gcc 12.1 (Compiler #3) ×
- ASM generation compiler returned: 0
- Execution build compiler returned: 0
- Program returned: 0
- 241287285840

cppinsights

Concept

Reveals the actual code behind C++ syntactic sugar

- lambdas
- range-based loops
- templates
- initializations
- auto
- ...

Typical usage

- understand how things work behind the C++ syntax
- debug some non working pieces of code



cppinsights by example

Check how range-based loop work

<https://cppinsights.io/s/b886aa76>

- See how they map to regular iterators
 - And how operators are converted to function calls

The screenshot shows a C++ code editor interface. The top bar includes tabs for 'Source' and 'Insight', and a status bar indicating 'C++ Standard: C++ 20'. The 'Source' tab displays the following code:

```
1 #include <array>
2 #include <iostream>
3
4 int main() {
5     std::array<int, 5> arr{2,4,6,8,10};
6     for(auto const & c : arr) {
7         std::cout << "c=" << c << "\n";
8     }
9 }
```

The 'Insight' tab shows a generated C++11 equivalent for the code above:

```
1 #include <array>
2 #include <iostream>
3
4 int main()
5 {
6     std::array<int, 5> arr = {2, 4, 6, 8, 10};
7     {
8         std::array<int, 5> __rangel = arr;
9         int * __begin = __rangel.begin();
10        int * __end = __rangel.end();
11        for(; __begin != __end; ++__begin) {
12            const std::array<int, 5>::value_type & c = *__begin;
13            std::operator<<(std::operator<<(std::operator<<(std::cout, "c="), c).operator<<("\n"));
14        }
15    }
16    return 0;
17 }
18
19
```

The bottom left corner shows a 'Console' tab and the message 'Insights exited with result code: 0'.



Debugging

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen



Debugging

The problem

- everything compiles fine (no warning)
- but crashes at run time
- no error message, no clue



Debugging

The problem

- everything compiles fine (no warning)
- but crashes at run time
- no error message, no clue

The solution: debuggers

- dedicated program able to stop execution at any time
- and show you where you are and what you have



Debugging

The problem

- everything compiles fine (no warning)
- but crashes at run time
- no error message, no clue

The solution: debuggers

- dedicated program able to stop execution at any time
- and show you where you are and what you have

Existing tools

► **gdb** THE main player

► **lldb** the debugger coming with clang/LLVM

► **gdb-oneapi** the Intel OneAPI debugger



gdb crash course

start gdb

- `gdb <program>`
- `gdb <program><core file>`
- `gdb --args <program><program arguments>`

inspect state

`bt` prints a backtrace

`print <var>` prints current content of the variable

`list` show code around current point

`up/down` go up or down in call stack

breakpoints

`break <function>` puts a breakpoint on function entry

`break <file>:<line>` puts a breakpoint on that line



gdb

Exercise: gdb

- go to exercises/debug
- compile, run, see the crash
- run it in gdb (or lldb on newer MacOS)
- inspect backtrace, variables
- find problem and fix bug
- try stepping, breakpoints



Debugging UIs

User interfaces for debuggers

- offer convenience on top of command line
- windows for variables, breakpoints, call stack, active threads, watch variables in-code, disassembly, run to cursor ...

Native gdb Try “tui enable” for a simple built-in UI

▶ VSCode Built-in support for gdb

▶ CodeLLDB VS Code plugin for LLDB

▶ GDB dashboard Poplar terminal UI for gdb

▶ GEF Modern terminal UI for gdb

- some editors and most IDEs have good debugger integration



Sanitizers

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- **Sanitizers**
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen



Address Sanitizer (ASan)

ASan introduction

- Compiler instrumentation
- Program stops on invalid memory access, e.g.
 - Invalid read/write on heap and stack
 - Double free/delete, use after free
 - Buffer overflow on stack (few tools can do this)
 - Only linux: memory leaks



Address Sanitizer (ASan)

ASan introduction

- Compiler instrumentation
- Program stops on invalid memory access, e.g.
 - Invalid read/write on heap and stack
 - Double free/delete, use after free
 - Buffer overflow on stack (few tools can do this)
 - Only linux: memory leaks

Usage (gcc/clang syntax)

- Compile with
 - `-fsanitize=address -fno-omit-frame-pointer -g`
- With clang, add optionally:
 - `-fsanitize-address-use-after-return=always`
 - `-fsanitize-address-use-after-scope`
- Link with `-fsanitize=address`
- Run the program



Address Sanitizer (ASan)

How it works

- Compiler adds run-time checks ($\sim 2x$ slow down)
- IsPoisoned(address) looks up state of address in asan's "shadow memory"
- Shadow memory: memory where 1 shadow byte tracks state of 8 application bytes (state = accessible, poisoned, ...)
- Functions that deal with memory (`new()` / `delete()` / strings / ...) update entries in shadow memory when called

asan instrumentation (mock code)

```
1 int i = *address;
```



Address Sanitizer (ASan)

How it works

- Compiler adds run-time checks ($\sim 2x$ slow down)
- IsPoisoned(address) looks up state of address in asan's "shadow memory"
- Shadow memory: memory where 1 shadow byte tracks state of 8 application bytes (state = accessible, poisoned, ...)
- Functions that deal with memory (`new()` / `delete()` / strings / ...) update entries in shadow memory when called

asan instrumentation (mock code)

```
1 if (IsPoisoned(address)) {  
2     ReportError(address, kAccessSize, kIsWrite);  
3 }  
4 int i = *address;
```

ASan red zones

- If adjacent data blocks are owned by the process, the operating system will allow an access

Illegal access (not detected without ASan)

```
1 void foo() {  
2     char a[8];  
3     char b[8];  
4     a[8] = '1';  
5 }
```

Memory layout



ASan red zones

- If adjacent data blocks are owned by the process, the operating system will allow an access
- ASan surrounds blocks of memory by poisoned red zones
- Program stops when accessing a red zone

Illegal access (not detected without ASan)

```
void foo() {  
+   char redzone1[32];  
    char a[8];  
+   char redzone2[24];  
    char b[8];  
+   char redzone3[24];  
+   // <poison redzones>  
    a[8] = '1';  
+   // <unpoison redzones>  
}
```

Memory layout

		0x300C
b[0-7]	redzone3	0x3008
a[0-7]	redzone2	0x3004
	redzone1	0x3000



```
==34015==ERROR: AddressSanitizer: stack-buffer-overflow on address 0x7ffee93ed968 a
WRITE of size 1 at 0x7ffee93ed968 thread T0
#0 0x106812df3 in foo() asan.cpp:4
#1 0x106812ed8 in main asan.cpp:9
#2 0x7fff6d3923d4 in start (libdyld.dylib:x86_64+0x163d4)
```

Address 0x7ffee93ed968 is located in stack of thread T0 at offset 40 in frame
#0 0x106812cdf in foo() asan.cpp:1

This frame has 2 object(s):

[32, 40) 'a' (line 2) <== Memory access at offset 40 overflows this variable
[64, 72) 'b' (line 3)

Shadow bytes around the buggy address:

```
=>0xfffffd27db20: 00 00 00 00 00 00 00 00 f1 f1 f1 f1 00[f2]f2 f2
 0xfffffd27db30: 00 f3 f3 f3 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0xfffffd27db40: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0xfffffd27db50: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
```

Shadow byte legend (one shadow byte represents 8 application bytes):

Addressable: 00

Partially addressable: 01 02 03 04 05 06 07

Heap left redzone: fa

Freed heap region: fd

Stack left redzone: f1

Stack mid redzone: f2

Stack right redzone: f3

Stack after return: f5

Finding memory leaks with ASan

- On linux, ASan can display memory leaks (“LeakSanitizer”)
- Might be active already, otherwise start executable with
`ASAN_OPTIONS=detect_leaks=1 ./myProgram`

```
==113262==ERROR: LeakSanitizer: detected memory leaks
```

Direct leak of 32 byte(s) in 1 object(s) allocated from:

```
#0 0x7f2671201647 in operator new(unsigned long) /build/dkonst/WORK/build/contrib
#1 0x4033c7 in memoryLeak[abi:cxx11]() /afs/cern.ch/user/s/shageboe/asan.cpp:33
#2 0x403633 in main /afs/cern.ch/user/s/shageboe/asan.cpp:40
#3 0x7f2670a15492 in __libc_start_main (/lib64/libc.so.6+0x23492)
```

Indirect leak of 22 byte(s) in 1 object(s) allocated from:

```
#0 0x7f2671201647 in operator new(unsigned long) /build/dkonst/WORK/build/contrib
#1 0x403846 in void std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char>>::operator=(const std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char>> &) /afs/cern.ch/user/s/shageboe/asan.cpp:33
#2 0x4033f4 in std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char>>>
#3 0x4033f4 in memoryLeak[abi:cxx11]() /afs/cern.ch/user/s/shageboe/asan.cpp:33
#4 0x403633 in main /afs/cern.ch/user/s/shageboe/asan.cpp:40
#5 0x7f2670a15492 in __libc_start_main (/lib64/libc.so.6+0x23492)
```

SUMMARY: AddressSanitizer: 54 byte(s) leaked in 2 allocation(s).



Address sanitizer (ASan)

Wrap up

- If a program crashes, run it with asan
- Should be part of every C++ continuous integration system
- It will also find bugs that by luck didn't crash the program
- It doesn't generate false positives

More info

- <https://github.com/google/sanitizers/wiki/AddressSanitizer>
- Compile with asan, and start executable using
`ASAN_OPTIONS=help=1 <executable>`



Address sanitizer (ASan)

Exercise: address sanitizer

- Go to [exercises/asan](#)
- Compile and run the program `./asan`
- There are two bugs and one memory leak. Use asan to trace them down.



Thread sanitizer (TSan)

TSan

- Thread sanitizer detects many data races in MT programs
- Recompile your program with e.g.
`clang++ -fsanitize=thread -g -O1 datarace.cpp`

```
% ./a.out
```

```
WARNING: ThreadSanitizer: data race (pid=19219)
```

```
Write of size 4 at 0x7fcf47b21bc0 by thread T1:
```

```
#0 Thread1 datarace.c:4 (exe+0x00000000a360)
```

```
Previous write of size 4 at 0x7fcf47b21bc0 by main thread:
```

```
#0 main datarace.c:10 (exe+0x00000000a3b4)
```

```
Thread T1 (running) created at:
```

```
#0 pthread_create tsan_interceptors.cc:705 (exe+0x00000000c790)
#1 main datarace.c:9 (exe+0x00000000a3a4)
```

<https://github.com/google/sanitizers/wiki/ThreadSanitizerCppManual>



Undefined Behaviour Sanitizer (UBSan)

UBSan

- Tracks uninitialized memory, broken arithmetic, wrong array indexing and other undefined behaviour
- Recompile your program with e.g.
`clang++ -fsanitize=undefined -g -O1 ub.cpp`

```
% ./a.out  
up.cpp:3:5: runtime error: signed integer overflow:  
    2147483647 + 1 cannot be represented in type 'int'
```

<https://clang.llvm.org/docs/UndefinedBehaviorSanitizer.html>



Undefined Behaviour Sanitizer (UBSan)

Exercise: Finding evil run-time bugs

- Go to `exercises/ubsan`
- Compile and run the program following the instructions in `README` or in the program
- The program should run without observable issues
- Recompile with UBSan and see that almost every second line contains evil bugs



The Valgrind family

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- **The Valgrind family**
- Static code analysis
- Profiling
- Doxygen



The valgrind family

Valgrind fundamentals

- valgrind is a framework for different tools
- a processor simulator allowing checks in between instructions
- slow (10-50 times slower than normal execution)
- easy to use : “valgrind <your executable>”
 - no recompilation
 - better with -g -O0, but not strictly needed
- it is free and open source



The valgrind family

Valgrind fundamentals

- valgrind is a framework for different tools
- a processor simulator allowing checks in between instructions
- slow (10-50 times slower than normal execution)
- easy to use : “valgrind <your executable>”
 - no recompilation
 - better with -g -O0, but not strictly needed
- it is free and open source

Main tools

memcheck a memory checker (default tool) and leak detector

callgrind a call graph builder

helgrind a race condition detector



memcheck

- keeps track of all memory allocations and deallocations
- is able to detect accesses to unallocated memory
- and even tell you when it was deallocated if it was
- or what is the closest array in case of overflow
- is able to list still allocated memory when program exits (memory leaks detection)



valgrind

Exercise: valgrind

- go to `exercises/valgrind`
- compile, run, it should work
- run with valgrind, see the problem
- fix the problem

- go back to the `exercises/debug` exercise
- check it with valgrind
- analyze the issue, see that the variance was biased
- fix the issue



memcheck

Exercise: memcheck

- go to `exercises/memcheck`
- compile, run, it should work
- run with valgrind, see LEAK summary
- run with `--leak-check=full` to get details
- analyze and correct it



callgrind and kcachegrind

callgrind

- keeps track of all function calls
- and time spent in each function
- build statistics on calls, CPU usages and more
- outputs flat statistics file, quite unreadable

kcachegrind

- a gui exploiting statistics built by callgrind
- able to browse graphically the program calls
- able to “graph” CPU usage on the program structure



callgrind

Exercise: callgrind

- go to exercises/callgrind
- compile, run, it will be slow
- change nb iterations to 20
- run with valgrind --tool=callgrind
- look at output with kcachegrind
- change fibo call to fibo2
- observe the change in kcachegrind



helgrind

- keeps track of all pthreads activity
- in particular keeps track of all mutexes
- builds a graph of dependencies of the different actions
- works on the resulting graph to detect:
 - possible dead locks
 - possible data races



helgrind

- keeps track of all pthreads activity
- in particular keeps track of all mutexes
- builds a graph of dependencies of the different actions
- works on the resulting graph to detect:
 - possible dead locks
 - possible data races

Note the “possible”. It finds future problems!



helgrind

Exercise: helgrind

- go to exercises/helgrind
- compile, run
- check it with valgrind. You may see strange behavior or it will be perfectly fine
- check it with valgrind --tool=helgrind
- understand issue and fix



Static code analysis

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- **Static code analysis**
- Profiling
- Doxygen



Static analysis

The problem

- all the tools discussed so far work on binaries
- they analyze the code being run
- so there is a coverage problem (e.g. for error cases)



Static analysis

The problem

- all the tools discussed so far work on binaries
- they analyze the code being run
- so there is a coverage problem (e.g. for error cases)

A (partial) solution : analyzing the source code

- build a graph of dependencies of the calls
- use graph tools to detect potential memory corruptions, memory leaks or missing initializations



Static analysis

The problem

- all the tools discussed so far work on binaries
- they analyze the code being run
- so there is a coverage problem (e.g. for error cases)

A (partial) solution : analyzing the source code

- build a graph of dependencies of the calls
- use graph tools to detect potential memory corruptions, memory leaks or missing initializations

Existing tools

- ▶ Coverity proprietary tool, the most complete
- ▶ cppcheck free and opensource, but less complete
- ▶ clang-tidy clang-based “linter”, includes clang static analyzer



cppcheck

Exercise: cppcheck

- go to exercises/cppcheck
- compile, run, see that it works
- use valgrind: no issue
- use cppcheck, see the problem
- analyze the issue, and fix it
- bonus: understand why valgrind did not complain
and how the standard deviation could be biased
hint : use gdb and check addresses of v and diffs



clang-tidy

Documentation and supported checks

- <https://clang.llvm.org/extra/clang-tidy/>

Run clang-tidy

- clang-tidy <file.cpp> -checks=...
- clang-tidy <file.cpp> (checks from .clang-tidy file)
- clang-tidy <file.cpp> --fix (applies fixes)

Compilation flags

- clang-tidy needs to know exactly how your program is built
- clang-tidy ... -- <all compiler flags>

.clang-tidy file

- describes which checks to run
- usually checked in at repository root



Automatically collecting compilation flags

- clang-tidy looks for a file called `compile_commands.json`
- contains the exact build flags for each `.cpp` file
- generate with CMake:
`cmake -DCMAKE_EXPORT_COMPILE_COMMANDS=ON ...`
- for Makefiles try ▶ Bear
- allows to run clang-tidy in bulk on all files:
 - `run-clang-tidy -checks ...`
 - `run-clang-tidy (checks from .clang-tidy)`
 - `run-clang-tidy -fix (applies fixes)`



clang-tidy

Exercise: clang-tidy

- go to any example which compiles (e.g. exercises/cppcheck)
- `mkdir build && cd build`
- `cmake -DCMAKE_EXPORT_COMPILE_COMMANDS=ON ..`
- `clang-tidy <../file.cpp> -checks=*`
- inspect output
- run with `--fix` flag
- revert changes using `git checkout <../file.cpp>`



Profiling

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- **Profiling**
- Doxygen



Profiling

Conceptually

- take a measurement of a performance aspect of a program
 - where in my code is most of the time spent?
 - is my program compute or memory bound?
 - does my program make good use of the cache?
 - is my program using all cores most of the time?
 - how often are threads blocked and why?
 - which API calls are made and in which order?
 - ...
- the goal is to find performance bottlenecks
- is usually done on a compiled program, not on source code



perf, VTune and uProf

perf

- perf is a powerful command line profiling tool for linux
- compile with `-g -fno-omit-frame-pointer`
- `perf stat -d <prg>` gathers performance statistics while running `<prg>`
- `perf record -g <prg>` starts profiling `<prg>`
- `perf report` displays a report from the last profile
- More information in [this wiki](#), [this website](#) or [this talk](#).

Intel VTune and AMD uProf

- Graphical profilers from CPU vendors with rich features
- Needs vendor's CPU for full experience
- More information on [Intel's website](#) and [AMD's website](#)



Doxxygen

6 Useful tools

- C++ editor
- Version control
- Code formatting
- The Compiling Chain
- Web tools
- Debugging
- Sanitizers
- The Valgrind family
- Static code analysis
- Profiling
- Doxygen



Doxxygen

Doxxygen

- Generates documentation from source code comments
- Output formats: HTML, LaTeX, XML, ...
 - May be input for further generators, e.g. Sphinx
- Doxygen uses a config file, usually called Doxyfile
- Run doxygen -g to generate an initial Doxyfile
- Edit Doxyfile (e.g. output format, source code location, etc.)
- Run doxygen to (re-)generate documentation
- View e.g. HTML documentation using a standard web browser
- More on the [doxygen website](#)



Doxxygen

Comment blocks

```

1  /**
2   * text/directives ...
3   */
4 some_cpp_entity
5
6 */
7   * text/directives ...
8   */
9 some_cpp_entity
10  ///
11  /// text/directives ...
12  ///
13 some_cpp_entity
14
15  //!
16  //! text/directives ...
17  //!
18 some_cpp_entity

```

- More details available [here](#)

Comment blocks for members

```
int a_class_member; //!<< trailing documentation
```



Doxxygen

Comment directives

```
1  /**
2   * Long description here. May include HTML etc.
3   *
4   * \brief Checks whether i is odd
5   * \param Integral The integral type of the input
6   * \param i Input value
7   * \return True if i is odd, otherwise false
8   * \throw std::out_of_range if i is larger than 100
9   * \see isEven
10  */
11 template <typename Integral>
12 bool isOdd(Integral i);
```

- All directives can also start with @ instead of \
- A list of all commands can be found [here](#)

Concurrency

1 History and goals

2 Language basics

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

6 Useful tools

7 Concurrency

- Threads and async
- Mutexes
- Atomic types
- Thread-local storage
- Condition Variables

8 C++ and python

Threads and async

7

Concurrency

- Threads and async
- Mutexes
- Atomic types
- Thread-local storage
- Condition Variables



Basic concurrency

C++ 11 Adv

Threading

- C++ 11 added `std::thread` in `<thread>` header
- takes a function as argument of its constructor
- must be detached or joined before the main thread terminates

Example code

```
1 void foo() {...}
2 void bar() {...}
3 int main() {
4     std::thread t1{foo};
5     std::thread t2{bar};
6     for (auto t: {&t1,&t2}) t->join();
7     return 0;
8 }
```

The thread constructor

C++ 11 Adv

Can take a function and its arguments

```
1 void function(int j, double j) {...}  
2 std::thread t1{function, 1, 2.0};
```



The thread constructor

Can take a function and its arguments

```
1 void function(int j, double j) {...}  
2 std::thread t1{function, 1, 2.0};
```

Can take any function-like object

```
1 struct AdderFunctor {  
2     AdderFunctor(int i): m_i(i) {}  
3     int operator() (int j) const { return m_i+j; }  
4     int m_i;  
5 };  
6 std::thread t2{AdderFunctor{2}, 5};  
7 int a;  
8 std::thread t3{[] (int i) { return i+2; }, a};  
9 std::thread t4{[a] { return a+2; }};
```

Good practice: Prefer std::jthread if C++ 20 is available

- In <thread> header since C++ 20
- Automatically joins on destruction
- Can be stopped from outside

Example with jthread

```
1 void foo() {...}
2 void bar() {...}
3 int main() {
4     std::jthread t1{foo};
5     std::jthread t2{bar};
6     // No join required
7     return 0;
8 }
```

Using jthread's stop_token

- Like `std::thread`, jthread can take a function or functors/lambdas and their arguments
- New: Can take functions that handle `stop_tokens`

stop_token in a function

```
1 void function(std::stop_token st, int j) {  
2     while (!st.stop_requested()) {  
3         // ...  
4     }  
5     std::jthread t1{function, 1};  
6     t1.request_stop();
```

A lambda with a stop_token

```
1 std::jthread t1{  
2     [](std::stop_token st, int i){ ... }, 1};
```



Basic asynchronicity

C++ 11 Adv

Concept

- separation of the specification of what should be done and the retrieval of the results
- “start working on this, and let me check if it’s done”



Basic asynchronicity

C++ 11 Adv

Concept

- separation of the specification of what should be done and the retrieval of the results
- “start working on this, and let me check if it’s done”

Practically

- `std::async` function launches an asynchronous task
- `std::future<T>` allows to retrieve the result



Basic asynchronicity

C++ 11 Adv

Concept

- separation of the specification of what should be done and the retrieval of the results
- “start working on this, and let me check if it’s done”

Practically

- `std::async` function launches an asynchronous task
- `std::future<T>` allows to retrieve the result

Example code

```
1 int f() {...}  
2 std::future<int> res = std::async(f);  
3 std::cout << res.get() << "\n";
```

Mixing the two

C++ 11 Adv

Is async running concurrently ?

- it depends!
- you can control this with a launch policy argument

`std::launch::async` start executing immediately on a separate
thread (may be a new thread, a thread pool, ...)

`std::launch::deferred` causes lazy execution in current thread

- execution starts when `get()` is called on the returned future
- default is not specified, but tries to use existing concurrency!

Mixing the two

C++ 11 Adv

Is async running concurrently ?

- it depends!
- you can control this with a launch policy argument

`std::launch::async` start executing immediately on a separate
thread (may be a new thread, a thread pool, ...)

`std::launch::deferred` causes lazy execution in current thread

- execution starts when `get()` is called on the returned future
- default is not specified, but tries to use existing concurrency!

Usage - godbolt

```
1 int f() {...}  
2 auto res1 = std::async(std::launch::async, f);  
3 auto res2 = std::async(std::launch::deferred, f);
```

Fine-grained control on asynchronous execution C++ 11 Adv

`std::packaged_task` template

- creates an asynchronous version of any function-like object
 - identical arguments
 - returns a `std::future`
- provides access to the returned future
- associated with threads, gives full control on execution



Fine-grained control on asynchronous execution C++ 11 Adv

std::packaged_task template

- creates an asynchronous version of any function-like object
 - identical arguments
 - returns a std::future
- provides access to the returned future
- associated with threads, gives full control on execution

Usage

```
1 int f() { return 42; }
2 std::packaged_task<int()> task{f};
3 auto future = task.get_future();
4 task();
5 std::cout << future.get() << std::endl;
```

Mutexes

7

Concurrency

- Threads and async
- **Mutexes**
- Atomic types
- Thread-local storage
- Condition Variables



Example code - godbolt

```
1 int a = 0;
2 void inc() { a++; };
3 void inc100() {
4     for (int i=0; i < 100; i++) inc();
5 };
6 int main() {
7     std::thread t1{inc100};
8     std::thread t2{inc100};
9     for (auto t: {&t1,&t2}) t->join();
10    std::cout << a << "\n";
11 }
```



Example code - godbolt

```
1 int a = 0;
2 void inc() { a++; };
3 void inc100() {
4     for (int i=0; i < 100; i++) inc();
5 };
6 int main() {
7     std::thread t1{inc100};
8     std::thread t2{inc100};
9     for (auto t: {&t1,&t2}) t->join();
10    std::cout << a << "\n";
11 }
```

What do you expect? (Exercise exercises/race)



Example code - godbolt

```
1 int a = 0;
2 void inc() { a++; };
3 void inc100() {
4     for (int i=0; i < 100; i++) inc();
5 };
6 int main() {
7     std::thread t1{inc100};
8     std::thread t2{inc100};
9     for (auto t: {&t1,&t2}) t->join();
10    std::cout << a << "\n";
11 }
```

What do you expect? (Exercise exercises/race)

Anything between 100 and 200 !!!

Atomicity

C++ 11 Adv

Definition (wikipedia)

- an operation (or set of operations) is atomic if it appears to the rest of the system to occur instantaneously

Practically

- an operation that won't be interrupted by other concurrent operations
- an operation that will have a stable environment during execution



Atomicity

C++ 11 Adv

Definition (wikipedia)

- an operation (or set of operations) is atomic if it appears to the rest of the system to occur instantaneously

Practically

- an operation that won't be interrupted by other concurrent operations
- an operation that will have a stable environment during execution

Is ++ operator atomic ?



Definition (wikipedia)

- an operation (or set of operations) is atomic if it appears to the rest of the system to occur instantaneously

Practically

- an operation that won't be interrupted by other concurrent operations
- an operation that will have a stable environment during execution

Is ++ operator atomic ?

Usually not. It behaves like :


```
1 eax = a          // memory to register copy
2 increase eax    // increase (atomic CPU instruction)
3 a = eax         // copy back to memory
```

Timing

Code

```
1 eax = a          // memory to register copy  
2 increase eax    // increase (atomic CPU instruction)  
3 a = eax         // copy back to memory
```

For 2 threads



Mutexes and Locks

C++ 17 Adv

Concept

- Use locks to serialize access to a non-atomic piece of code



Mutexes and Locks

C++ 17 Adv

Concept

- Use locks to serialize access to a non-atomic piece of code

The objects

`std::mutex` in the `mutex` header. **Mutual exclusion**

`std::scoped_lock` RAII to lock and unlock automatically

`std::unique_lock` same, but can be released/relocked explicitly



Mutexes and Locks

C++ 17 Adv

Concept

- Use locks to serialize access to a non-atomic piece of code

The objects

`std::mutex` in the `mutex` header. **Mutual exclusion**

`std::scoped_lock` RAII to lock and unlock automatically

`std::unique_lock` same, but can be released/relocked explicitly

Practically - godbolt

```
1 int a = 0;
2 std::mutex m;
3 void inc() {
4     std::scoped_lock lock{m};
5     a++;
6 }
```



Mutexes and Locks

C++ 17 Adv

Good practice: Locking

- Generally, use `std::scoped_lock`
 - Before C++ 17 use `std::lock_guard`.
- Hold as short as possible
 - Consider wrapping critical section in block statement `{ }`
- Only if manual control needed, use `std::unique_lock`

```
1 void function(...) {
2     // uncritical work ...
3     {
4         std::scoped_lock myLocks{mutex1, mutex2, ...};
5         // critical section
6     }
7     // uncritical work ...
8 }
```

Mutexes and Locks

Exercise: Mutexes and Locks

- Go to [exercises/race](#)
- Look at the code and try it
See that it has a race condition
- Use a mutex to fix the issue
- See the difference in execution time



Scenario

- 2 mutexes, 2 threads
- locking order different in the 2 threads




Dead locks

C++ 11 Adv

Scenario

- 2 mutexes, 2 threads
- locking order different in the 2 threads

Sequence diagram



How to avoid dead locks

C++ 11 Adv

Possible solutions

- C++ 17: `std::scoped_lock lock{m1, m2};` comes with deadlock-avoidance algorithm
- Never take several locks
 - Or add master lock protecting the locking phase
- Respect a strict locking order across all threads
- Do not use locks
 - Use other techniques, e.g. lock-free queues



Shared mutex / locks

Sharing a mutex

- Normal std::mutex objects cannot be shared
- std::shared_mutex to the rescue, but can be slower
- It locks *either* in exclusive or in shared mode; never both

```
1 Data data; std::shared_mutex mutex;
2 auto reader = [&](){
3     std::shared_lock lock{mutex};
4     read(data); // Many can read
5 };
6 std::thread r1{reader}, r2{reader}, ...;
7
8 std::thread writer([&](){
9     std::scoped_lock lock{mutex}; // exclusive
10    modify(data); // Only one can write
11});
```



Synchronised cout

Parallel writing to streams

- Parallel writes to streams become garbled when not synchronised
- `std::osyncstream` wraps and synchronises output to streams
- Multiple instances of the `osyncstream` synchronise globally

```
1 void worker(const int id, std::ostream & os);
2 void func() {
3     std::osyncstream synccout1{std::cout};
4     std::osyncstream synccout2{std::cout};
5     std::jthread t1{worker, 0, std::ref(synccout1)};
6     std::jthread t2{worker, 1, std::ref(synccout2)};
7 }
```

Atomic types

7

Concurrency

- Threads and async
- Mutexes
- Atomic types
- Thread-local storage
- Condition Variables



Atomic types in C++

C++ 11 Adv

std::atomic template

- Any trivially copyable type can be made atomic in C++
- Most useful for integral types
- May internally use locks for custom types

```
1 std::atomic<int> a{0};  
2 std::thread t1([&](){ a++; });  
3 std::thread t2([&](){ a++; });  
4 a += 2;  
5 t1.join(); t2.join();  
6 assert( a == 4 ); // Guaranteed to succeed
```




Warning: Expressions using an atomic type are *not* atomic!

```
1 std::atomic<int> a{0};  
2 std::thread t1([&]{ a = a + 2; });  
3 std::thread t2([&]{ a = a + 2; });
```

- Atomic load; value+2; atomic store

Sequence diagram




Solution: Use atomic member functions

- The member functions of `std::atomic` are thread safe
- `fetch_add` and `operator+=`: `atomic { load; add; store }`
- But don't confuse "`a += 2`" and "`a = a + 2`"

```
1 std::atomic<int> a{0};  
2 std::thread t1([&]{ a.fetch_add(2); });  
3 std::thread t2([&]{ a += 2; });
```

Sequence diagram



Atomic references

C++ 20 Adv

std::atomic_ref template

- Wraps a `T&` and makes access to it atomic
- Like `std::atomic<T>`, but does not contain the `T`

```
1 int a{0};  
2 std::thread t1([&]{ std::atomic_ref<int> r{a}; r++;});  
3 std::thread t2([&]{ std::atomic_ref{a}++; });  
4 t1.join(); t2.join();  
5 a += 2; // non-atomic (fine, threads joined before)  
6 assert( a == 4 ); // Guaranteed to succeed
```

Don't mix concurrent atomic and non-atomic access

```
1 std::thread t3([&]{ std::atomic_ref{a}++; });  
2 a += 2; // data race  
3 t3.join();
```



Atomic types in C++

Exercise: Atomics

- Go to exercises/atomic
- You'll find a program with the same race condition as in race
- Fix it using std::atomic



Thread-local storage

7

Concurrency

- Threads and async
- Mutexes
- Atomic types
- Thread-local storage
- Condition Variables



Thread-local storage

C++ 20 Adv

thread_local keyword

- A variable can be declared **thread-local**
- Then every thread will have its own copy
- Most useful for “working memory” in each thread
- Note: Can also be declared directly on the stack of a thread
 - and will be faster, thus should be preferred when possible

```
1 thread_local int a{0};  
2 {  
3     std::jthread t1([&] { a++; });  
4     std::jthread t2([&] { a++; });  
5     a += 2;  
6 }  
7 assert( a == 2 ); // Guaranteed to succeed
```



Condition Variables

7

Concurrency

- Threads and async
- Mutexes
- Atomic types
- Thread-local storage
- Condition Variables

Condition variables

C++ 11 Adv

Communicating between threads

- Take the case where threads are waiting for other thread(s)
- `std::condition_variable`
 - from `<condition_variable>` header
- Allows for a thread to sleep (= conserve CPU time) until a given condition is satisfied



Condition variables

C++ 11 Adv

Communicating between threads

- Take the case where threads are waiting for other thread(s)
- `std::condition_variable`
 - from `<condition_variable>` header
- Allows for a thread to sleep (= conserve CPU time) until a given condition is satisfied

Usage

- Use RAII-style locks to protect shared data
- `wait()` will block until the condition is met
 - you can have several waiters sharing the same mutex
- `notify_one()` will wake up one waiter
- `notify_all()` will wake up all waiters

Using condition variables: notify

C++ 17 Adv

Producer side: providing data to waiting threads

- Protect data with a mutex, and use condition variable to notify consumers
- Optimal use: don't hold lock while notifying
 - waiting threads would be blocked

```
1 std::mutex mutex;
2 std::condition_variable cond;
3 Data data;
4 std::thread producer([&](){
5 {
6     std::scoped_lock lock{mutex};
7     data = produceData(); // may take long ...
8 }
9 cond.notify_all();
10});
```



Using condition variables: wait

C++ 11 Adv

Mechanics of wait

- Many threads can wait for shared data
- Pass `unique_lock` and optional predicate to `wait()`
- `wait()` keeps threads asleep while predicate is `false`
 - Threads might wake up spuriously, but `wait()` returns only when lock available *and* predicate `true`
- `wait()` will only lock when necessary; unlocked while sleeping

Naïve waiting

```
1 auto processData = [&](){
2     std::unique_lock<std::mutex> lock{mutex};
3     cond.wait(lock, [&](){ return data.isReady(); });
4     process(data);
5 };
```

Using condition variables: wait

C++ 11 Adv

Waiting / waking up

- When `notify_all()` is called, threads wake up
- Threads try to lock mutex, and evaluate predicate
- One thread succeeds to acquire mutex, starts processing data
- Problem:** One thread holds mutex, other threads are blocked!

Naïve waiting

```
1 auto processData = [&](){
2     std::unique_lock<std::mutex> lock{mutex};
3     cond.wait(lock, [&](){ return data.isReady(); });
4     process(data);
5 };
```

Using condition variables: correct wait

C++ 11 Adv

Waiting / waking up

- **Solution:** Put locking and waiting in a scope
- Threads will one-by-one wake up, acquire lock, evaluate predicate, release lock

Correct waiting

```
1 auto processData = [&] (){  
2     {  
3         std::unique_lock<std::mutex> lock{mutex};  
4         cond.wait(lock, [&](){ return data.isReady(); });  
5     }  
6     process(data);  
7 };  
8 std::thread t1{processData}, t2{processData}, ...;  
9 for (auto t : {&producer, &t1, &t2, ...}) t->join();
```

Condition variables

Exercise: Condition variables

- Go to `exercises/condition_variable`
- Look at the code and run it
See that it has a race condition
- Fix the race condition in the usage of the condition variable
- Try to make threads process data in parallel



C++ and python

1 History and goals

2 Language basics

3 Object orientation (OO)

4 Core modern C++

5 Expert C++

6 Useful tools

7 Concurrency

8 C++ and python

- Writing a module
- Marrying C++ and C
- The ctypes module
- The cppyy project



Writing a module

8

C++ and python

- Writing a module
- Marrying C++ and C
- The ctypes module
- The cppyy project



How to build a python 3 module around C++ code

C++ code : mandel.hpp

```
1 int mandel(Complex const & a);
```



Basic Module(1): wrap your method

mandelModule.cpp - see exercises/python exercise

```
1 #include <Python.h>
2 #include "mandel.hpp"
3 PyObject * mandel_wrapper(PyObject * self,
4                             PyObject * args) {
5     // Parse Input
6     float r, i;
7     if (!PyArg_ParseTuple(args, "ff", &r, &i))
8         return nullptr;
9     // Call C++ function
10    int result = mandel(Complex(r, i));
11    // Build returned objects
12    return PyLong_FromLong(result);
13 }
```

Basic Module(2): create the python module

mandelModule.cpp - see exercises/python exercise

```
1 // declare the modules' methods
2 PyMethodDef mandelMethods[] = {
3     {"mandel", mandel_wrapper, METH_VARARGS,
4      "computes nb of iterations for mandelbrot set"},
5     {nullptr, nullptr, 0, nullptr}
6 };
7 // declare the module
8 struct PyModuleDef mandelModule = {
9     PyModuleDef_HEAD_INIT,
10    "mandel", nullptr, -1, mandelMethods
11 };
12 PyMODINIT_FUNC PyInit_mandel() {
13     return PyModule_Create(&mandelModule);
14 }
```

Basic Module(3): use it

First compile the module

- as a regular shared library
- with '-I \\$PYTHON_INCLUDE'

mandel.py - see exercises/python exercise

```
from mandel import mandel
v = mandel(0.7, 1.2)
```



Marrying C++ and C

8

C++ and python

- Writing a module
- **Marrying C++ and C**
- The ctypes module
- The cppyy project



A question of mangling

Mangling

the act of converting the name of variable or function to a symbol name in the binary code

C versus C++ symbol names

- C uses bare function name
- C++ allows overloading of functions by taking the signature into account
- so C++ mangling has to contain signature



C mangling

Source : file.c

```
1 float sum(float a, float b);
2 int square(int a);
3 // won't compile : conflicting types for 'square'
4 // float square(float a);
```

Binary symbols : file.o

```
# nm file.o
0000000000000001a T square
0000000000000000 T sum
```



C++ mangling

Source : file.cpp

```
1 float sum(float a, float b);
2 int square(int a);
3 // ok, signature is different
4 float square(float a);
```

Binary symbols : file.o

```
# nm file.o
0000000000000000 T _Z3sumff
0000000000000002a T _Z6squaref
0000000000000001a T _Z6squarei
```



Forcing C mangling in C++

extern "C"

These functions will use C mangling :

```
1 extern "C" {  
2     float sum(float a, float b);  
3     int square(int a);  
4 }
```



Forcing C mangling in C++

extern "C"

These functions will use C mangling :

```
1 extern "C" {  
2     float sum(float a, float b);  
3     int square(int a);  
4 }
```

You can now call these C++ functions from C code



Forcing C mangling in C++

extern "C"

These functions will use C mangling :

```
1 extern "C" {  
2     float sum(float a, float b);  
3     int square(int a);  
4 }
```

You can now call these C++ functions from C code

Limitations

- no C++ types should go out
- no exceptions either (use `noexcept` here)
- member functions cannot be used
 - they need to be wrapped one by one

The ctypes module

8

C++ and python

- Writing a module
- Marrying C++ and C
- **The ctypes module**
- The cppyy project



The ctypes python module

From the documentation

- provides C compatible data types
- allows calling functions in DLLs or shared libraries
- can be used to wrap these libraries in pure Python



ctypes: usage example

C++ code : mandel.hpp

```
1 int mandel(Complex const & a);
```

"C" code : mandel_cwrapper.hpp

```
1 extern "C" {
2     int mandel(float r, float i) {
3         return mandel(Complex(r, i));
4     };
5 }
```

calling the mandel library

```
from ctypes import *
libmandel = CDLL('libmandelc.so')
v = libmandel.mandel(c_float(0.3), c_float(1.2))
```

Marrying C++ and python

Exercise: C++ and python

- go to exercises/python
- look at the original python code mandel.py
- time it ('time python3 mandel.py')
- look at the code in mandel.hpp/cpp
- look at the python module mandel_module.cpp
- compile and modify mandel.py to use it
- see the gain in time
- look at the C wrapper in mandel_cwrapper.cpp
- modify mandel.py to use libmandelc directly with ctypes

Note : you may have to add '.' to LD_LIBRARY_PATH and PYTHONPATH



The cppyy project

8 C++ and python

- Writing a module
- Marrying C++ and C
- The ctypes module
- The cppyy project



Automatic Python-C++ bindings

The [cppyy](#) project

- originated from the ROOT project
- still young, version 1.0 from Mid 2018
- but very active, current version 2.1.0
- extremely powerful for interfacing C++ and python

How it works

- uses Just in Time compilation through [cling](#)
 - an interactive C++ interpreter



cppyy crash course(1)

Shamelessly copied from the cppyy documentation

```
>>> import cppyy
>>> cppyy.include('zlib.h')           # bring in C++ definitions
>>> cppyy.load_library('libz')        # load linker symbols
>>> cppyy.gbl.zlibVersion()         # use a zlib API
'1.2.11'

>>> import cppyy
>>> cppyy.cppdef("""
... class MyClass {
... public:
...     MyClass(int i) : m_data(i) {}
...     virtual ~MyClass() {}
...     virtual int add_int(int i) { return m_data + i; }
...     int m_data;
... };""")
True
```



cppyy crash course(1)

```
>>> from cppyy.gbl import MyClass
>>> m = MyClass(42)
>>> cppyy.cppdef("""
... void say_hello(MyClass* m) {
...     std::cout << "Hello, the number is: " << m->m_data << std::endl;
... }""")
True
>>> MyClass.say_hello = cppyy.gbl.say_hello
>>> m.say_hello()
Hello, the number is: 42
>>> m.m_data = 13
>>> m.say_hello()
Hello, the number is: 13
```



This is the end

Questions ?

https://github.com/hsf-training/cpluspluscourse/raw/download/talk/C++Course_full.pdf
<https://github.com/hsf-training/cpluspluscourse>



Index of Good Practices

① C's memory management	27	⑯ Operator overloading	172
② Manual memory management	28	⑰ Throwing exceptions	191
③ Initialisation	31	⑱ Catching exceptions	193
④ Avoid unions	38	⑲ Exceptions	197
⑤ References	45	⑳ noexcept	200
⑥ Write readable functions	54	㉑ Specialization vs. overloading	238
⑦ for syntax	67	㉒ STL and lambdas	271
⑧ preprocessor	75	㉓ C random library	302
⑨ Header / source	77	㉔ Use std::fstream	313
⑩ Assert	87	㉕ std::unique_ptr	319
⑪ Implementing methods	94	㉖ Single responsibility principle	325
⑫ Rule of 3/5	109	㉗ Rule of zero	325
⑬ Prefer smart pointer	127	㉘ Initialization	335
⑭ Prefer containers	128	㉙ SFINAE vs. if constexpr	386
⑮ Virtual destructors	141	㉚ jthread vs. thread	510
⑯ Avoid multiple inheritance	148	㉛ Locking	520



Index of Exercises

① Functions	53	⑫ std::variant	289	㉓ valgrind	487
② Control structs	72	⑬ Smart pointers	326	㉔ memcheck	488
③ Loops, refs, auto	80	⑭ Variadic tpl	357	㉕ callgrind	490
④ Polymorphism	144	⑮ Concepts	399	㉖ helgrind	492
⑤ Virtual OO	149	⑯ Modules	424	㉗ cppcheck	495
⑥ Operators	161	⑰ Header units	424	㉘ clang-tidy	498
⑦ Constness	179	⑲ clang-format	452	㉙ Mutexes/Locks	521
⑧ Move semantics	220	⑳ Compiler chain	462	㉚ Atomics	531
⑨ Templates	242	㉑ gdb	471	㉛ Condition vars	539
⑩ STL	276	㉒ address sanitizer	480	㉜ C++ and python	554
⑪ std::optional	284	㉓ Finding evil run-time bugs	483		

Index of Godbolt code examples

① assert - godbolt	86	⑳ shared_ptr quiz - godbolt	322
② Unary constructor - godbolt	111	㉑ shared/weak_ptr - godbolt	324
③ casting - godbolt	152	㉒ <=> output - godbolt	341
④ function objects - godbolt	164	㉓ various orderings - godbolt	343
⑤ Nested templates - godbolt	230	㉔ tuple code - godbolt	353
⑥ CTAD - godbolt	239	㉕ hasFoo failing - godbolt	378
⑦ lambda - godbolt	257	㉖ hasFoo working - godbolt	380
⑧ lambda C++ 20 - godbolt	258	㉗ enable_if - godbolt	383
⑨ STL - godbolt	261	㉘ SFINAE failure - godbolt	389
⑩ Iterator example - godbolt	267	㉙ concept failure - godbolt	390
⑪ span - godbolt	281	㉚ Coroutine resumption - godbolt	433
⑫ expected - godbolt	285	㉛ co_yield - godbolt	437
⑬ variant visitor - godbolt	287	㉜ Infinite generator - godbolt	438
⑭ variant λ visitor - godbolt	288	㉝ Fibonacci generator - godbolt	439
⑮ Structured binding - godbolt	294	㉞ Thread switching - godbolt	442
⑯ timing - godbolt	297	㉟ std::async - godbolt	513
⑰ Views - godbolt	305	㉟ Race - godbolt	516
⑱ Lazy view - godbolt	307	㉟ mutex - godbolt	519
⑲ unique_ptr copy - godbolt	315	㉟ atomic - godbolt	520



Books



A Tour of C++, Third Edition

Bjarne Stroustrup, Addison-Wesley, Sep 2022
ISBN-13: 978-0136816485



Effective Modern C++

Scott Meyers, O'Reilly Media, Nov 2014
ISBN-13: 978-1491-90399-5



C++ Templates - The Complete Guide, 2nd Edition

David Vandevoorde, Nicolai M. Josuttis, and Douglas Gregor
ISBN-13: 978-0321-71412-1



C++ Best Practices, 2nd Edition

Jason Turner
<https://leanpub.com/cppbestpractices>



Clean Architecture

Robert C. Martin, Pearson, Sep 2017
ISBN-13: 978-013449416-6



The Art of UNIX Programming

Eric S. Raymond, Addison-Wesley, Sep 2002
ISBN-13: 978-0131429017



Introduction to Algorithms, 4th Edition

T. H. Cormen, C. E. Leiserson, R. L. Rivest, C. Stein, Apr 2022
ISBN-13: 978-0262046305

Conferences

- CppCon — cppcon.org —  CppCon
- C++ Now — cppnow.org —  BoostCon
- Code::Dive — codedive.pl —  codediveconference
- ACCU Conference — accu.org —  ACCUConf
- Meeting C++ — meetingcpp.com —  MeetingCPP
- See link below for more information
<https://isocpp.org/wiki/faq/conferences-worldwide>